



# Master in Computer Vision *Barcelona*

**Module: 1**

**Lecture 8:** Feature extraction

**Lecturer:** Verónica Vilaplana

# Motivation

- Two pairs of images to be matched. What kind of feature might one use to establish a set of correspondences between these images?



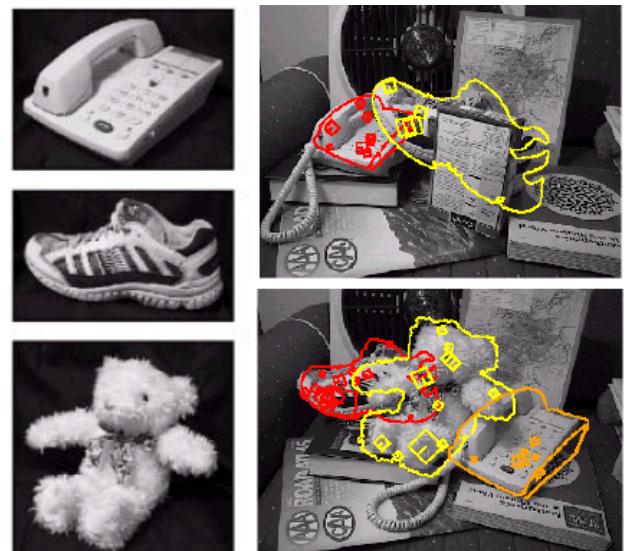
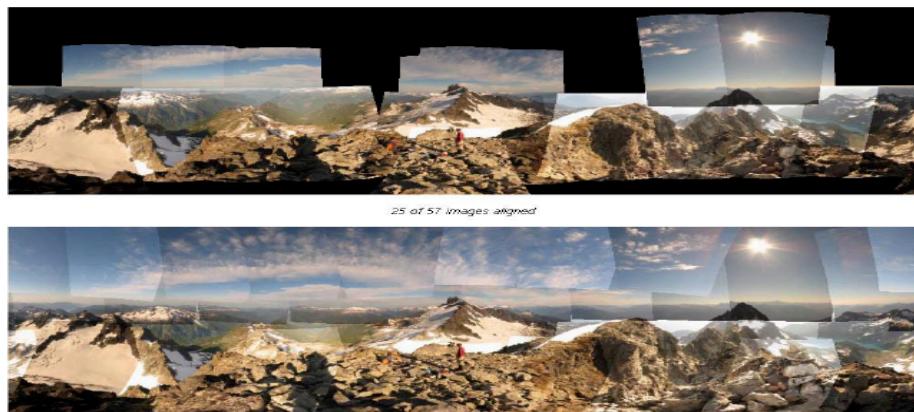
**Specific locations (keypoints):**  
mountain peaks,  
building corners,  
some patches of  
snow



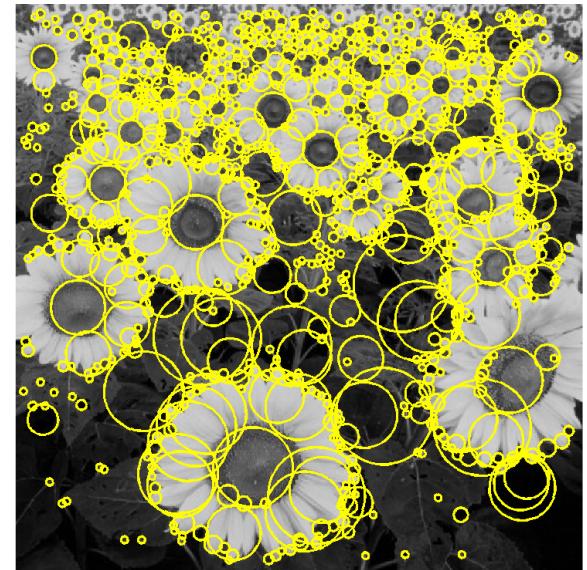
**Edges:** profile of mountains, window frame

# More applications

- Image retrieval and indexing
- Panorama stitching
- Object detection
- Automate object tracking
- Point matching for computing disparity
- Stereo calibration: estimation of the fundamental matrix
- Motion-based segmentation
- 3D object reconstruction
- Robot navigation
- Structure from motion....



# Local features: edges, corners, blobs



# Local Features:

- Edges
  - Image gradient. Effect of noise
  - Derivative of Gaussian
  - Laplacian of Gaussian
  - Canny detector
- Corners
  - Local invariant features. Requirements
  - Harris detector
- Blobs (regions)
  - Scale invariant region selection
  - Automatic scale selection
  - Laplacian of Gaussian (LoG)
  - Difference of Gaussian (DoG)
- Local descriptors
  - SIFT



# Edge detection

- **Goal:** identify sudden changes (discontinuities) in an image
  - Intuitively, most semantic and shape information from the image can be encoded in the edges
  - More compact representation than pixels

Ideal:

artist's line drawing

(artist is also using object-level knowledge)



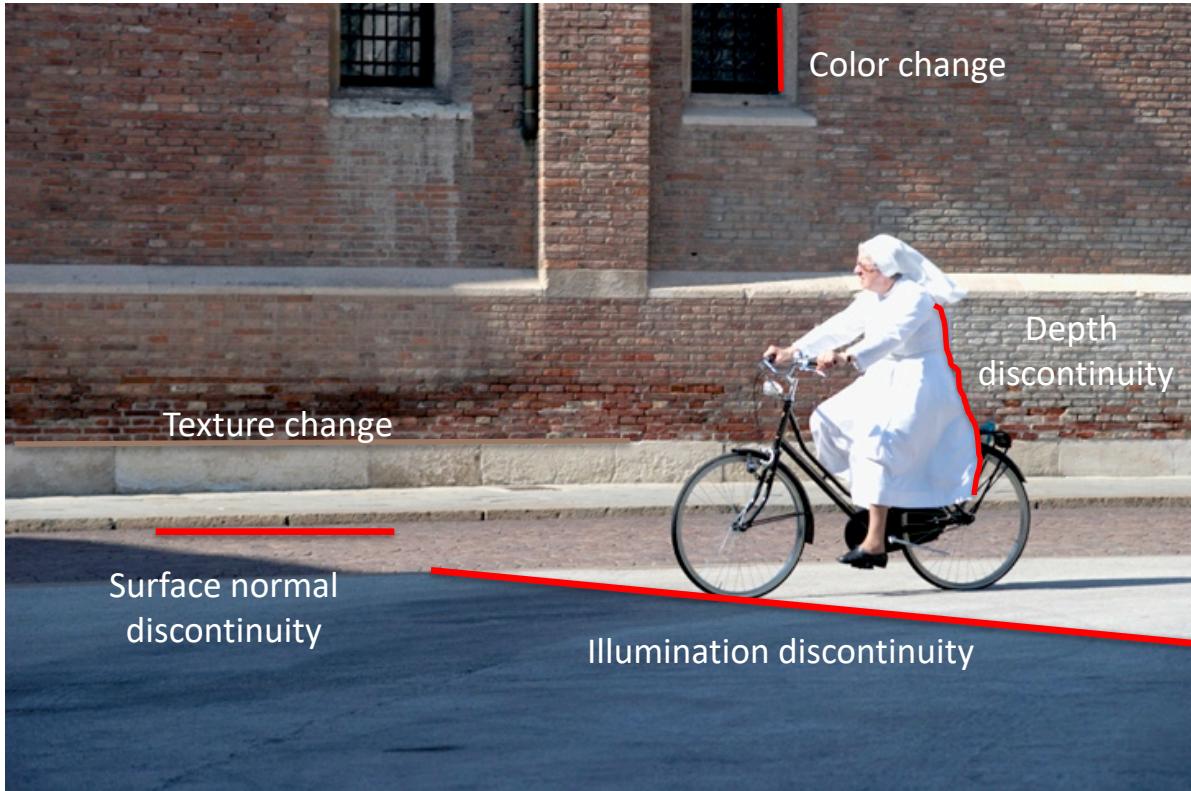
Reality:



Source: Lazebnik

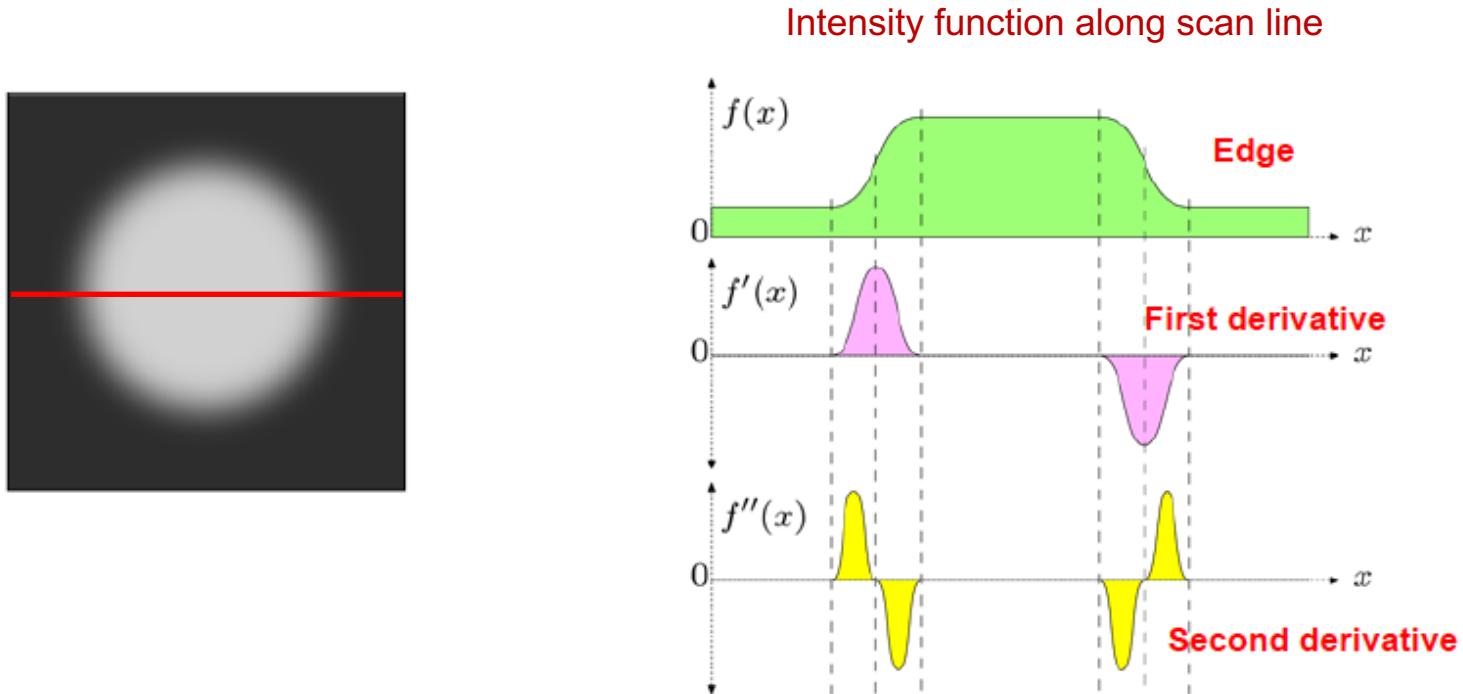
# Edge detection

- **Origin of edges:** edges are caused by a variety of factors: depth discontinuity, surface color discontinuity, surface normal discontinuity, illumination discontinuity, texture discontinuity



# Characterizing edges

- **What's an edge?** A rapid change in the image intensity
- How can we find large changes in intensity?



- Edges can be characterized by high abs values of the first derivative (**gradient**)
- Or by zero-crossings of the second derivative (**laplacian**)



# Image Gradient

Edge points = extreme points of the gradient

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f(x, y)}{\partial x} & \frac{\partial f(x, y)}{\partial y} \end{bmatrix}$$

The gradient of an image  
points in the direction of most rapid increase in intensity

$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$

$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

Edge strength is given by the magnitude:

$$|\nabla f| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

Gradient direction (orthogonal to edge direction)

$$\phi = \arctg \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

# Image Gradient

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f(x, y)}{\partial x} & \frac{\partial f(x, y)}{\partial y} \end{bmatrix}$$

Discrete case:

$$\nabla f[m, n] = \begin{bmatrix} f * h_x^s & f * h_y^s \end{bmatrix} = \begin{bmatrix} g_x[m, n] & g_y[m, n] \end{bmatrix}$$

Prewitt

$$h_x = \begin{bmatrix} 1 & 0 & -1 \\ 1 & \textcolor{red}{0} & -1 \\ 1 & 0 & -1 \end{bmatrix} \quad h_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & \textcolor{red}{0} & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

Sobel

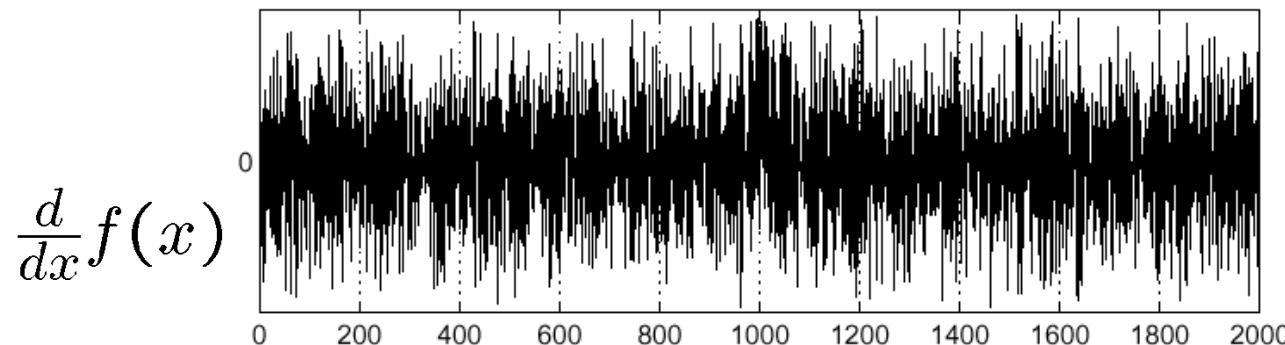
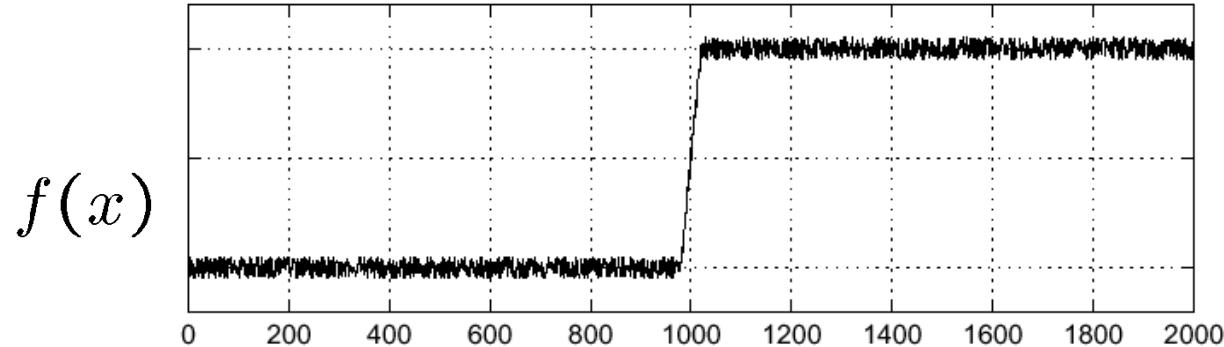
$$h_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & \textcolor{red}{0} & 2 \\ -1 & 0 & 1 \end{bmatrix} \quad h_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & \textcolor{red}{0} & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Roberts

$$h_x = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix} \quad h_y = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$

# Effects of noise

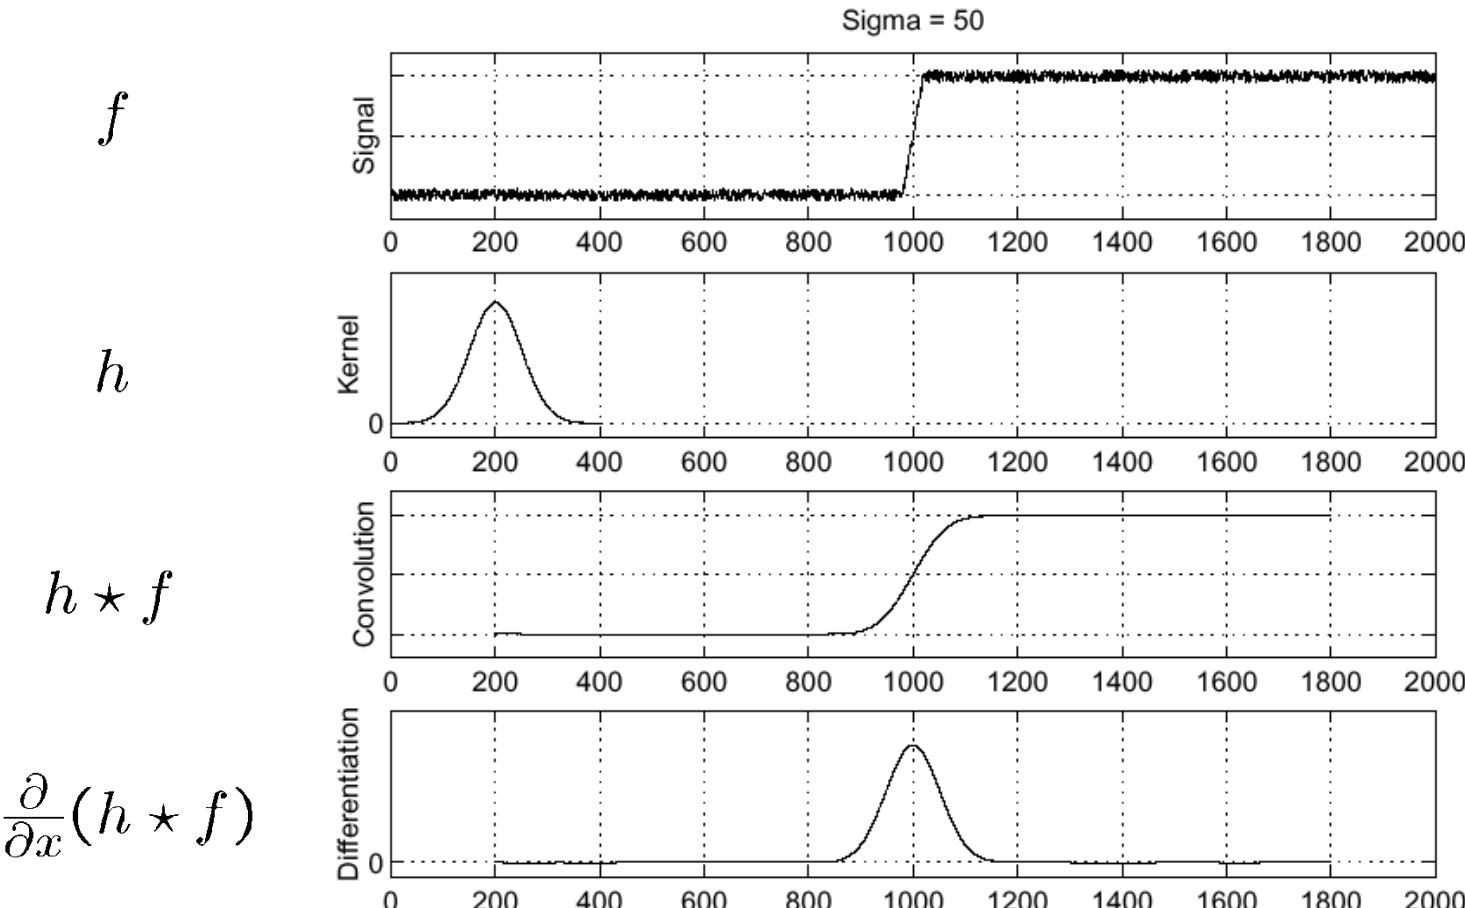
- Consider a single row or column of the image
  - Plotting intensity as a function of position gives a signal



- Where is the edge?

Source: S. Seitz

# Solution: smooth first



- Where is the edge? Look for peaks in

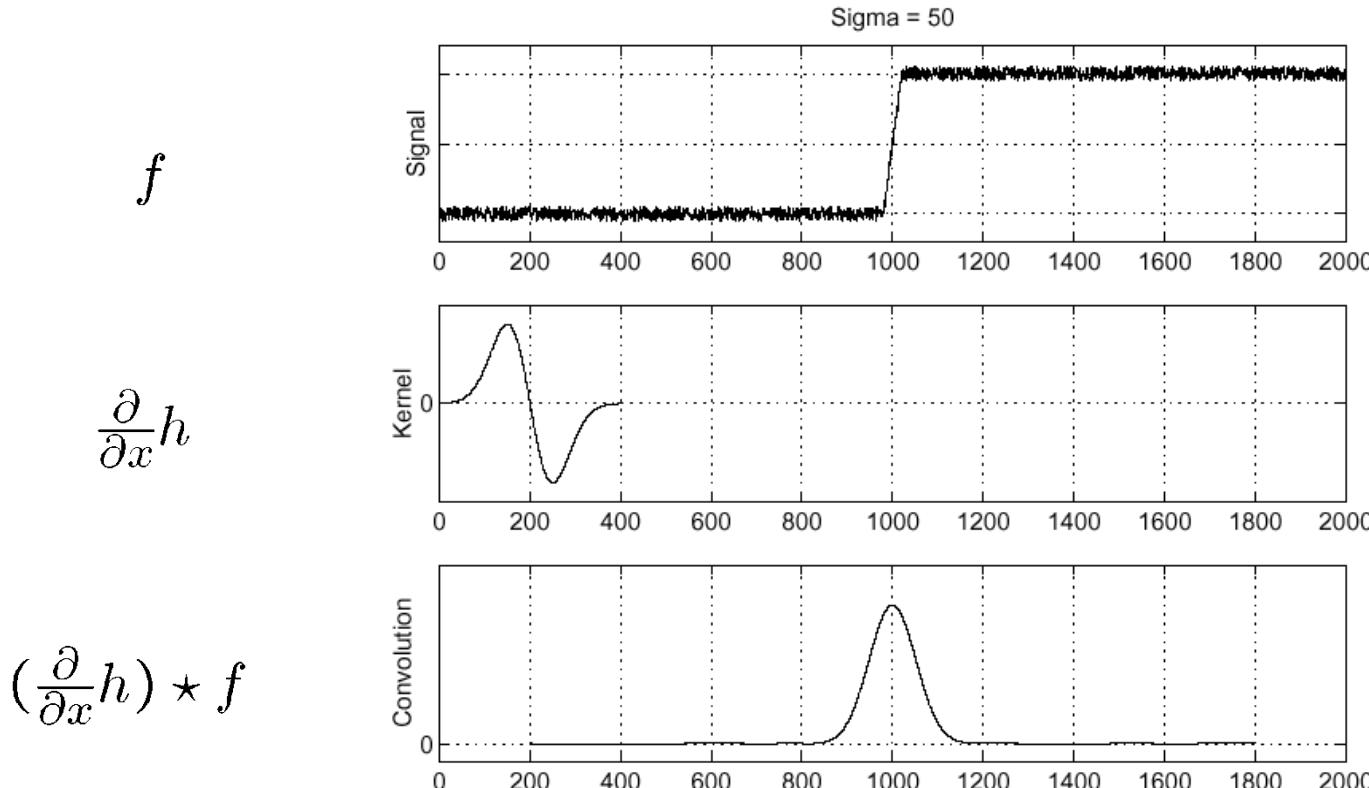
$$\frac{\partial}{\partial x}(h ∗ f)$$

Source: S. Seitz

# Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

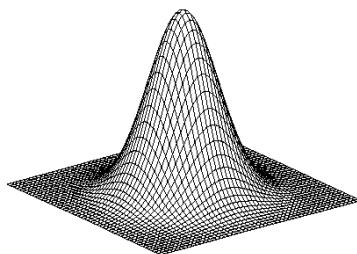
- This saves us one operation



Source: S. Seitz

# Derivative of Gaussian filter

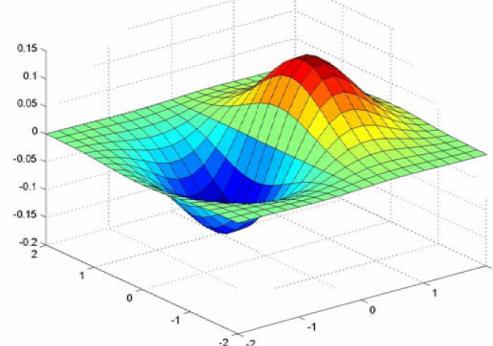
Gaussian



$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

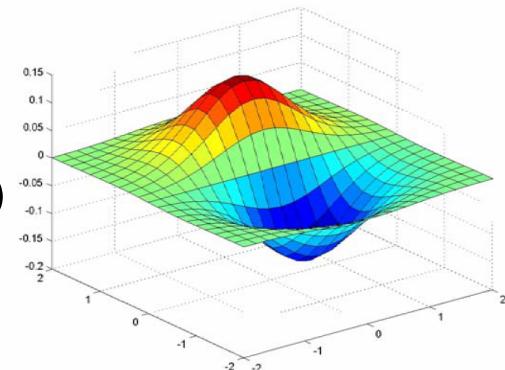
Derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$



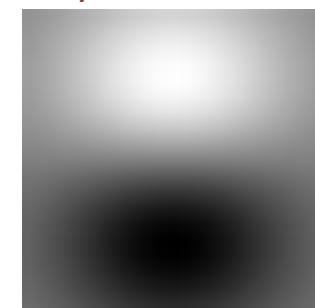
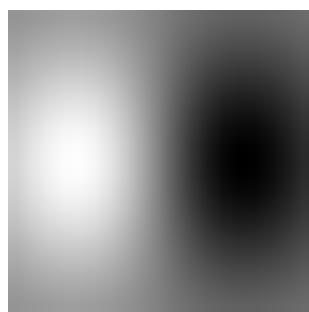
x direction

$$\frac{\partial}{\partial y} h_\sigma(u, v)$$

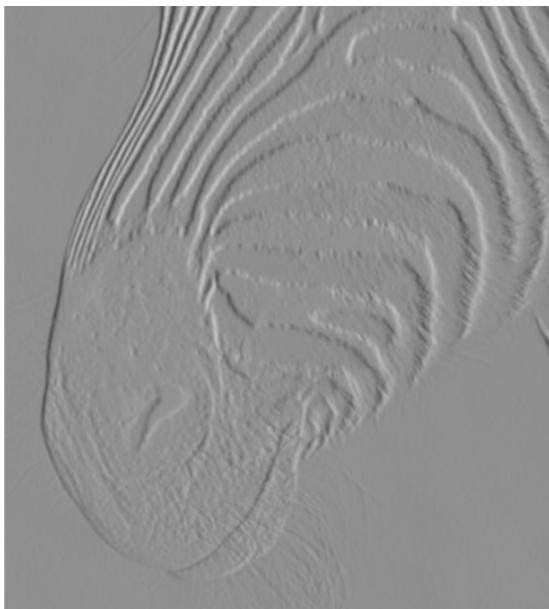


y direction

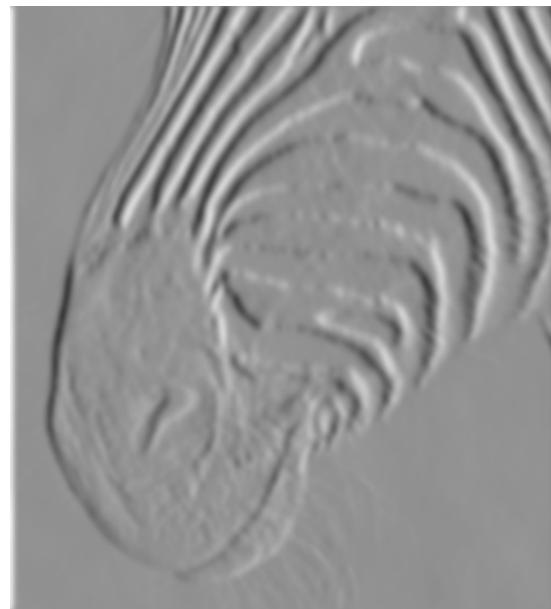
Which one finds horizontal/vertical edges?



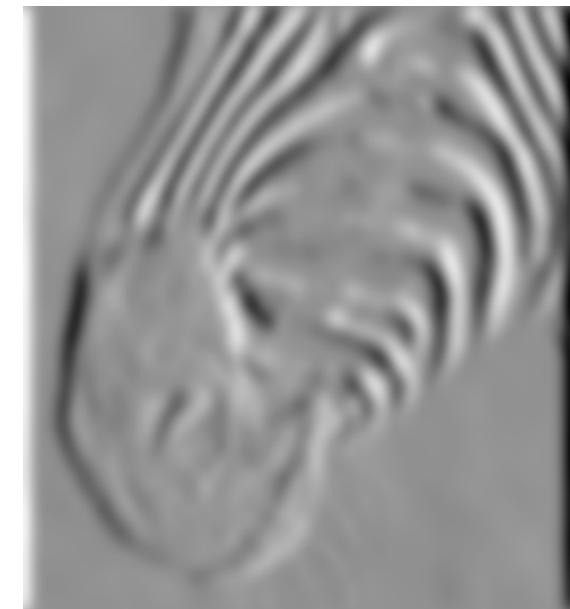
# Tradeoff between smoothing and localization



$\sigma = 1$  pixel



$\sigma = 3$  pixels



$\sigma = 7$  pixels

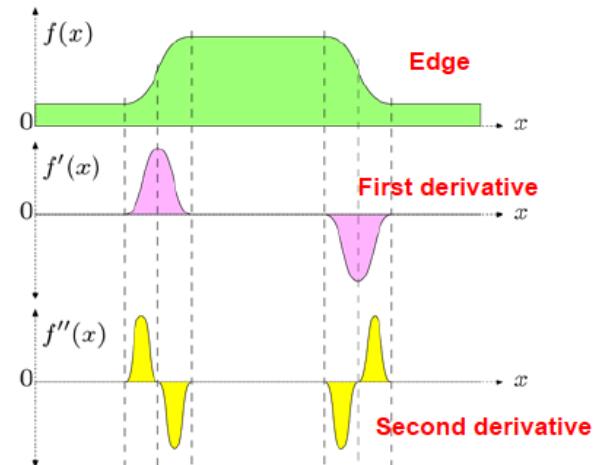
- Smoothed derivative removes noise, but blurs edges making it more difficult to localize
- The scale of the smoothing filter affects derivative estimates, and also the semantics of the edges recovered.

Source: D Forsyth

# Image Laplacian

- Edge points = zero crossings of Laplacian

$$\Delta f(x, y) = \nabla^2 f(x, y) = \frac{\partial^2 f}{\partial x^2}(x, y) + \frac{\partial^2 f}{\partial y^2}(x, y)$$



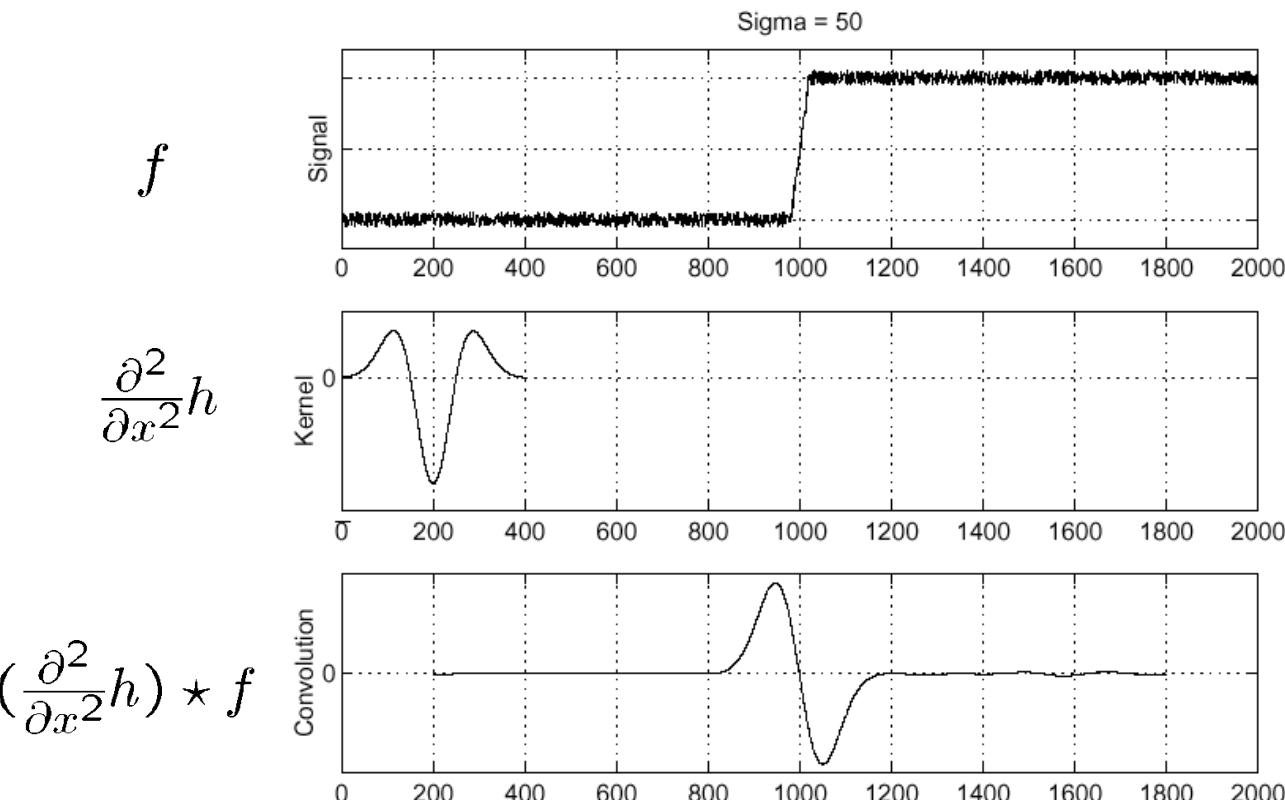
Discrete case:  $\Delta f[m, n] = f * h_L$

**Laplacian**

$$h_L = \left[ \begin{array}{c|c|c} 0 & -1 & 0 \\ \hline -1 & 4 & -1 \\ \hline 0 & -1 & 0 \end{array} \right] \quad h_L = \left[ \begin{array}{c|c|c} -1 & -1 & -1 \\ \hline -1 & 8 & -1 \\ \hline -1 & -1 & -1 \end{array} \right]$$

# Smooth first

- Consider  $\frac{\partial^2}{\partial x^2}(h \star f)$  h Gaussian filter



Laplacian of Gaussian operator

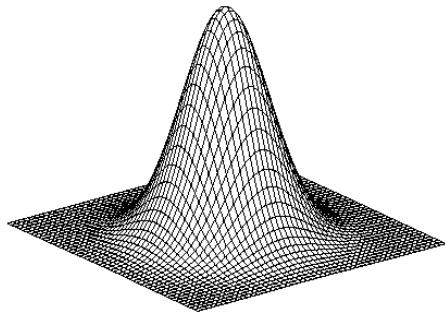
- Where is the edge? Zero-crossings of

$$(\frac{\partial^2}{\partial x^2} h) \star f$$

# Laplacian of Gaussian filter

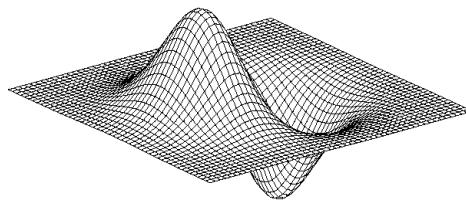
$\nabla^2$  is the **Laplacian** operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$



Gaussian

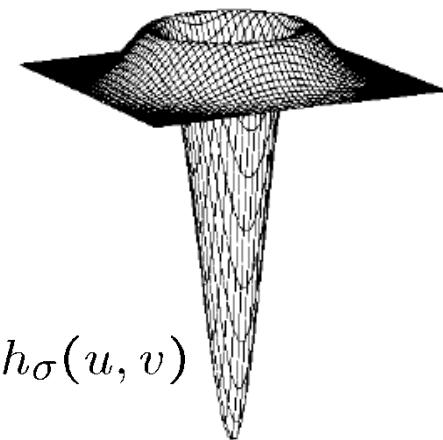
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



Derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian



$$\nabla^2 h_\sigma(u, v)$$

- Laplacian of Gaussian computes 2nd derivatives using one rotationally symmetric filter, it responds to edges at all orientations.
- **Interesting:** The Difference of two Gaussians can be a very close approximation to the Laplacian of Gaussian... (we will use this later)



# Edge detection

Example: Sobel



Original image

$$h_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & \textcolor{red}{0} & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

Vertical  
contours



$$g_1 = h_1 * f$$

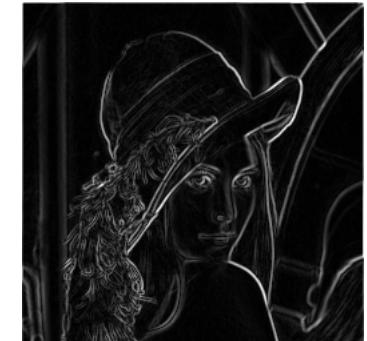
$$h_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & \textcolor{red}{0} & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Horizontal  
contours



$$g_2 = h_2 * f$$

Combined information: gradient



$$\begin{aligned} |\nabla f| &\cong \sqrt{g_1^2 + g_2^2} \\ |g_x| + |g_y| \end{aligned}$$

Straightforward strategy: gradient binarization

Issues to consider:

- Difficulty to define the threshold
- Noise may appear
- Contours may not be closed → [How to link edge points to form curves?](#)
- Contours may be thick → [How do we identify the actual edge points](#)

# Canny edge detector

- The most widely used edge detector in computer vision
- Theoretical model: Step-edges corrupted by additive Gaussian noise
- Gives single-pixel-wide images with good continuation between adjacent pixels

## Steps

1. **Noise reduction:** filter image with derivative of Gaussian
2. Find magnitude and orientation of **gradient**
3. **Thinning:** Non-maximum suppression
  - thin multi-pixel wide ridges down to single pixel width
4. **Linking and hysteresis thresholding**
  - Track high magnitude contours, removing weak small segments
  - Use a high threshold to start edge curves and a low threshold to continue them

J. Canny, *A Computational Approach To Edge Detection*, IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

# Example

## Compute Gradients

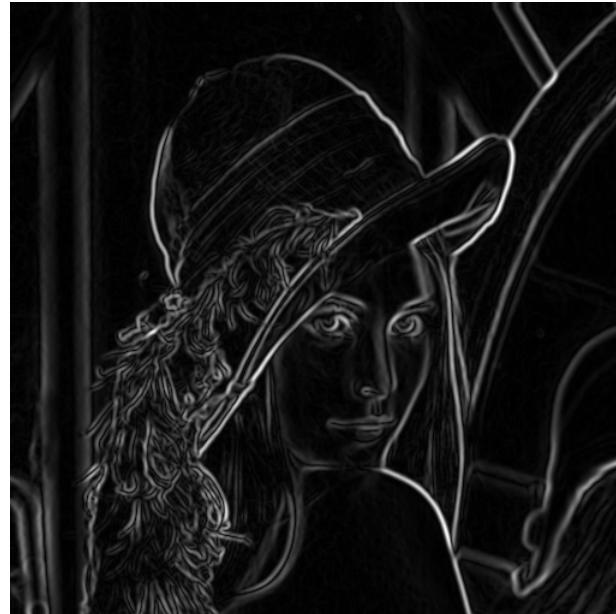
- noise-> derivative of Gaussian
- thick contours-> non-maximum suppression
- broken lines-> hysteresis thresholding



X-Derivative of Gaussian



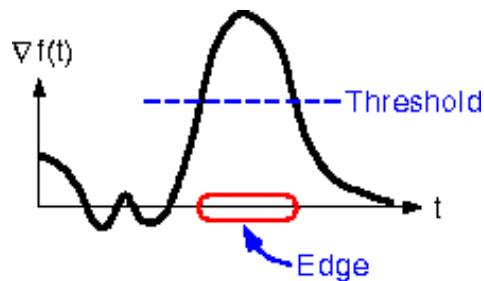
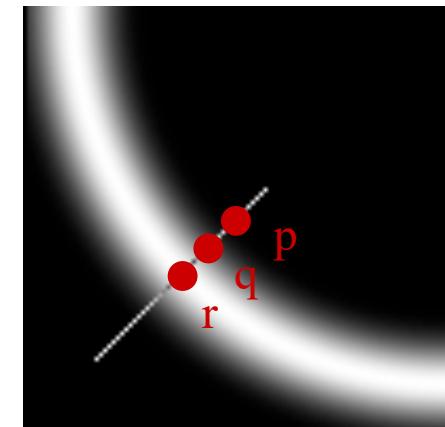
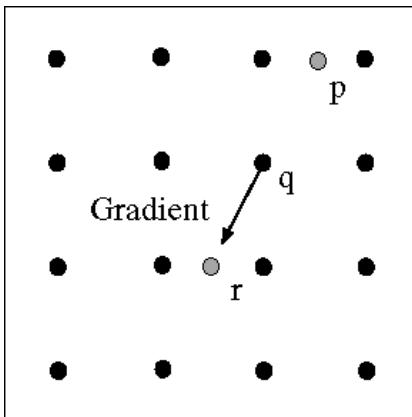
Y-Derivative of Gaussian



Gradient Magnitude

# Non-maximum suppression (thinning)

- How to turn thick regions of the gradient into curves?
- Check if a pixel is a local maximum along gradient direction



Keep q as an edge pixel only if it has the maximum gradient magnitude along the gradient direction

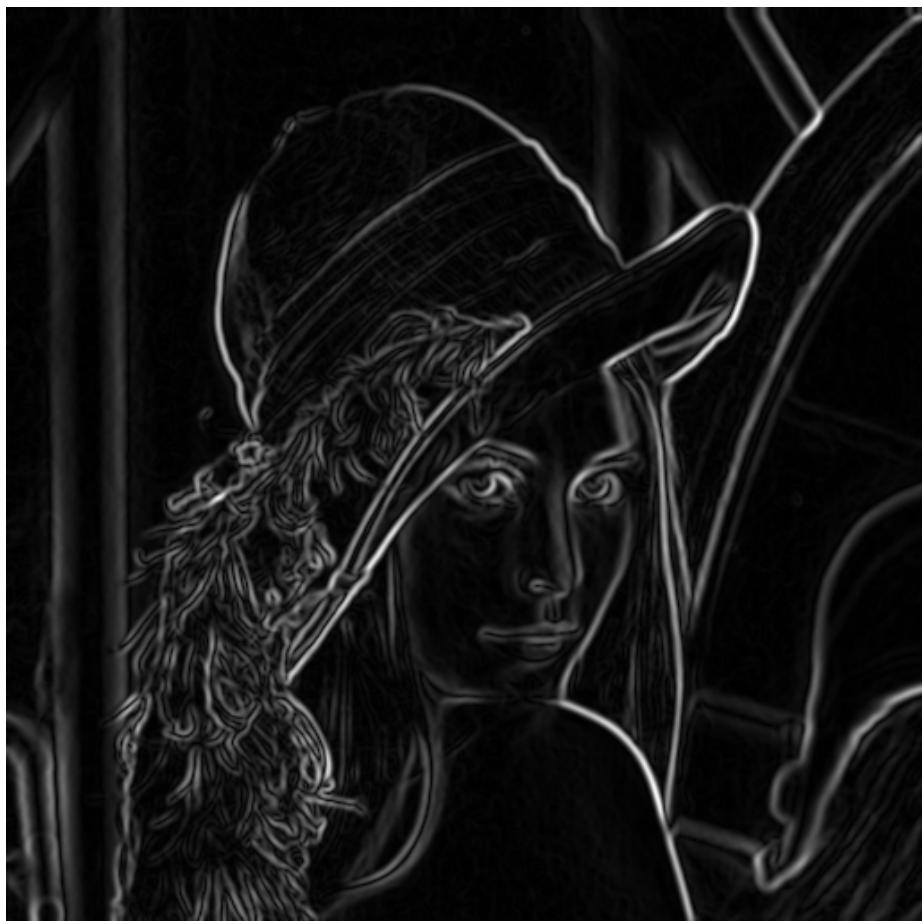
At q, we have a maximum if the value is larger than those at both p and at r.

Interpolate to get these values.

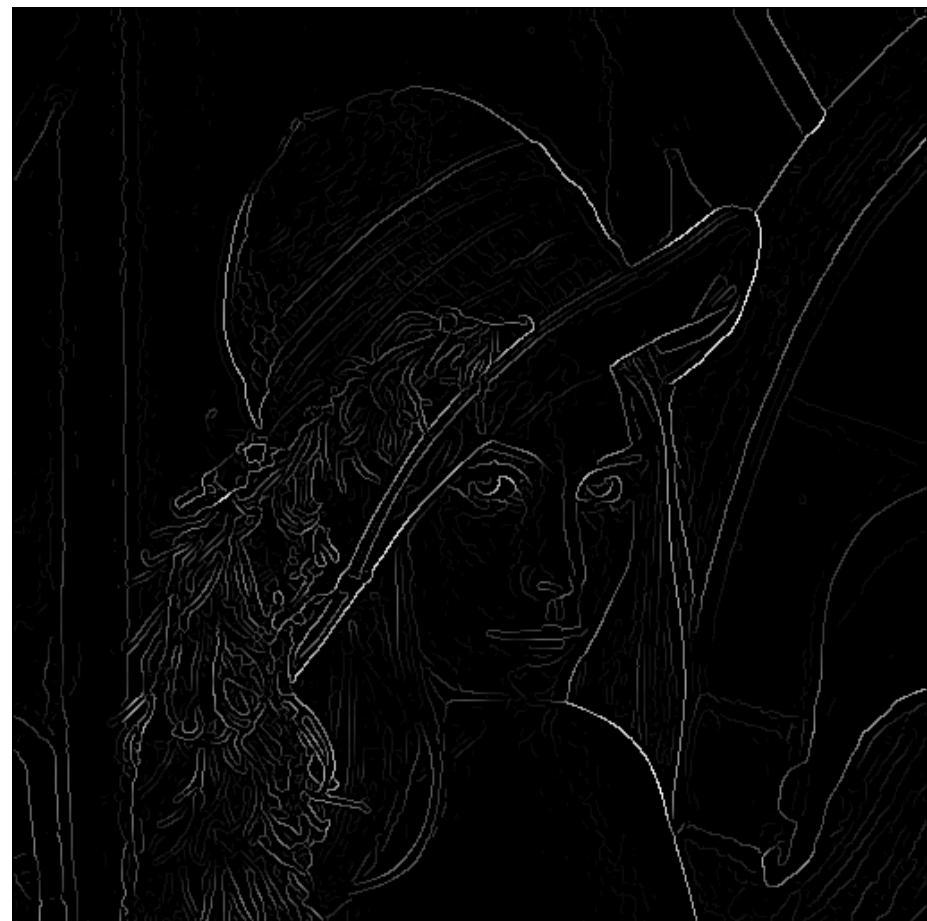
Source: S. Forsyth

# Example

## Non-max Suppression



Before



After

# Example

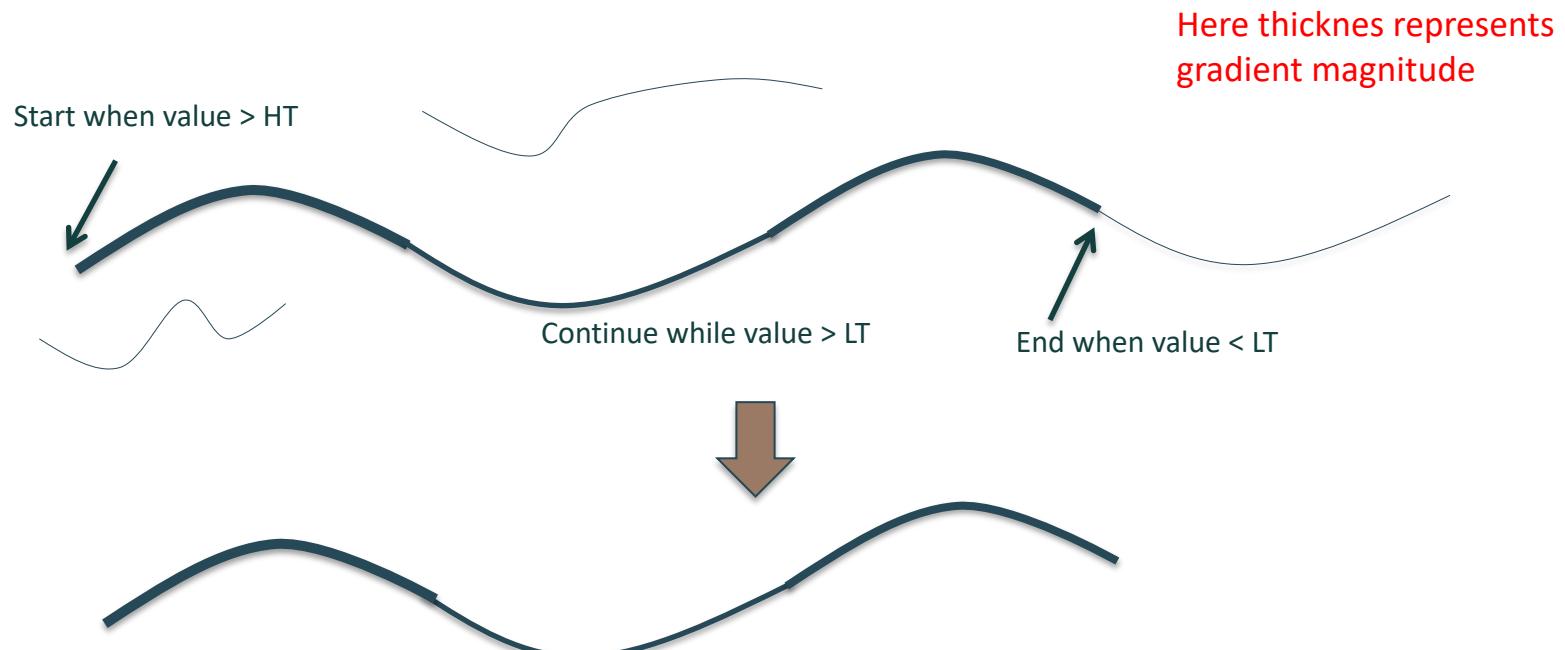
## Hysteresis thresholding

- Threshold at low/high levels to get weak/strong edge pixels
- Connect components, starting from strong edge pixels



# Hysteresis thresholding

- Use two thresholds: a high threshold (HT) to start edge curves and a low threshold (LT) to continue them
  - Iteratively label edges
    - Edges grow out from ‘strong edges’
    - Iterate until no change in image



# Example

## Final Canny Edges



### Matlab

```
contour = edge(ima, 'canny', threshold, sigma);
```

```
from skimage import feature
```

```
edges = feature.canny(im, sigma=3,low_threshold=0.1,  
high_threshold=0.2)
```

### Python

# Hysteresis thresholding example



Original image



High threshold  
(strong edges)



Low threshold  
(weak edges)



Hysteresis threshold

Source: F. F. Li

# Effect of scale $\sigma$

- Effect of the Gaussian kernel spread/size  $\sigma$



Original



Canny with  $\sigma = 1$



Canny with  $\sigma = 2$

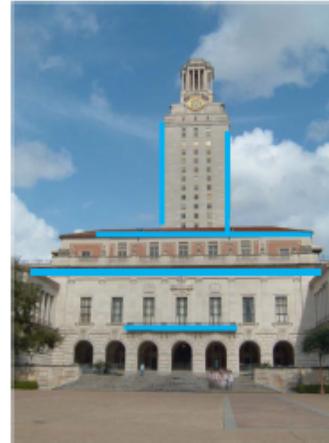
- The choice of  $\sigma$  depends on desired behavior
  - Large  $\sigma$  detects large scale edges
  - Small  $\sigma$  detects fine features

Source: S. Seitz

# ... Line fitting

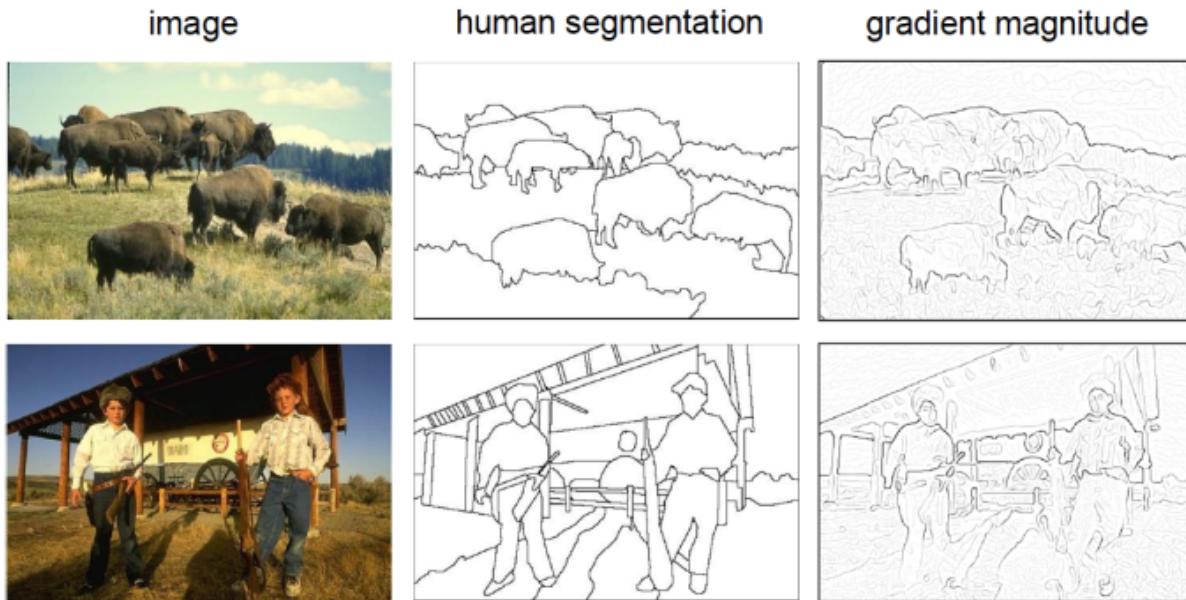
Many objects characterized by presence of straight lines

... are we done just by running edge detection?



- Problems
  - Extra edge points (clutter)
  - Only parts of each line detected, parts missing
  - Noise in measure edge points, orientations
- Solution: LS, Hough, Ransac ([Lecture 9](#))

# ... Image gradients vs. meaningful contours



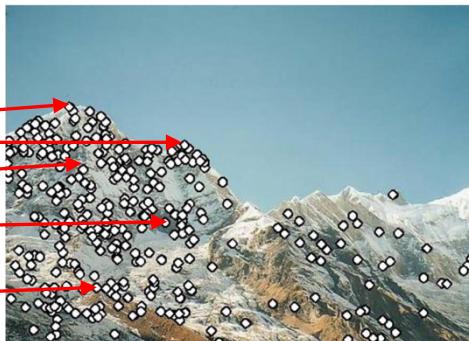
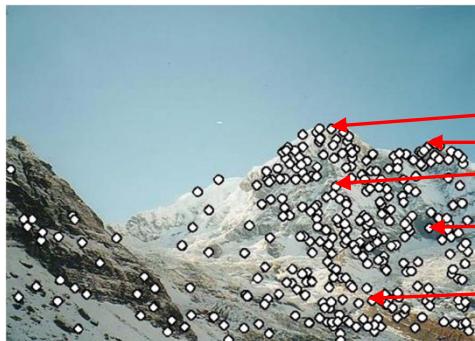
- Low level information is not enough: use higher level cues: texture, color,...
  - Contour Detection and Hierarchical Image Segmentation. P. Arbelaez, M. Maire, C. Fowlkes and J. Malik. IEEE TPAMI, Vol. 33, No. 5, pp. 898-916, May 2011.
  - **Sketch Tokens**: A Learned Mid-level Representation for Contour and Object Detection, CVPR13
  - **Structured Forests** for Fast Edge Detection, P. Dollár and C. Zitnick, ICCV 2013
  - **Deep Contour**: Wei Shen et al, DeepContour: A Deep Convolutional Feature Learned by Positive-sharing Loss for Contour Detection, CVPR 2015

# Local features

- Edges
  - Image gradient. Effect of noise
  - Derivative of Gaussian
  - Laplacian of Gaussian
  - Canny detector
- Corners
  - Local invariant features. Requirements
  - Harris corner detector
- Blobs
  - Scale invariant region selection
  - Automatic scale selection
  - Laplacian of Gaussian (LoG)
  - Difference of Gaussian (DoG)
  - Affine invariant region selection

# Application: panorama stitching

How do we combine the two images?



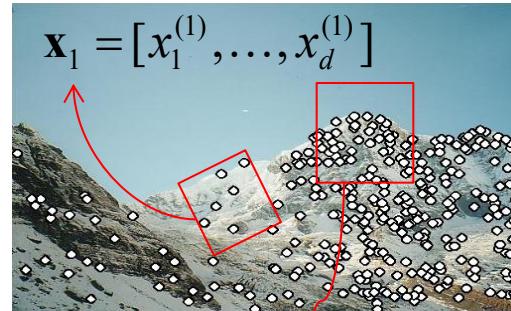
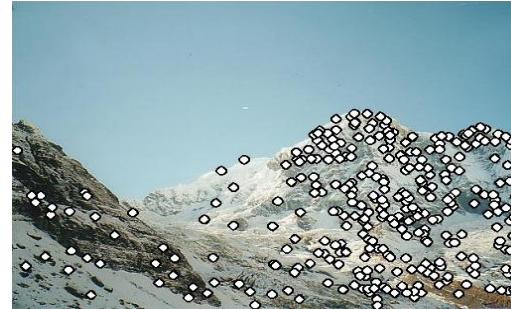
- Detect feature points in both images
- Find corresponding pairs
- Use these pairs to align the images

Source: Frolova, Simakov

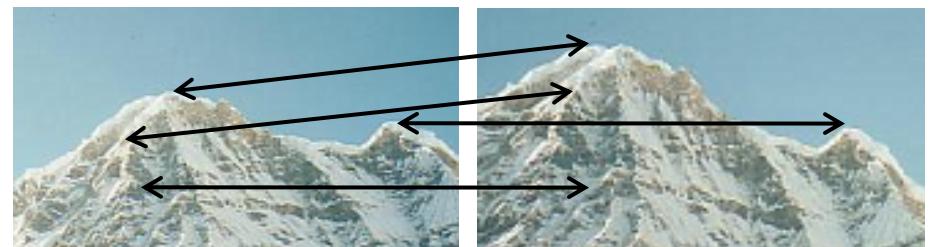
# Application: panorama stitching

Approach:

- **Detection** : Identify a set of distinctive *keypoints*
- **Description** : Define a region around each point, extract and normalize content, compute a *local descriptor*
- **Matching** : Determine the *correspondence* between descriptors in two views



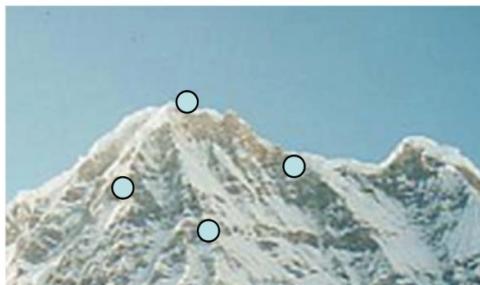
$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$



Source: K. Grauman

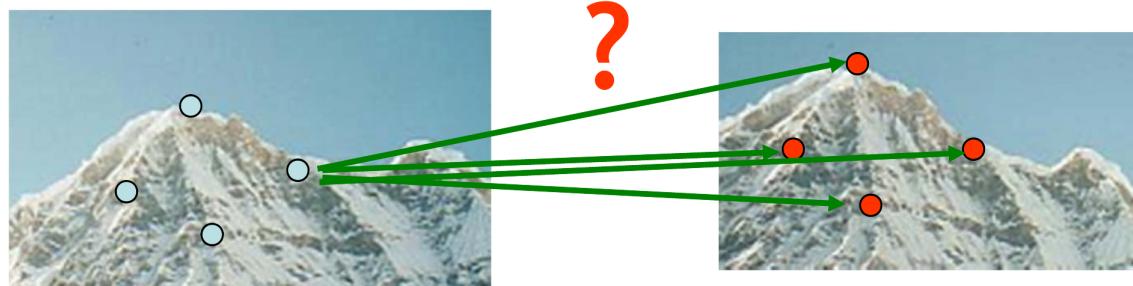
# Matching with features: requirements

- Detect (at least some of) the same points in both images: Run the detection independently per image
  - Repeatable detector



No chance to find true matches!

- For each point correctly recognize the corresponding one
  - Reliable and distinctive descriptor



- Provide some **invariance (covariance)** to photometric and geometric differences between the views

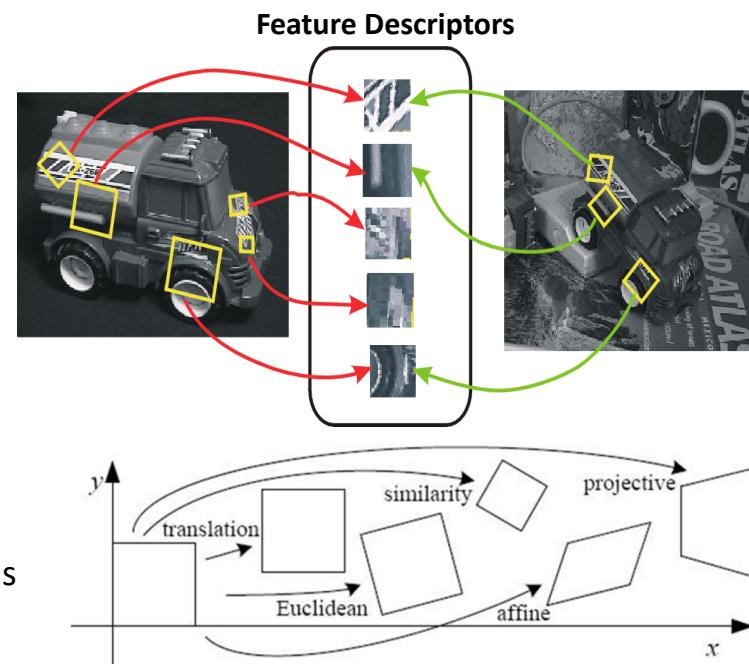
Source: K. Grauman

# Invariance of local features

- We want feature locations to be **invariant** to photometric transformations (brightness, exposure,...) and **covariant** to geometric transformations (translation, rotation, scale,...)
  - **Invariance**: image is transformed and feature locations do not change
  - **Covariance**: if we have two transformed versions of the same image, features should be detected in corresponding locations



photometric transformations



Geometric  
transformations

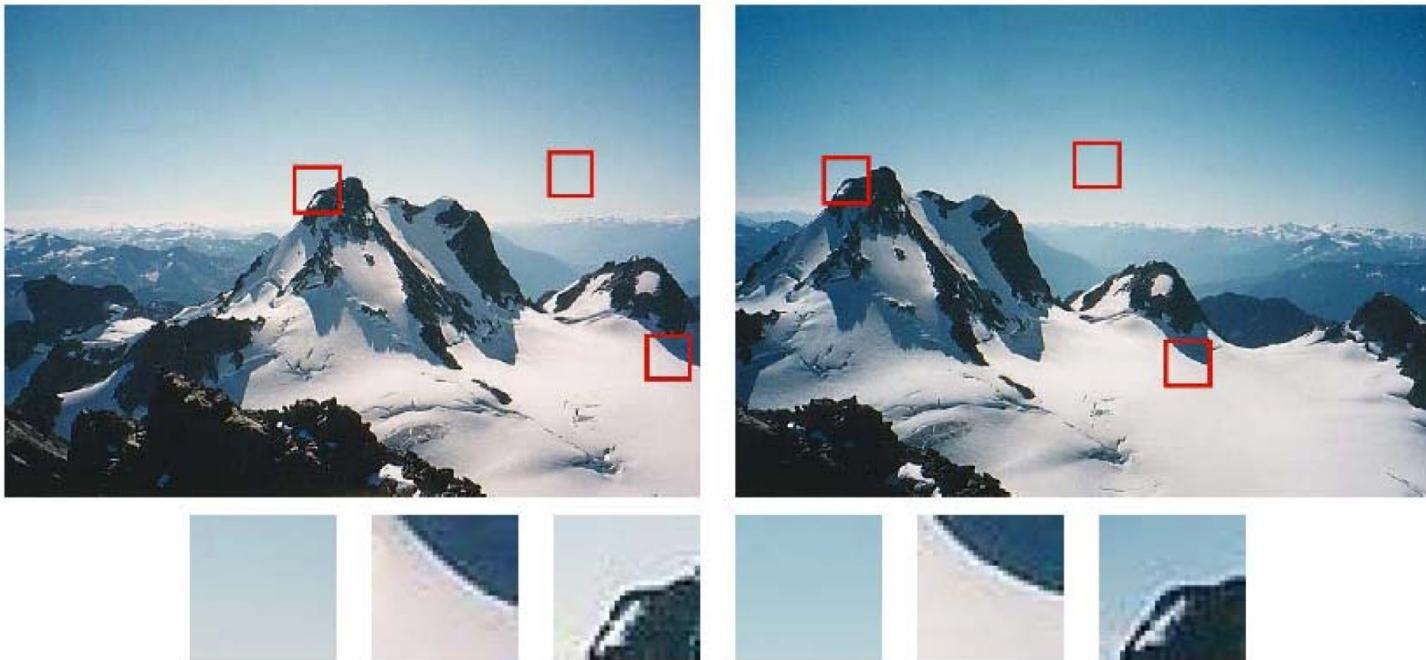
# Many existing detectors available

- Hessian & Harris [Beaudet '78], [Harris '88]
- Laplacian, DoG [Lindeberg '98], [Lowe 1999]
- Harris-/Hessian-Laplace [Mikolajczyk & Schmid '01]
- Harris-/Hessian-Affine [Mikolajczyk & Schmid '04]
- EBR and IBR [Tuytelaars & Van Gool '04]
- MSER [Matas '02]
- Salient Regions [Kadir & Brady '01]
- Others...

Those detectors have become a basic building block in many applications in Computer Vision

# Keypoints

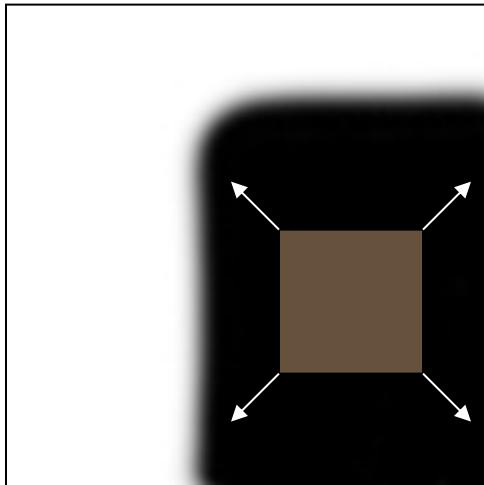
- Which points would you chose?
- Some patches can be localized or matched with higher accuracy than others



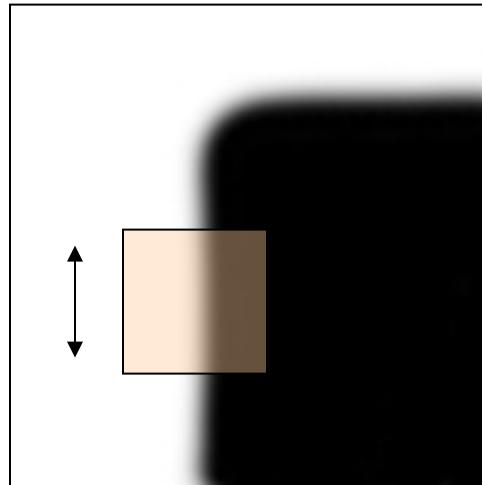
- Points worth matching are **corners**, because a corner can be localized:
  - intersection of two edges
  - point with two dominant (different) edge directions in a neighborhood

# Corner detection

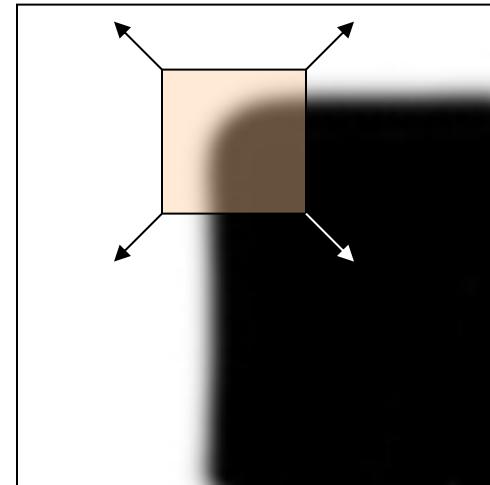
- We should easily recognize the point by looking through a small window
- Shifting a window  $w$  in *any direction* should give a *large change* in intensity



“flat” region:  
no change in all  
directions



“edge”:  
no change along the  
edge direction



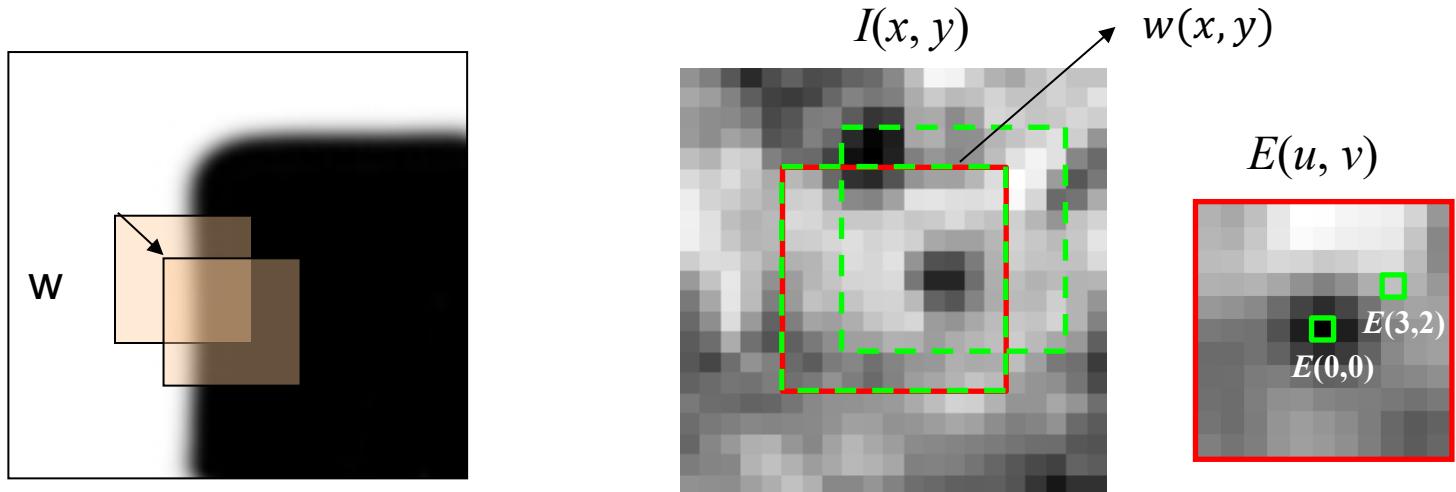
“corner”:  
significant change in all  
directions

- **Key property:** in the region around a corner, image gradient has two or more dominant directions
- Corners are repeatable and distinctive

Source: A. Efros

# A simple matching criterion

- Consider shifting the window  $w$  by  $(u, v)$ 
  - How do the pixels in  $w$  change?
  - Compare each pixel before and after by summing up the squared differences (SSD)



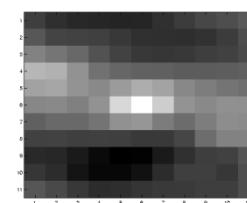
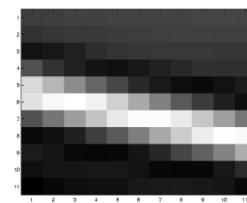
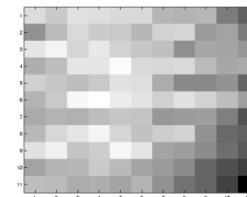
- This defines an SSD error (change in window for shift  $(u, v)$ )

$$E(u, v) = \sum_{(x,y) \in w} [I(x + u, y + v) - I(x, y)]^2 = \sum_{(x,y)} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

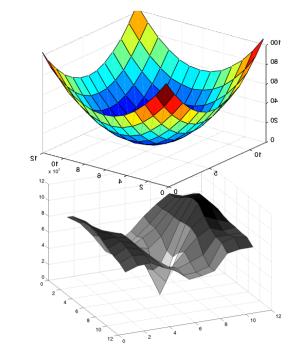
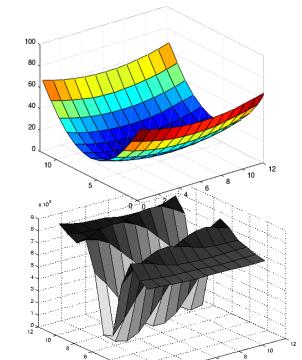
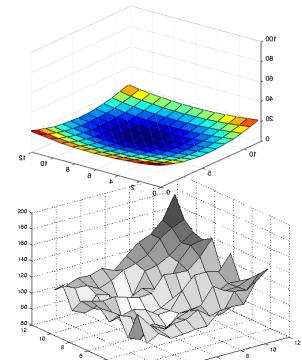
# Corner Detection

Note: inverted gray-scale!!



At a corner we expect two effects:

- large gradients
- In a small neighborhood the gradient orientation should swing sharply



Source: R. Szeliski

# Corner Detection

- We assume small motion
- $E(u,v)$  can be approximated in a neighborhood of a point by a *second-order Taylor expansion*:

$$\begin{aligned} E(u,v) &\approx \sum_{x,y} w(x,y) \left[ I(x,y) + uI_x + vI_y - I(x,y) \right]^2 \\ &= \sum_{x,y} w(x,y) [uI_x + vI_y]^2 = \\ &= \sum_{x,y} w(x,y) (u, v) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix} \begin{pmatrix} u \\ v \end{pmatrix} \end{aligned}$$

Gradient with respect to x,  
times gradient with respect to y  
*'Second moment matrix' M*

The quadratic approximation simplifies to

$$E(u,v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Sum over image region –the area we are checking for corner

Source: R. Szeliski

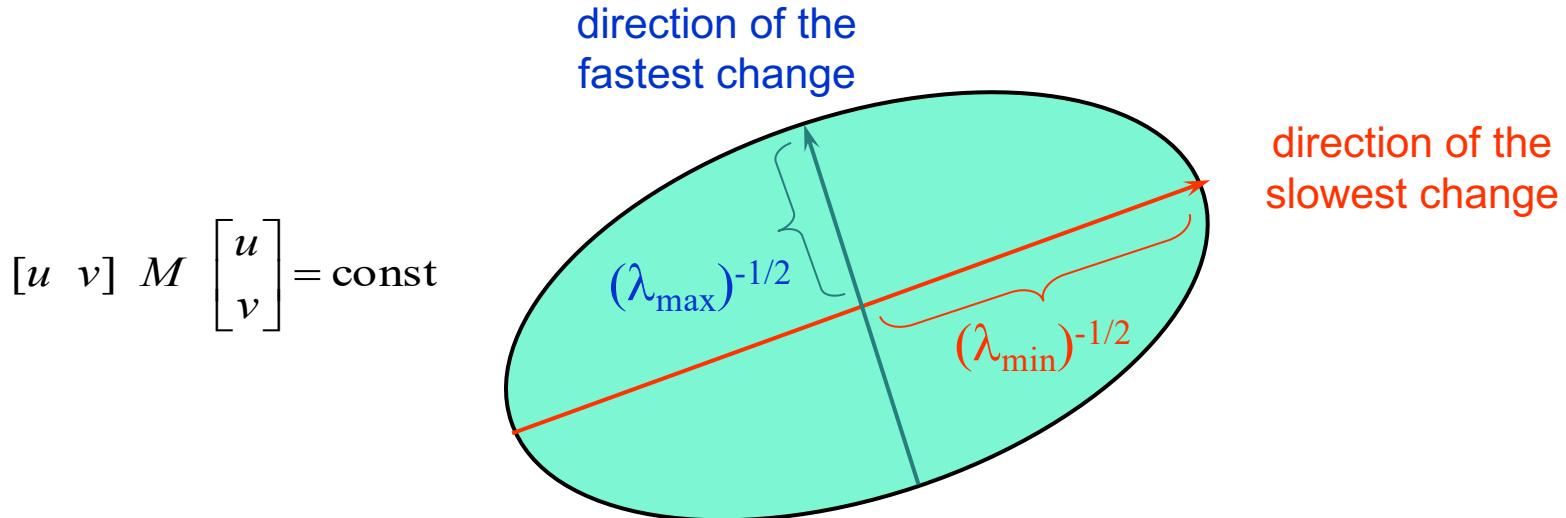


# Interpreting of the second moment matrix

- $M$  can be diagonalized

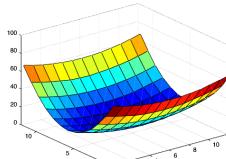
$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

- Points are characterized in terms of the eigenvalues of  $M$ : the axis lengths of the ellipse are determined by the eigenvalues and the orientation is determined by  $R$



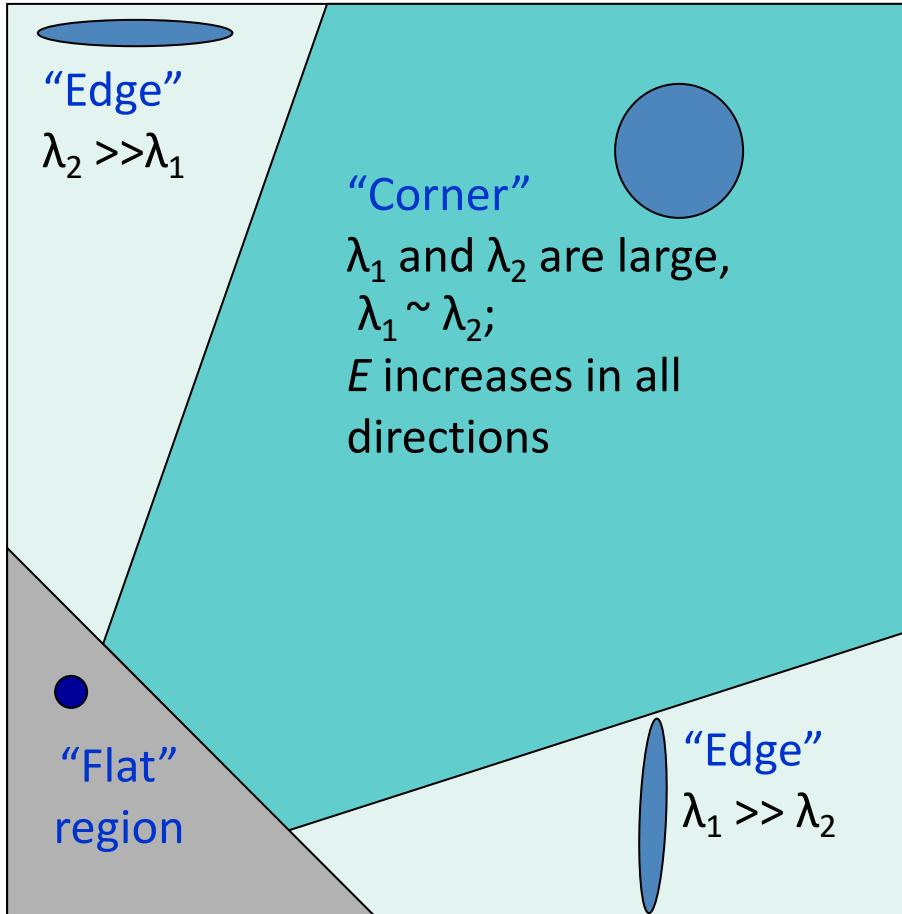
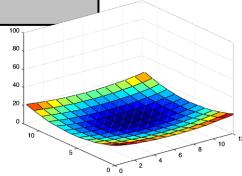
# Interpreting the eigenvalues

Classification of image points using eigenvalues of  $M$ :



$\lambda_2$

$\lambda_1$  and  $\lambda_2$  are small;  
 $E$  is almost constant in  
all directions



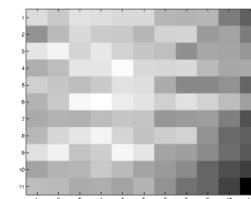
$\lambda_1$

Source: K. Grauman

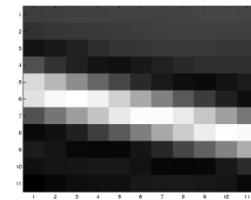
# Corner Detection



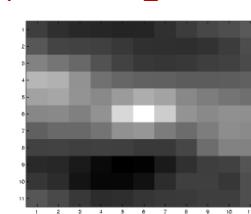
Small  $\lambda_1$ , small  $\lambda_2$



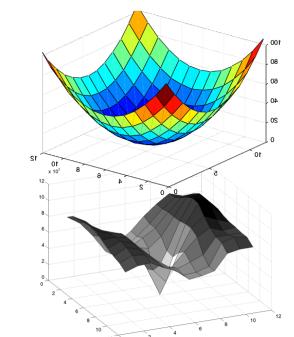
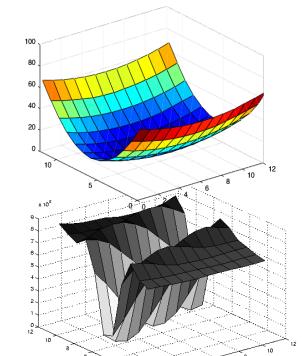
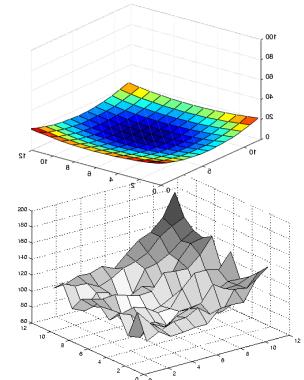
Large  $\lambda_1$ , small  $\lambda_2$



$\lambda_1$  and  $\lambda_2$  are large



Note: inverted gray-scale!!



Source: R. Szeliski

# Corner response function R

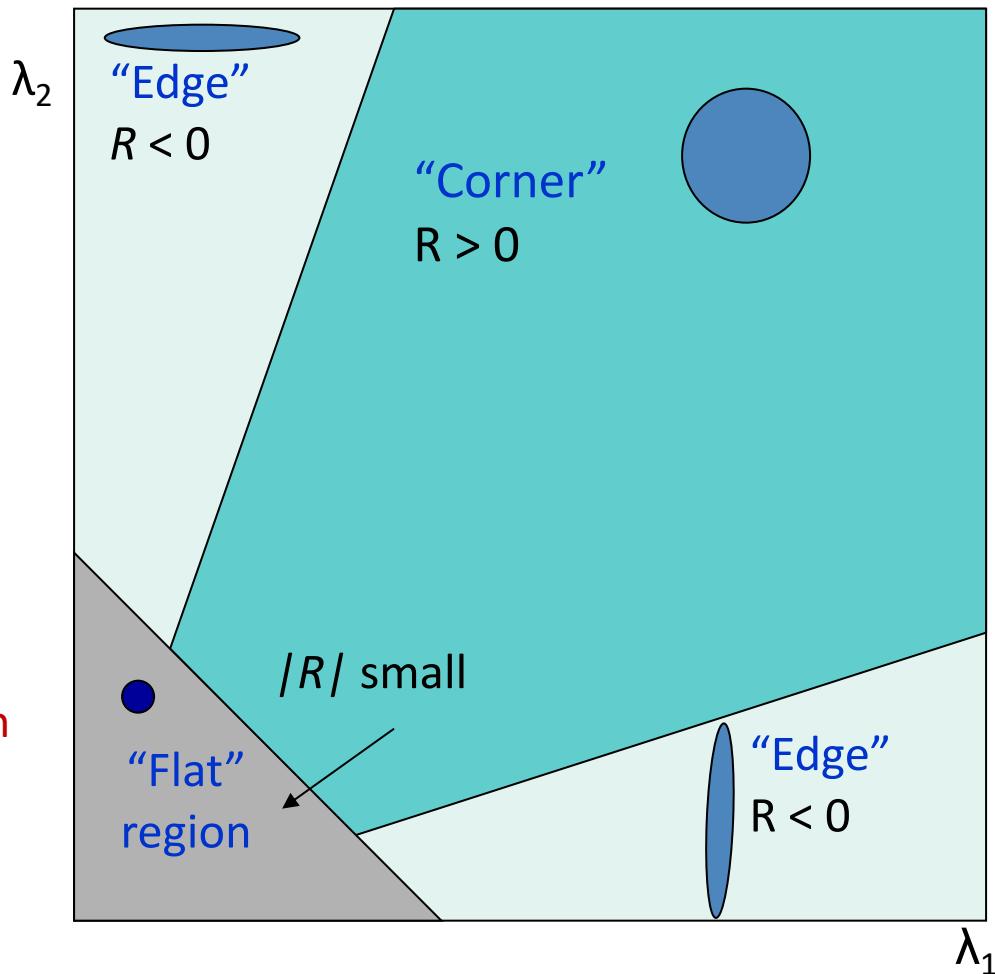
A fast approximation to avoid computing eigenvalues: a function of the eigenvalues

$$R = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2 = \\ = \det(M) - \alpha \operatorname{trace}(M)^2$$

$\alpha$ : constant (0.04 to 0.06)

$R$  depends only on eigenvalues of  $M$

- $R$  is large for a **corner**
- $R$  is negative with large magnitude for an **edge**
- $|R|$  is small for a **flat** region



Source: K. Grauman

# Window function $w(x,y)$

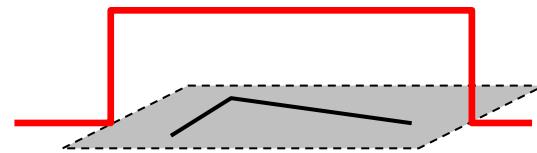
- Option 1: uniform window

- Sum over square window

$$M = \sum_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Problem: not rotation invariant

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

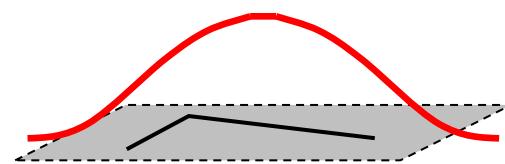


1 in window, 0 outside

- Option 2: smooth with Gaussian
- Gaussian already performs weighted sum

$$M = g(\sigma) * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Result is rotation invariant



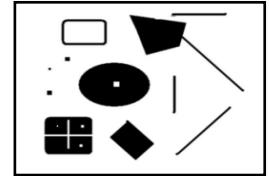
Gaussian

# Harris corner detector [Harris 88]

## Algorithm:

1. Compute M, the second moment matrix (**optionally blur first**)
2. Compute cornerness function – both eigenvalues are strong
3. Threshold result: remove points with low cornerness value
4. Suppress non-maxima: remove points with a neighbor with a higher cornerness value

# Harris corner detector [Harris 88]



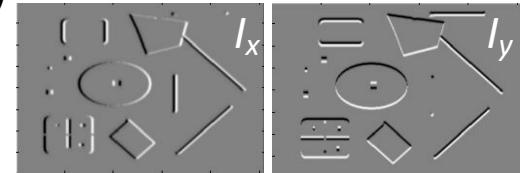
1. Compute M, the second moment matrix (optionally blur first)

$$M(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

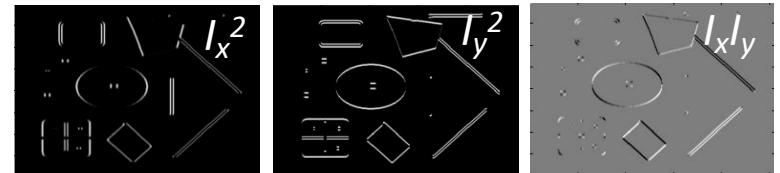
$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

- 1.a. Image derivatives (optionally, blur first)



- 1.b. Products of derivatives



- 1.c. Gaussian filter  $g(\sigma_I)$



2. Cornerness function – both eigenvalues are strong

$$R = \det[M(\sigma_I, \sigma_D)] - \alpha[\text{trace}(M(\sigma_I, \sigma_D))^2] = \\ g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha[g(I_x^2) + g(I_y^2)]^2$$

3. Apply threshold

4. Apply non-maxima suppression



# Harris detector: example

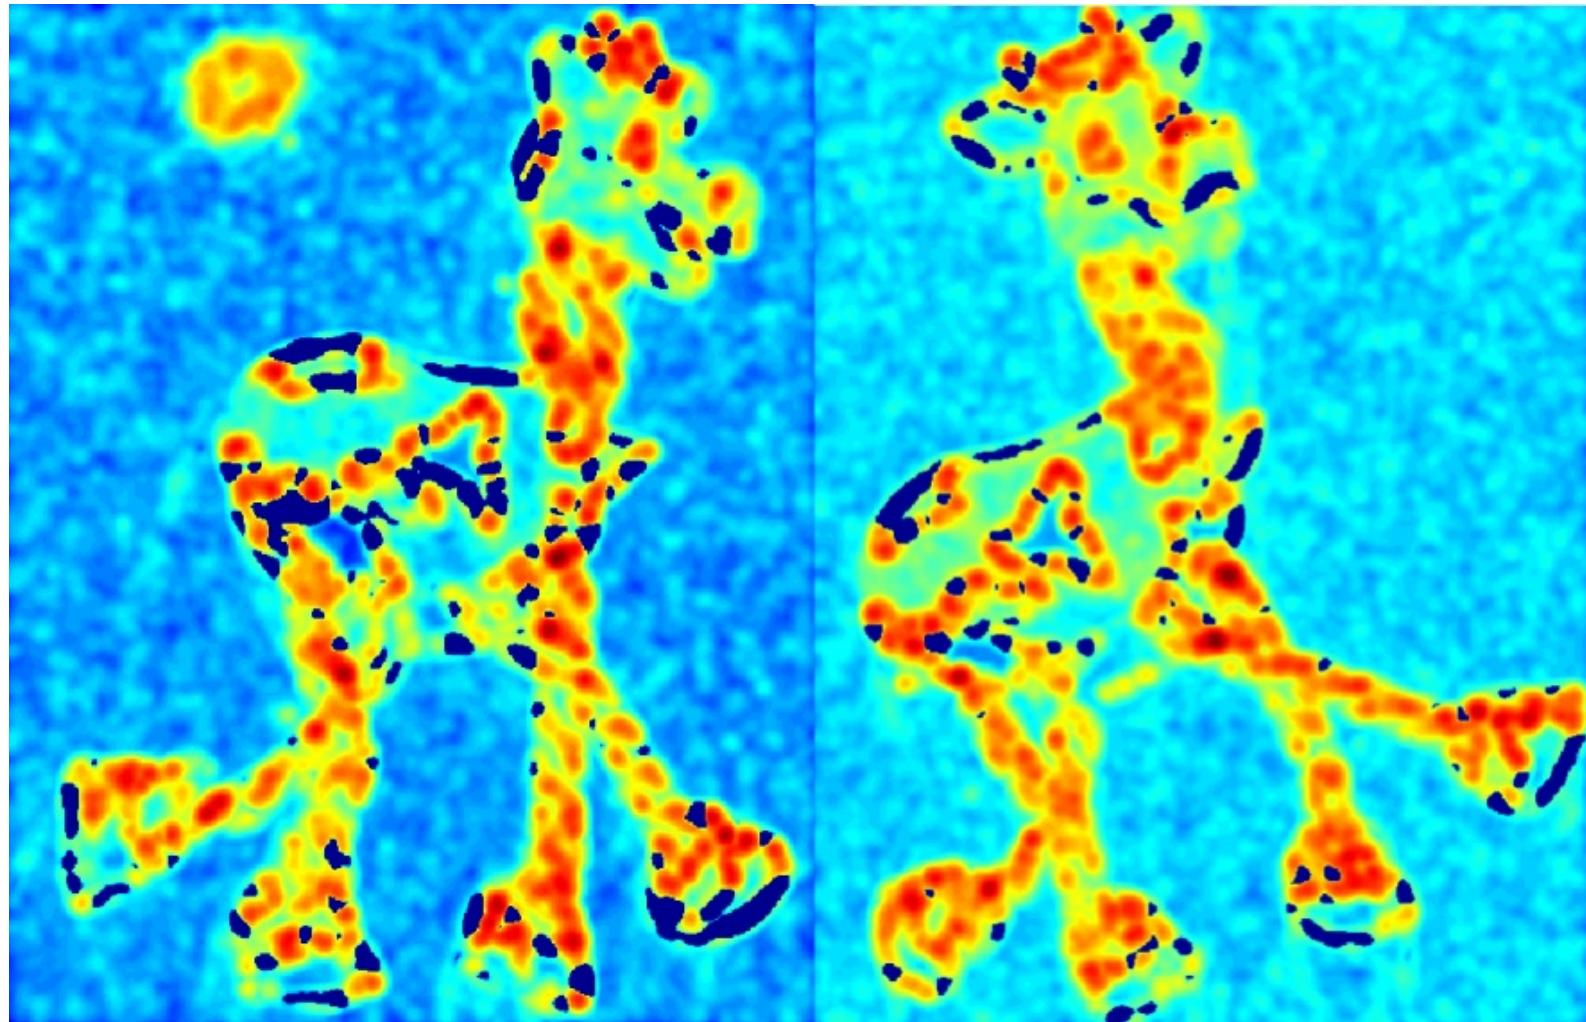
- Original images



Source: Frolova, Simakov

# Harris detector: example

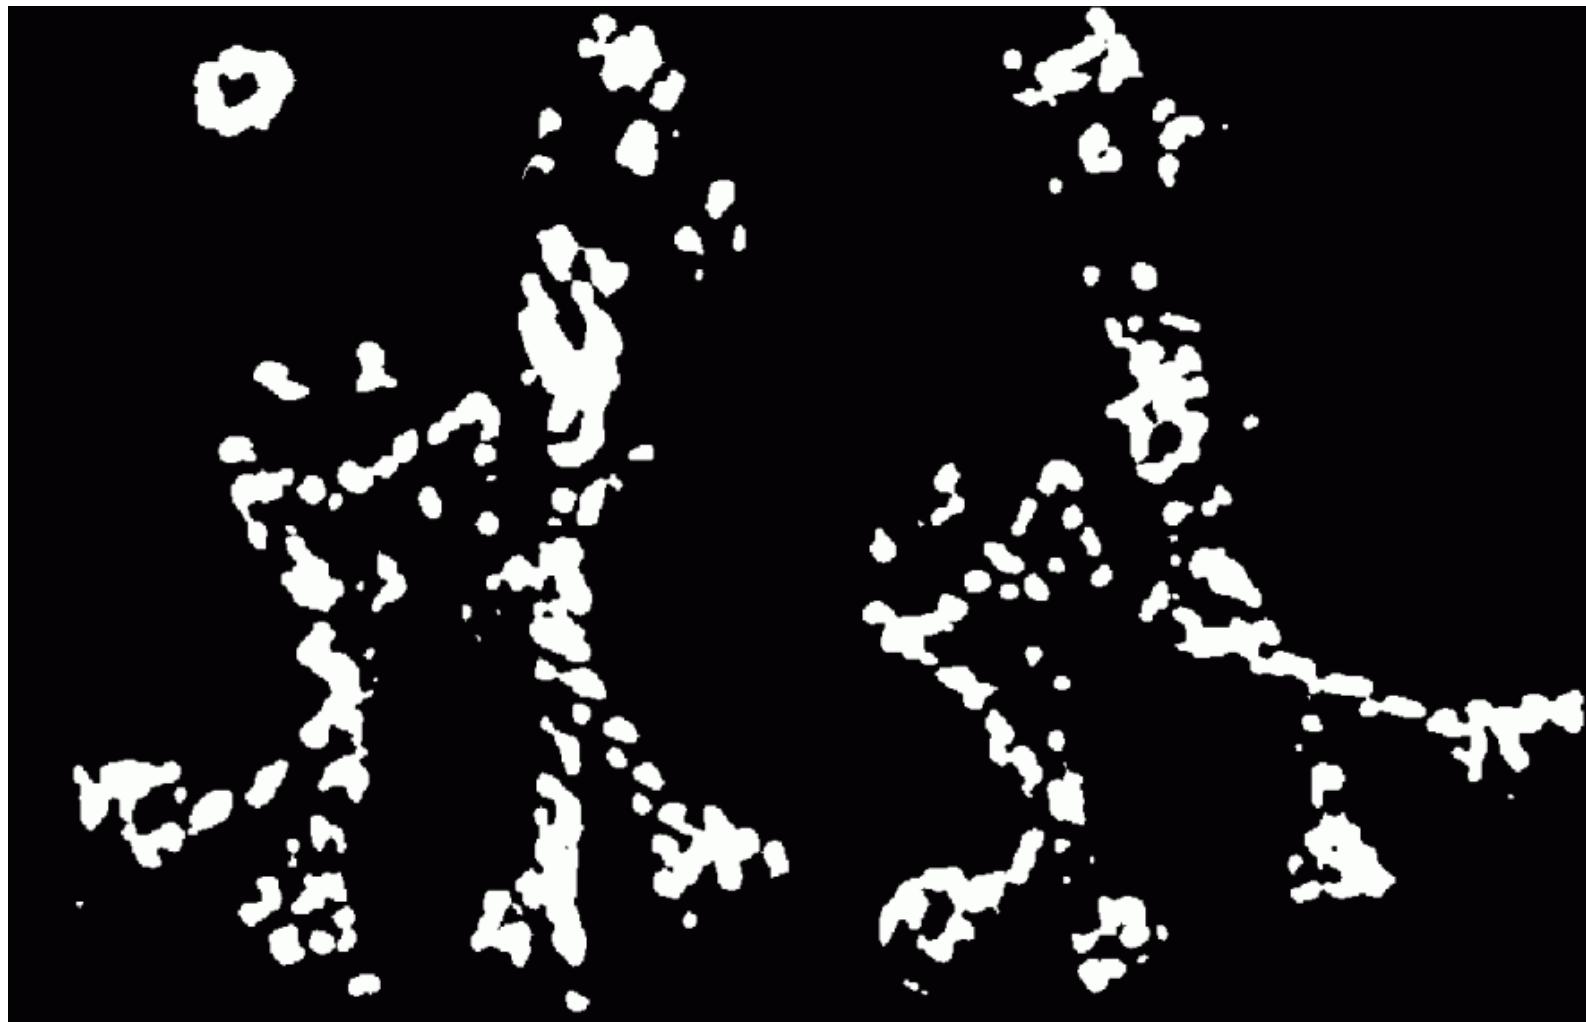
- Compute corner response  $R$



Source: Frolova, Simakov

# Harris detector: example

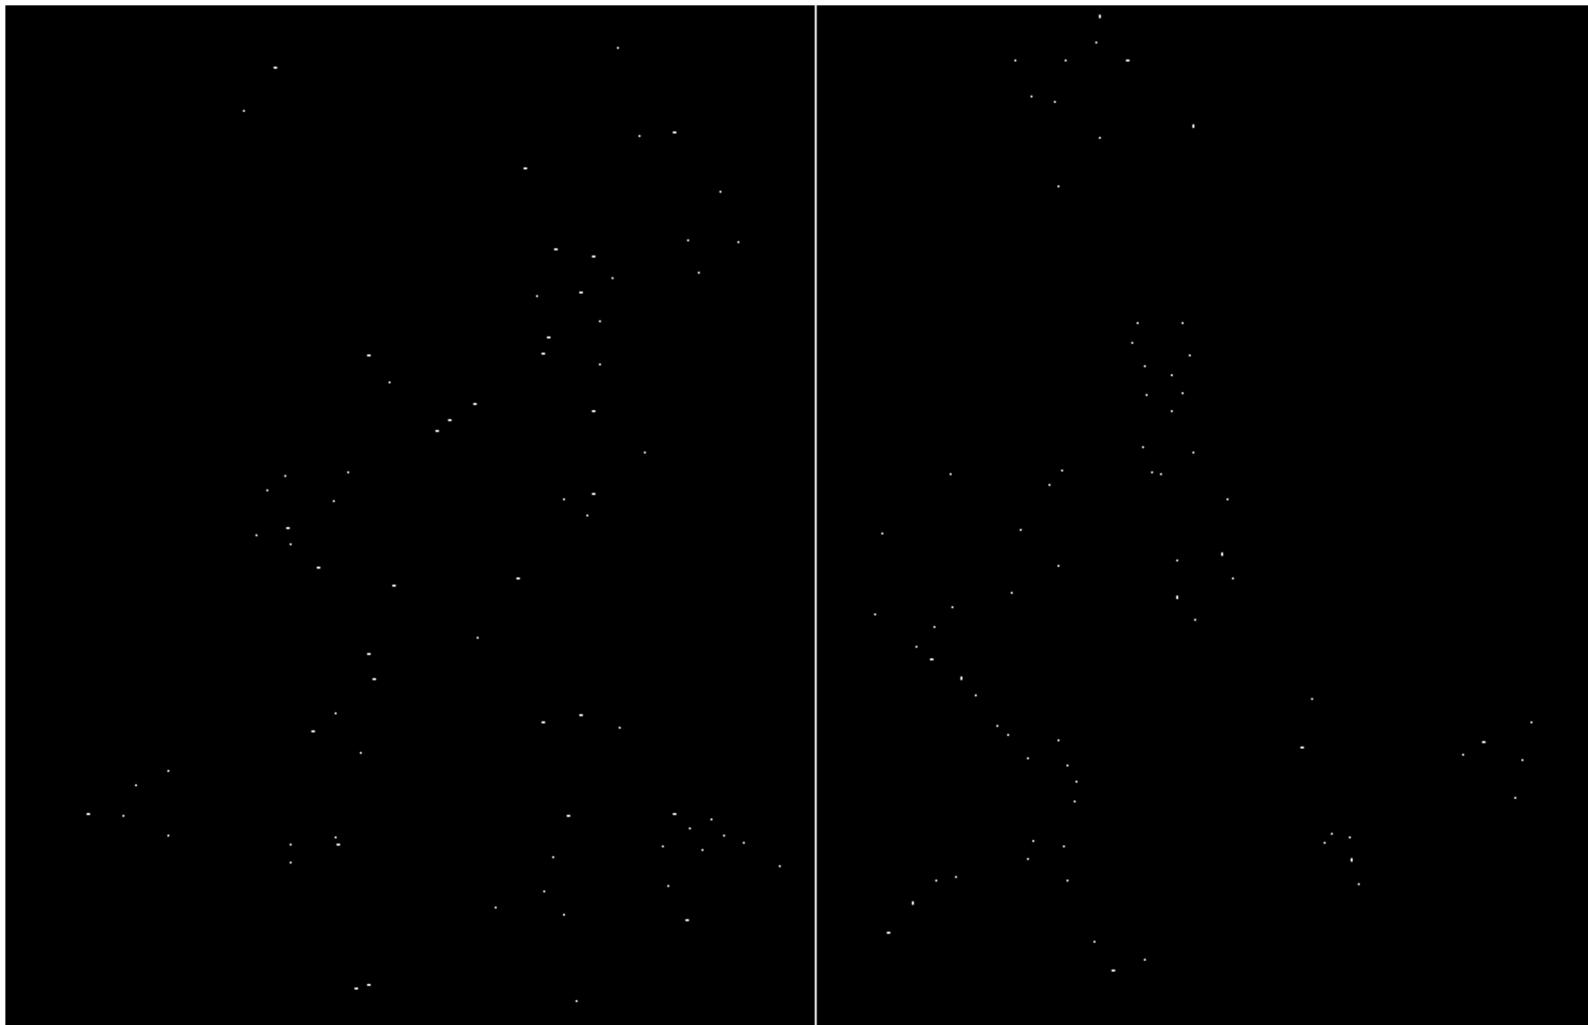
- Find points with large corner response:  $R > \text{threshold}$



Source: Frolova, Simakov

# Harris detector: example

- Take only the points of local maxima of  $R$



Source: Frolova, Simakov

# Harris detector: example

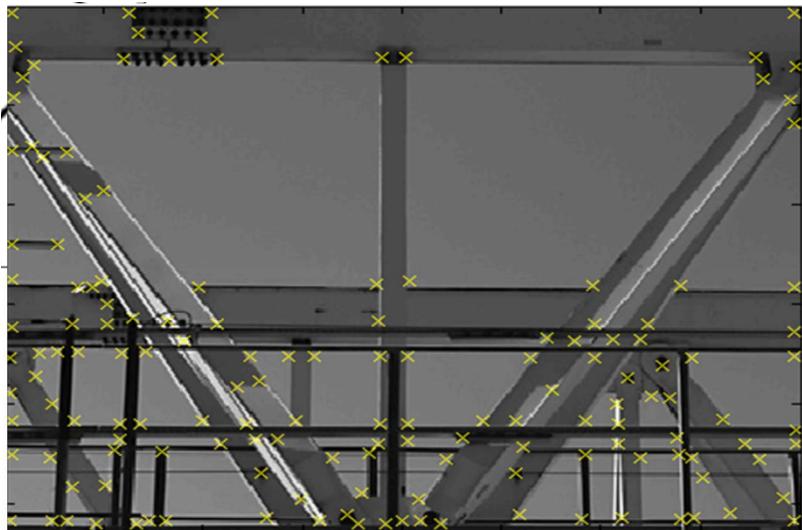
- Result



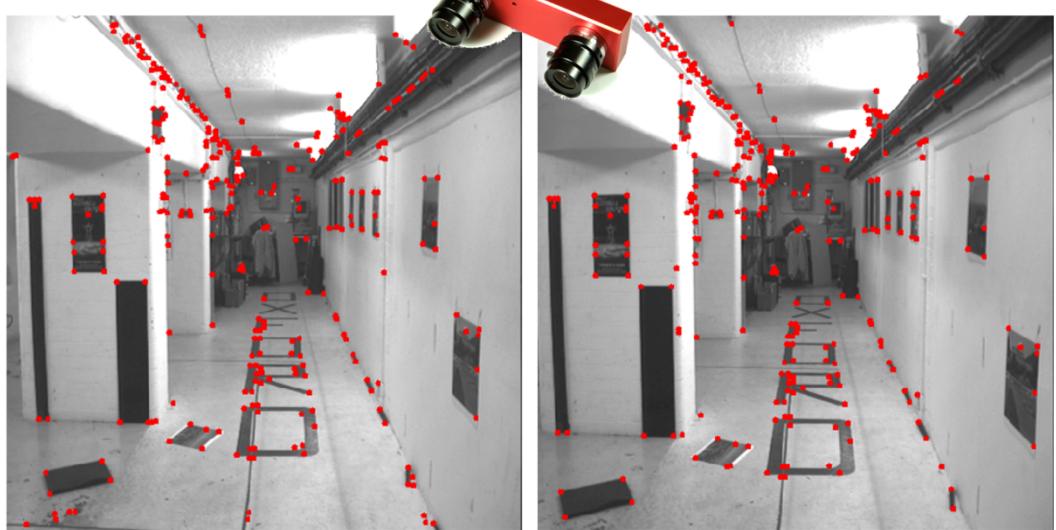
Source: Frolova, Simakov

# Harris detector: more examples

- Effect: a very precise corner detection



- Results are well suited for finding stereo correspondences

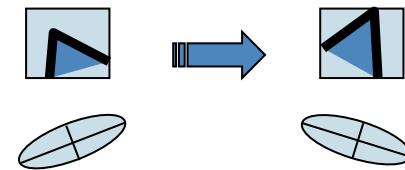


Source: K. Grauman

# Harris detector: Properties

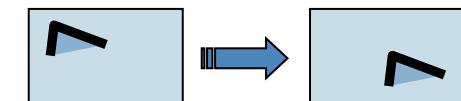
- Covariant wrt rotation:

Second moment ellipse rotates but its shape (i.e. eigenvalues) remains the same



- Covariant wrt translation

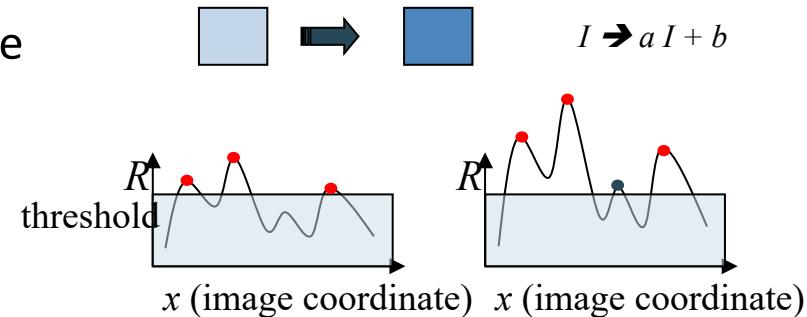
Derivatives and window function are shift-invariant



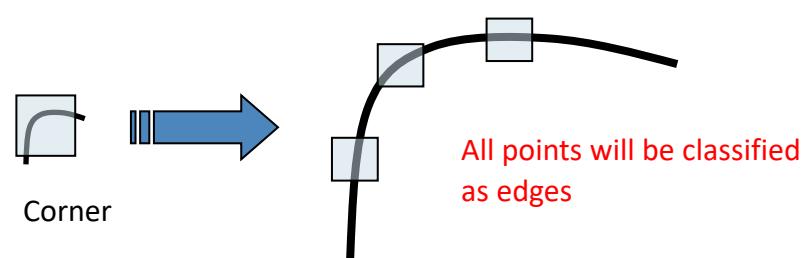
- Partially invariant wrt affine intensity change

Only derivatives are used: invariance to intensity shift  $I \rightarrow I + b$

Intensity scaling:  $I \rightarrow a I$



- Not covariant to scaling



# From points to regions

- How can we independently select interest points in each image, such that the detections are repeatable across different scales?



- How to find corresponding patch sizes, with only one image in hand?

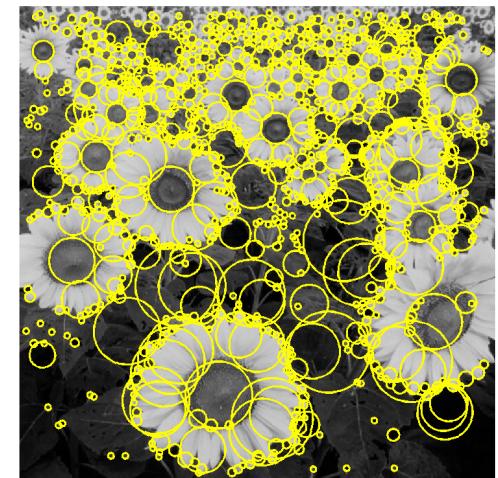
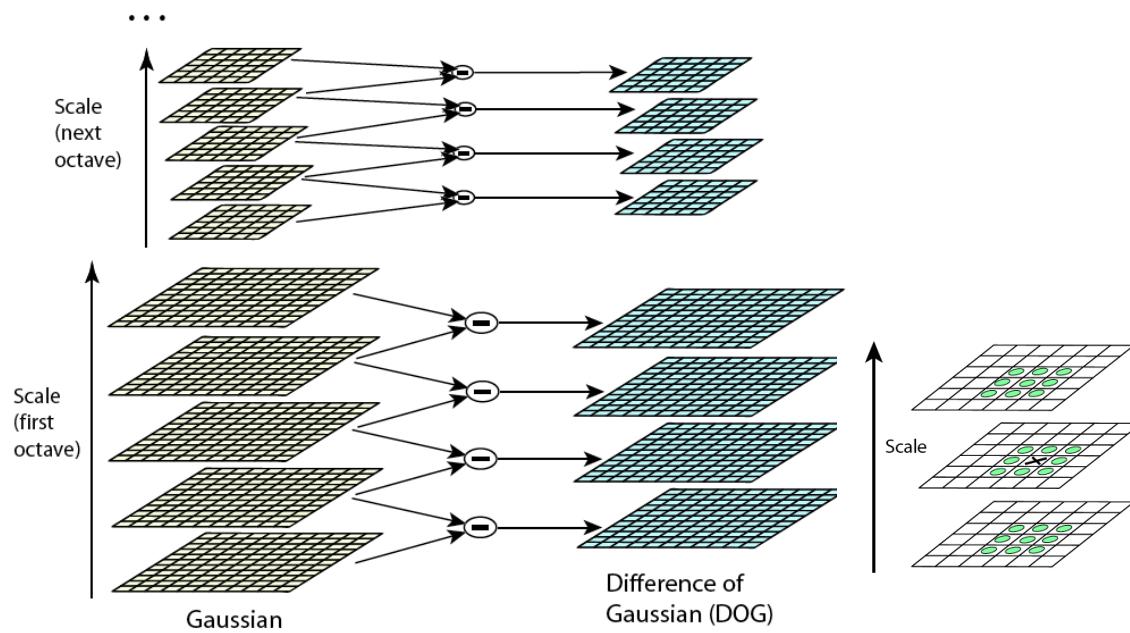


# Local features

- Edges
  - Image gradient. Effect of noise
  - Derivative of Gaussian
  - Laplacian of Gaussian
  - Canny detector
- Corners
  - Local invariant features. Requirements
  - Harris corner detector
- Blobs
  - Scale invariant region selection
  - Automatic scale selection
  - Laplacian of Gaussian (LoG)
  - Difference of Gaussian (DoG)

# Scale invariant features: blobs

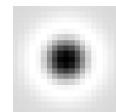
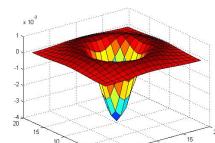
- The solution will be: to detect blobs, convolve the image with a ‘blob filter’ at multiple scales and look for extrema of filter response in the resulting *scale space*
  - Lindeberg (1998): extrema in the Laplacian of Gaussian (LoG).
  - Lowe (2004) proposed computing a set of sub-octave Difference of Gaussian filters looking for 3D (space+scale) maxima in the resulting structure (DoG: SIFT)



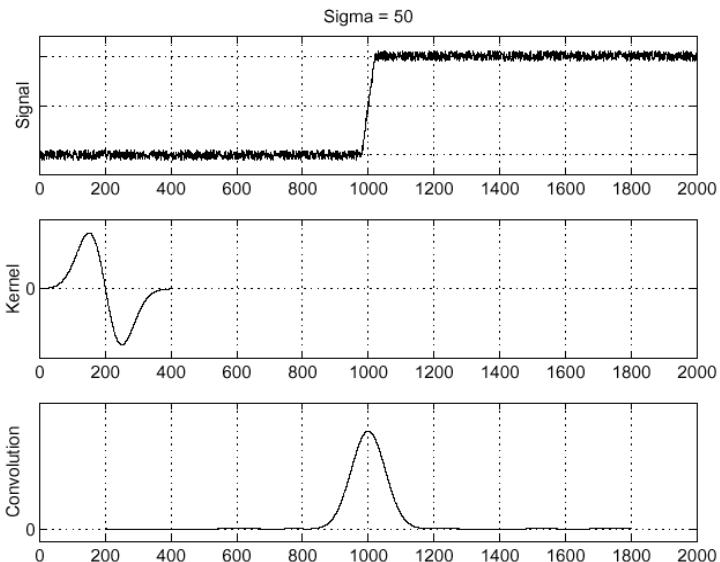
# Automatic scale selection

- Laplacian-of-Gaussian = “blob” detector

$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \quad g = \text{gaussian}$$

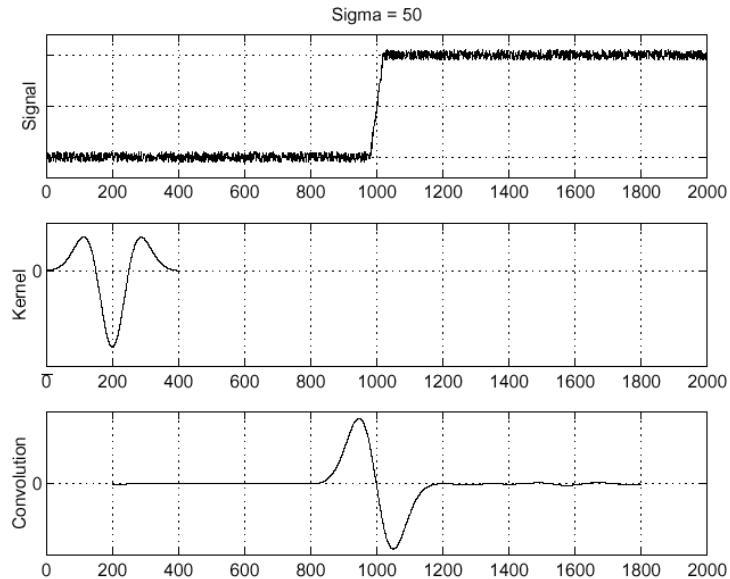


- Recall: edge detection



Derivative of Gaussian

Edge = maximum of derivative



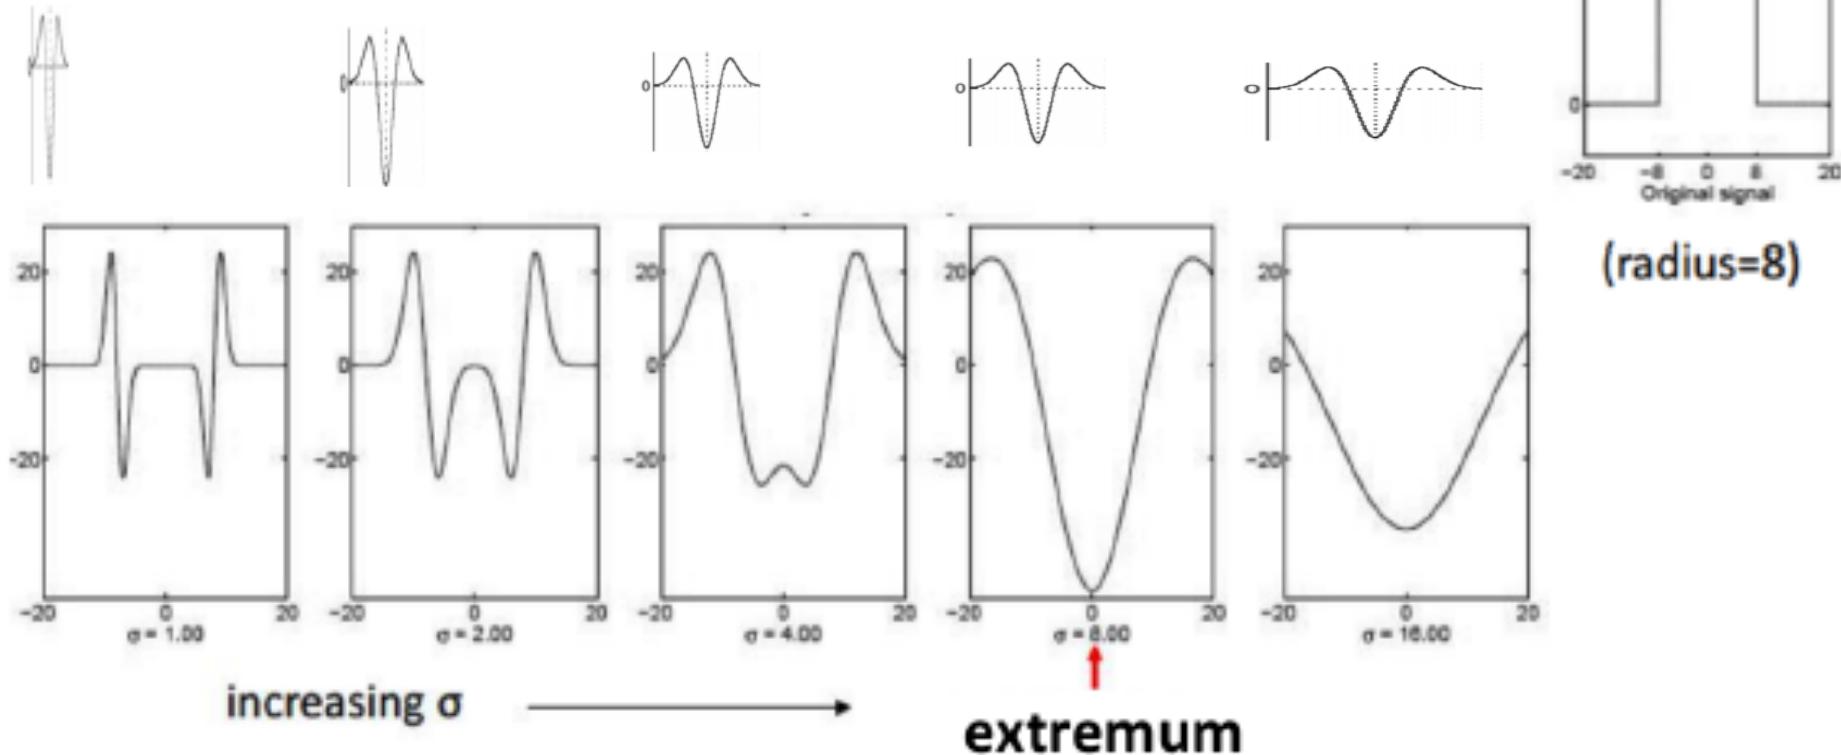
Second derivative of Gaussian (Laplacian)

Edge = zero crossing of second derivative

# From edges to blobs

Edge = ripple

Blob = superposition of two ripples

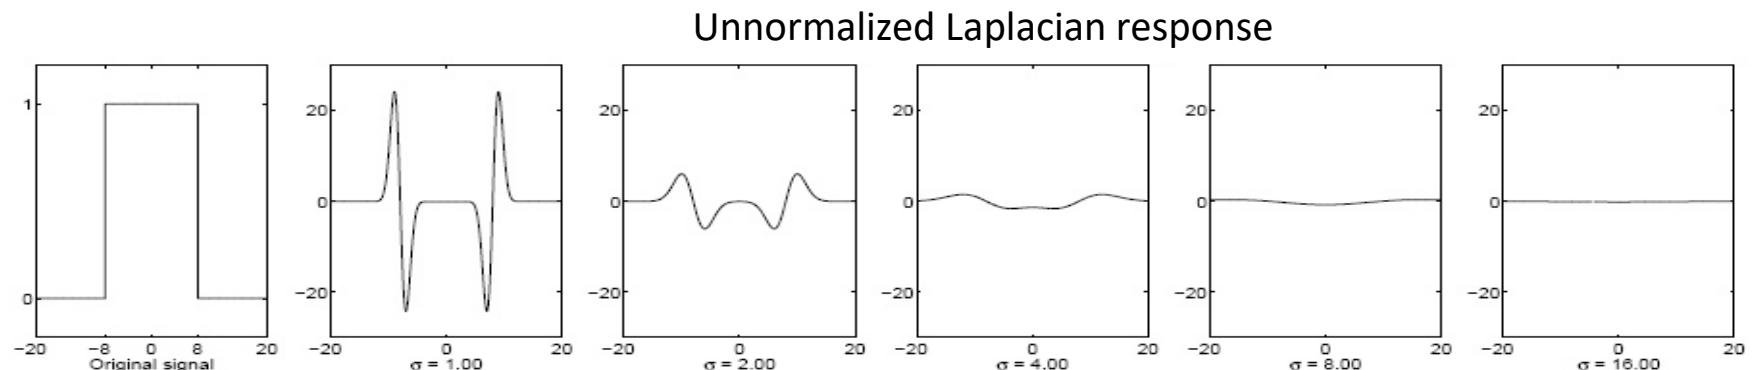


**Spatial selection:** magnitude of the Laplacian response will achieve a **maximum (or min)** at the center of the blob when the scale of the Laplacian is “matched” to the scale of the blob

# Scale selection

- We find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response
- However, Laplacian response decays as scale increases

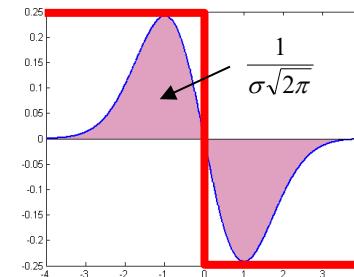
$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$



# Scale normalization

- Laplacian response decays as scale increases

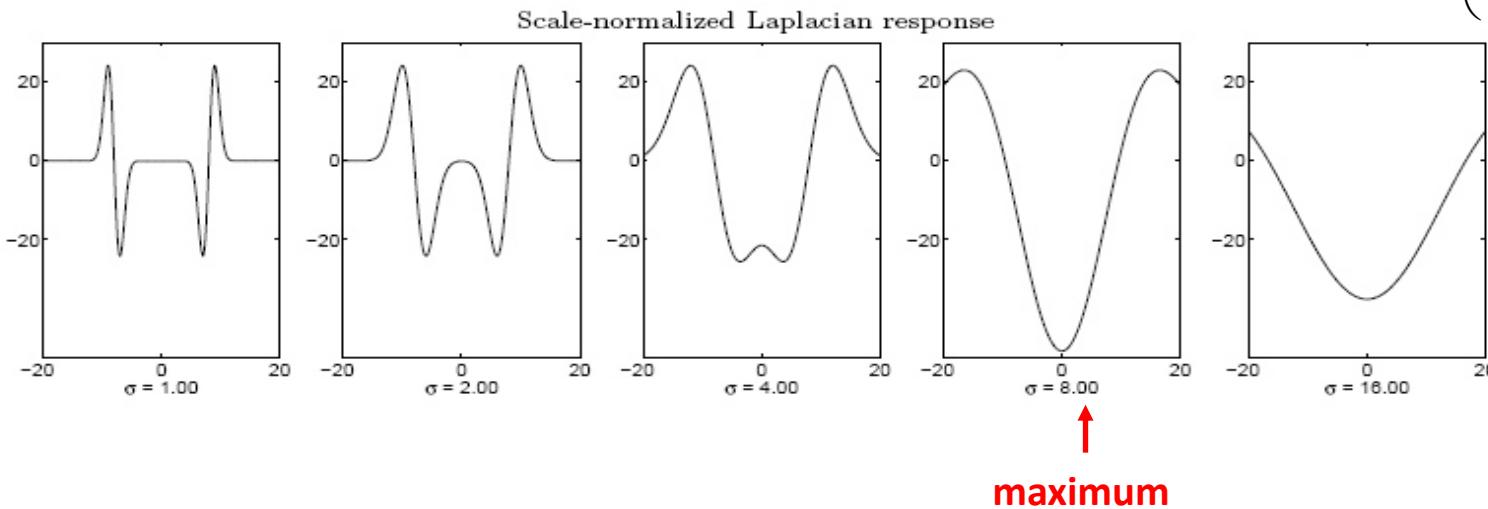
The response of the derivative of Gaussian filter to a perfect step edge decreases as  $\sigma$  increases



- To keep response the same (scale-invariant), we must multiply Gaussian derivative by  $\sigma$
- Laplacian is the second Gaussian derivative, so it must be multiplied by  $\sigma^2$

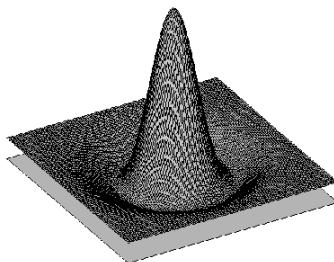
Scale-normalized Laplacian response

$$\nabla_{\text{norm}}^2 g = \sigma^2 \left( \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

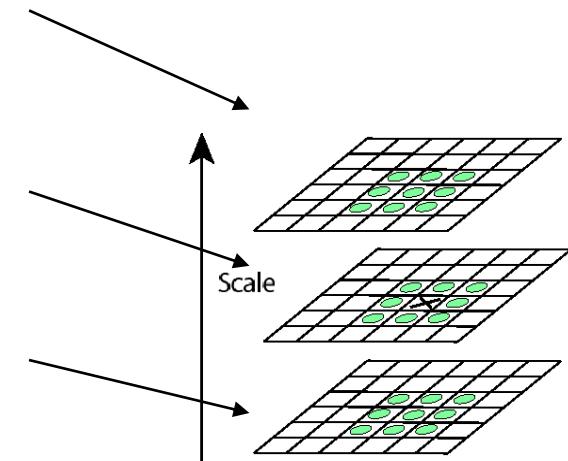
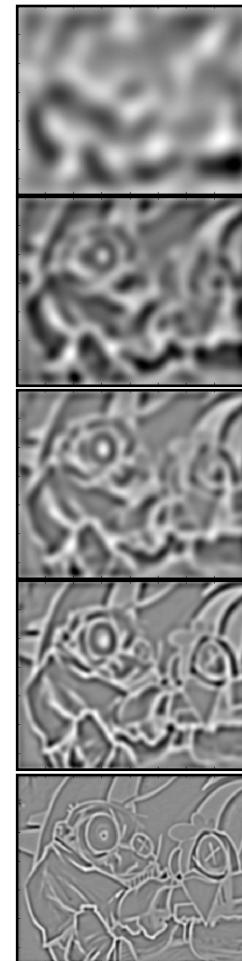


# Laplacian-of-Gaussian (LoG) blob detector

1. Convolve image with scale-normalized Laplacian of Gaussian at several scales
2. Find maxima of Laplacian response in both position and scale (3D)



$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma^5$$
$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma^4$$
$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma^3$$
$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma^2$$
$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma$$

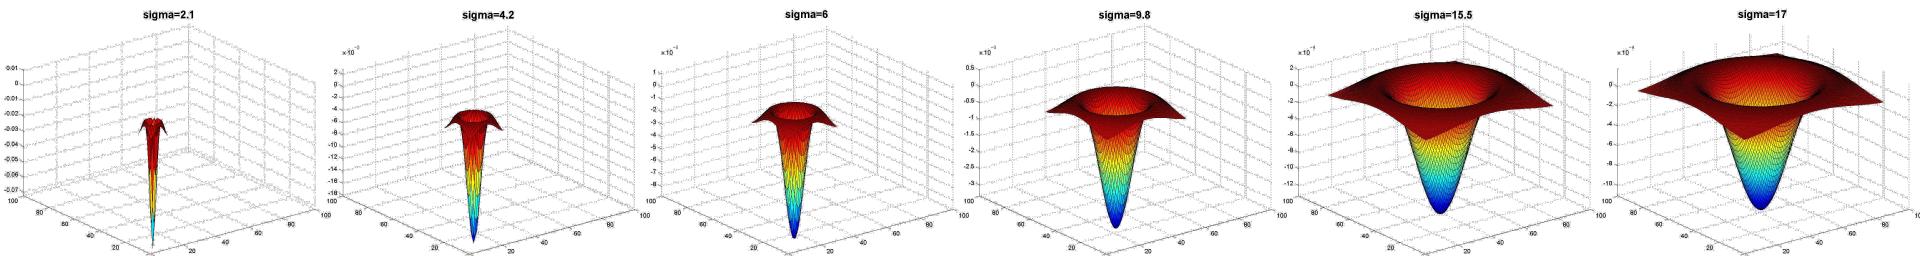


List of  $(x, y, \sigma)$   
position and scale

T. Lindeberg (1998). "Feature detection with automatic scale selection." *International Journal of Computer Vision*

Source: Mikolajczyk

# Example



Full size

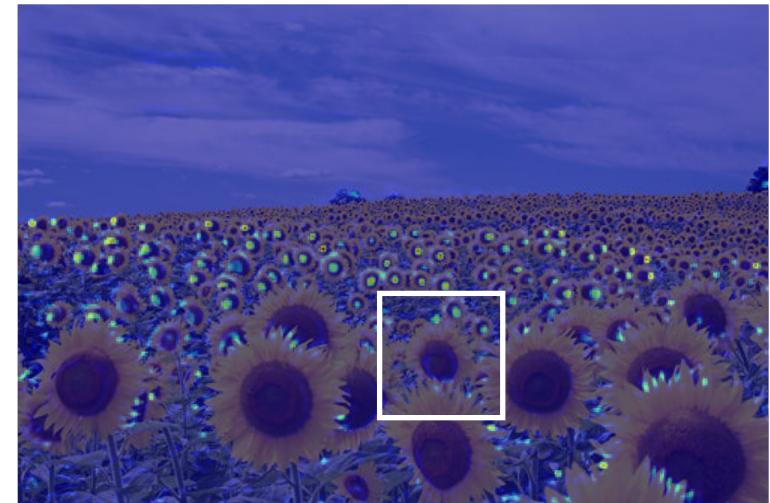
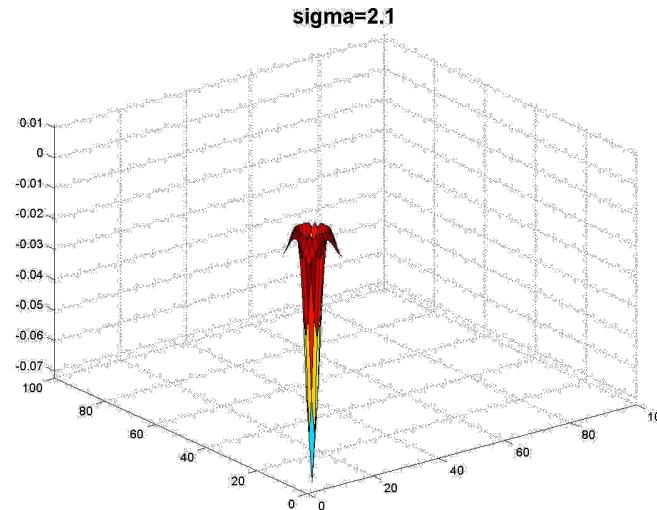


3/4 size

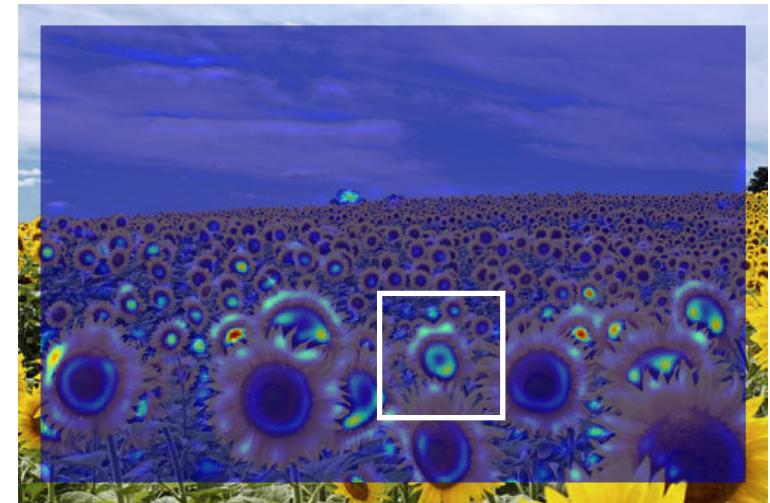
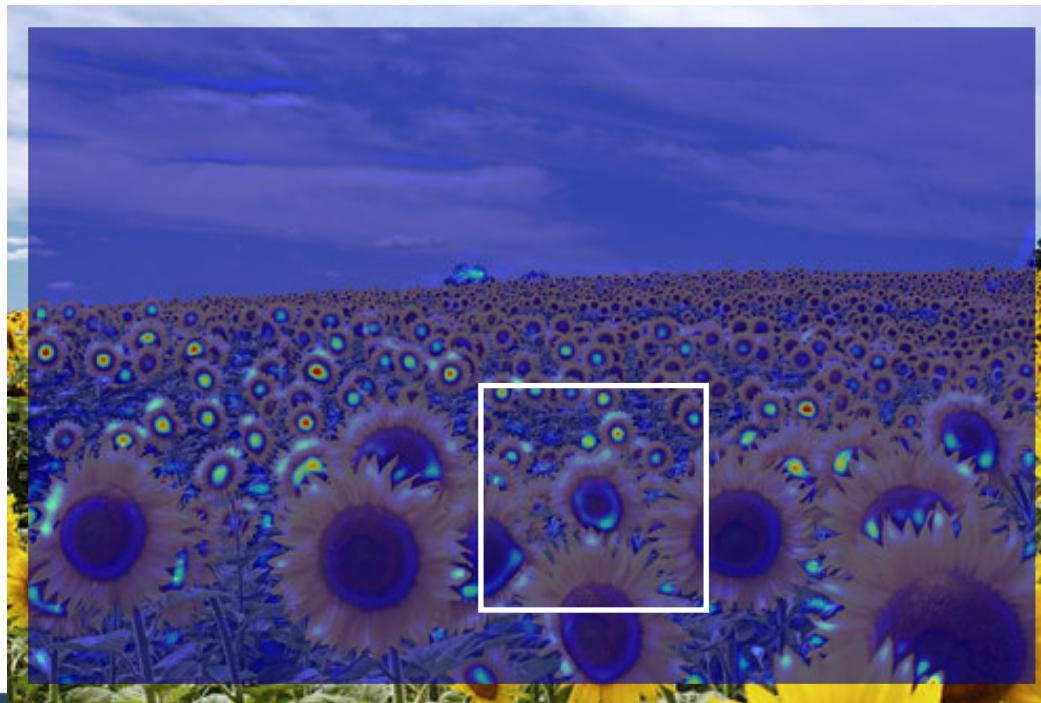
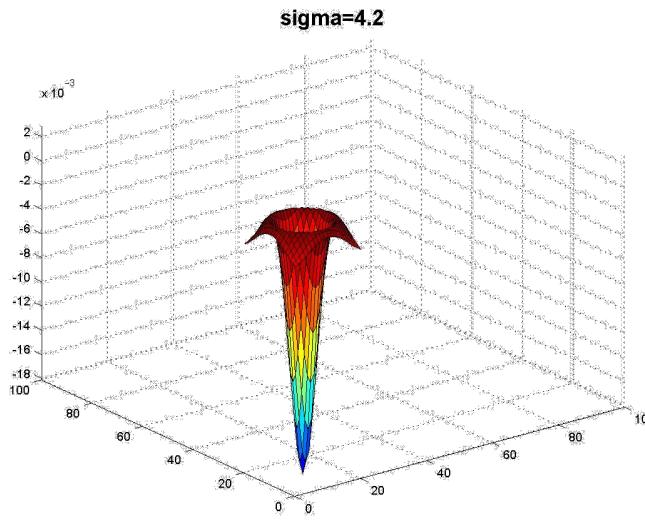


jet color scale  
blue: low, red: high

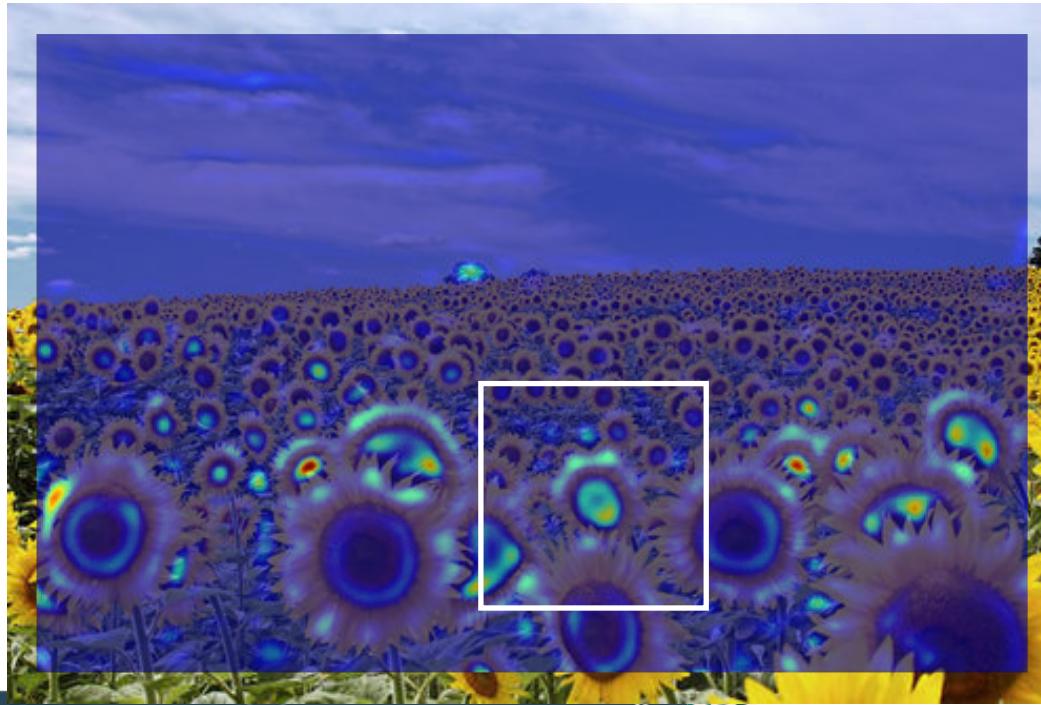
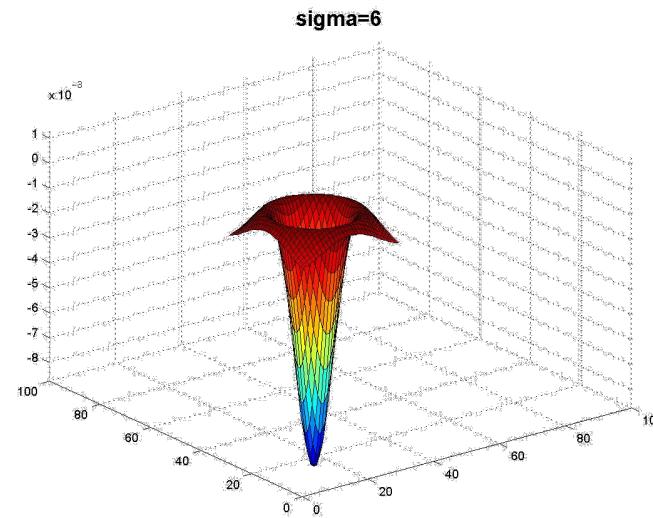
# Example



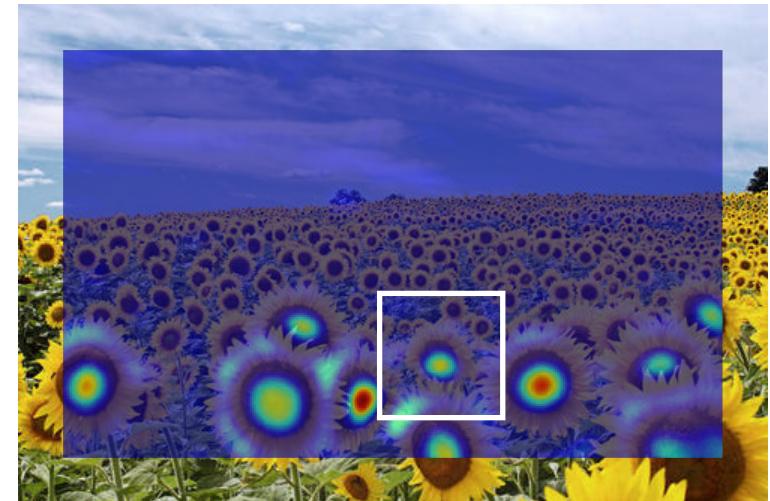
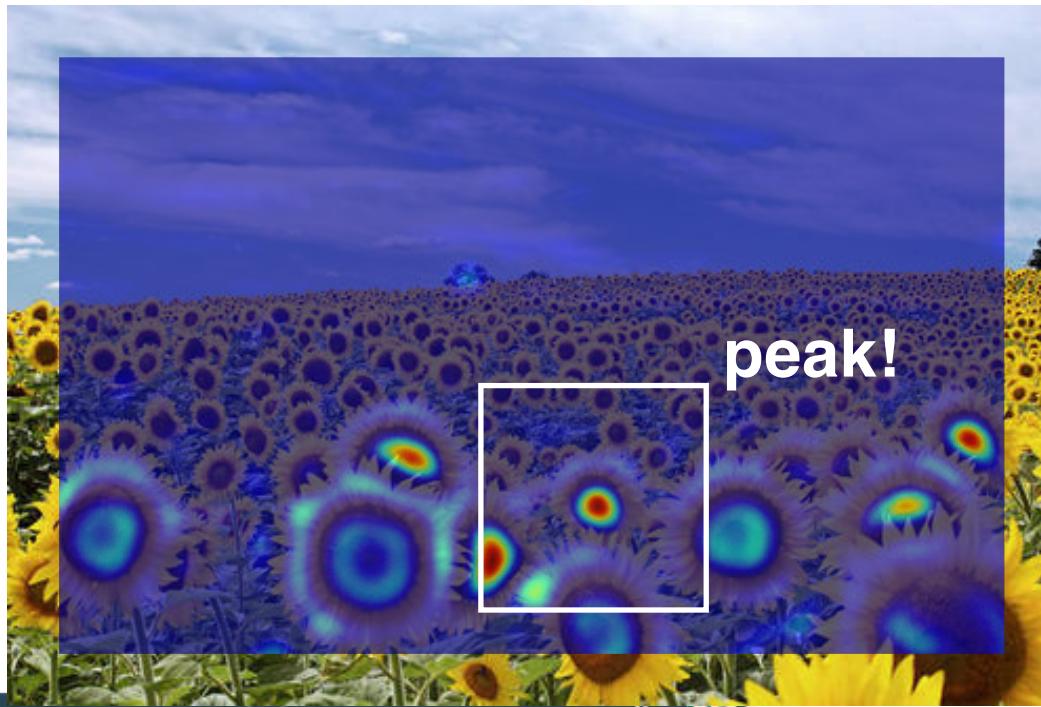
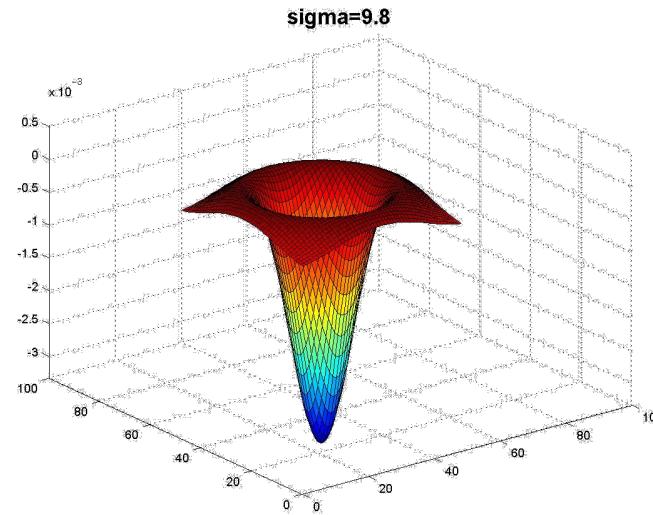
# Example



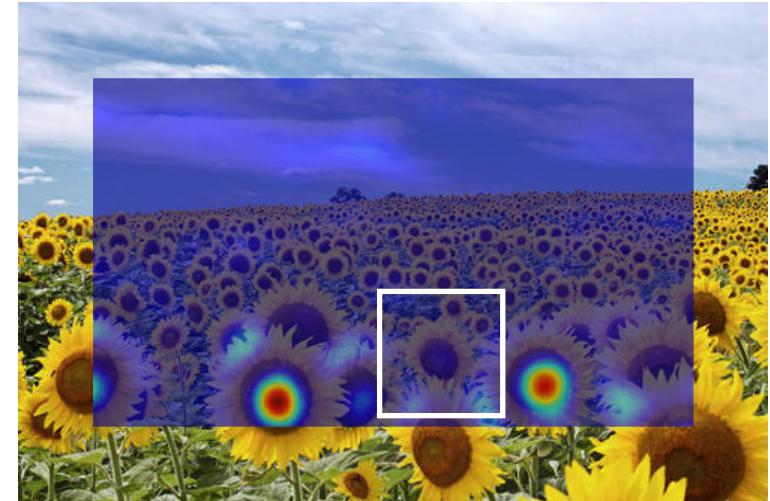
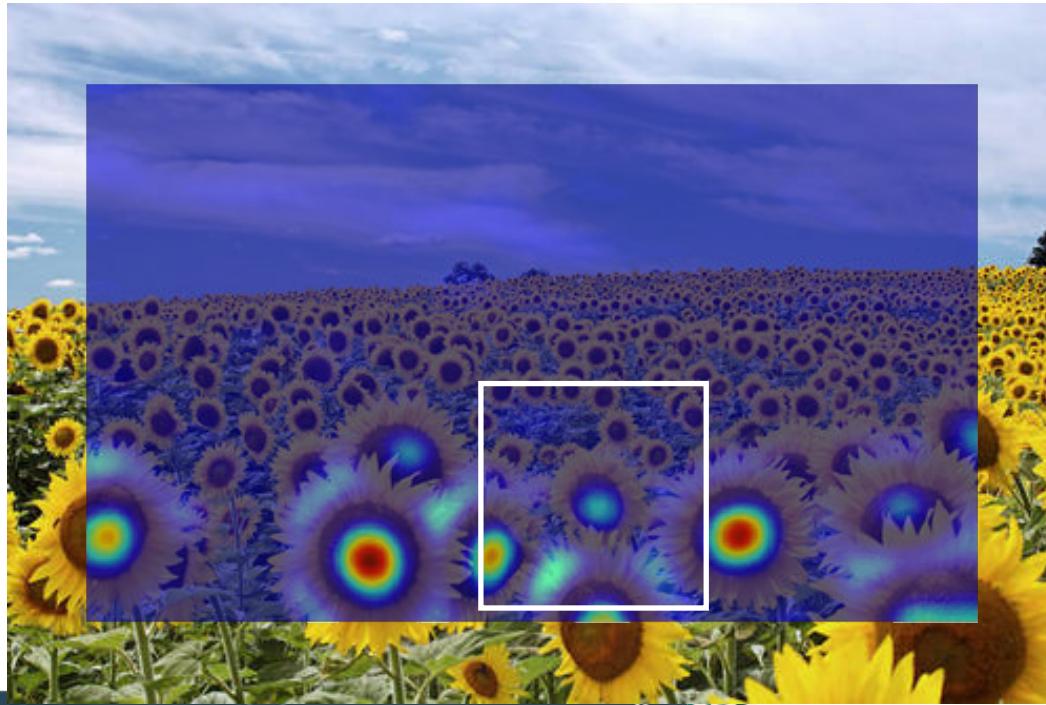
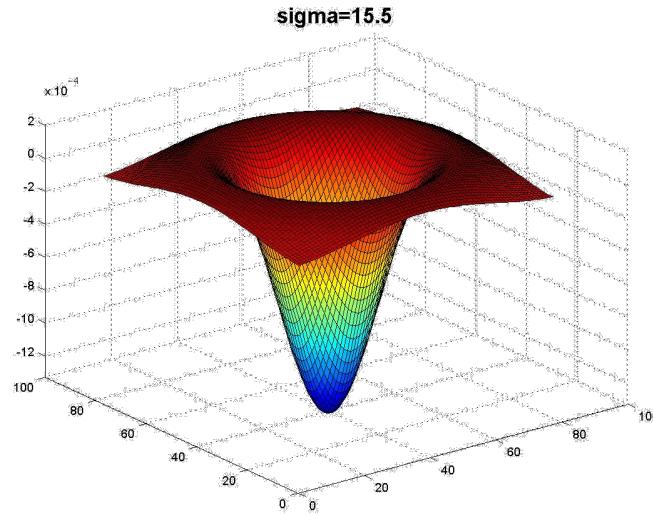
# Example



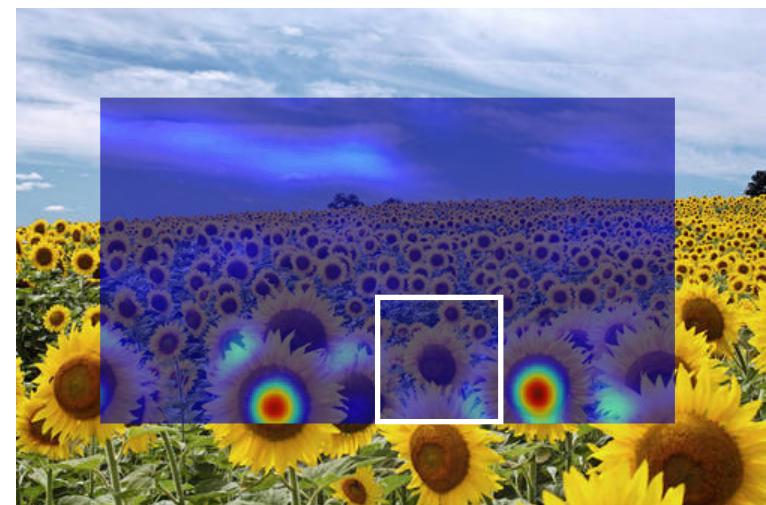
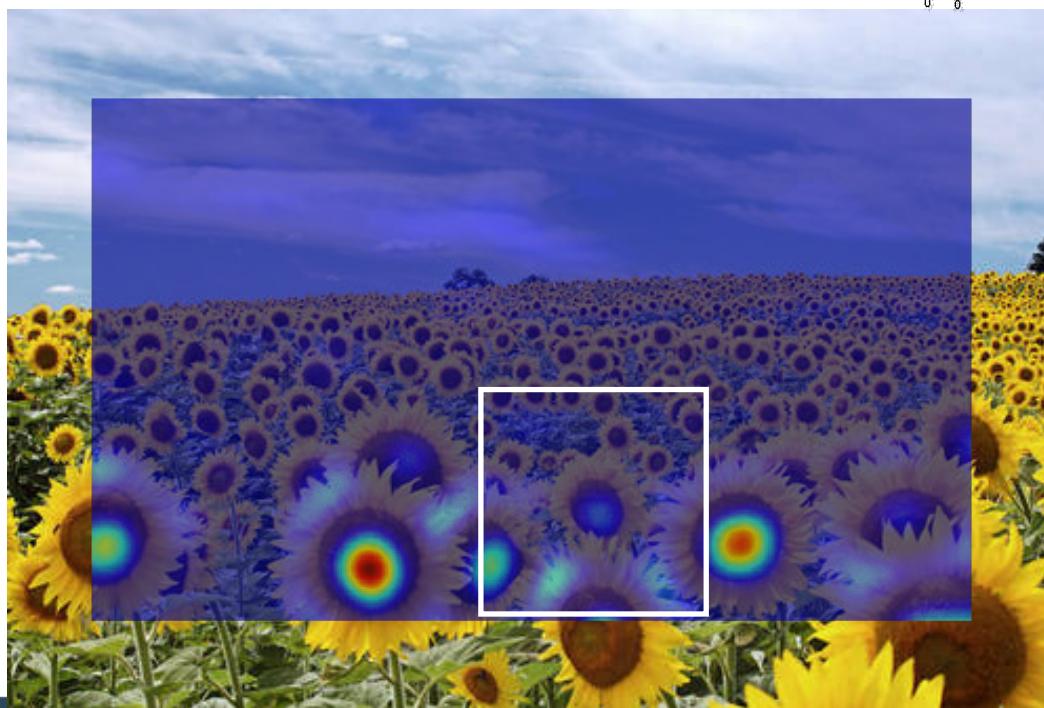
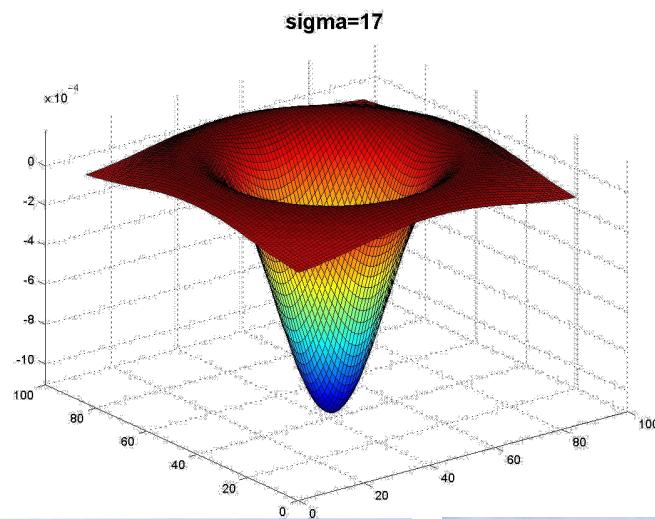
# Example



# Example

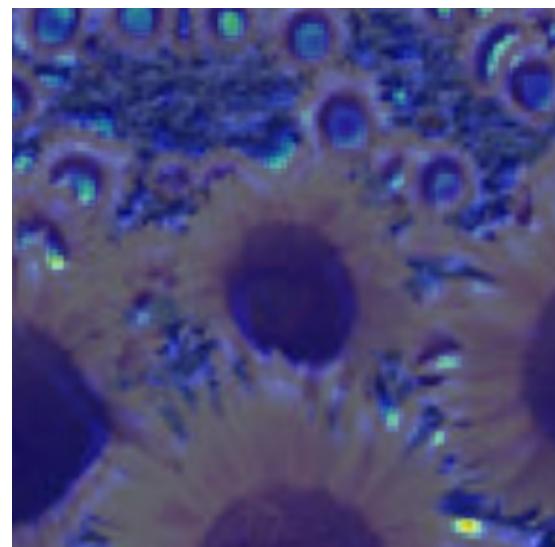


# Example

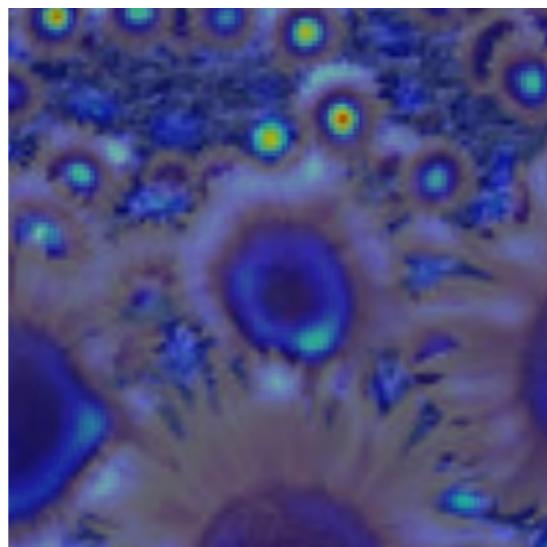


# Example

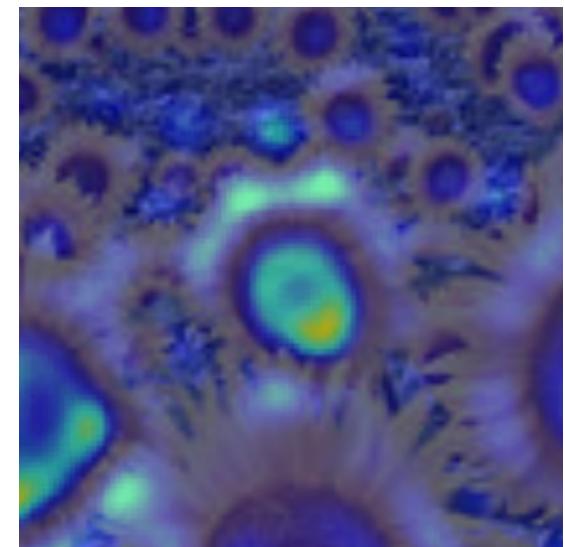
2.1



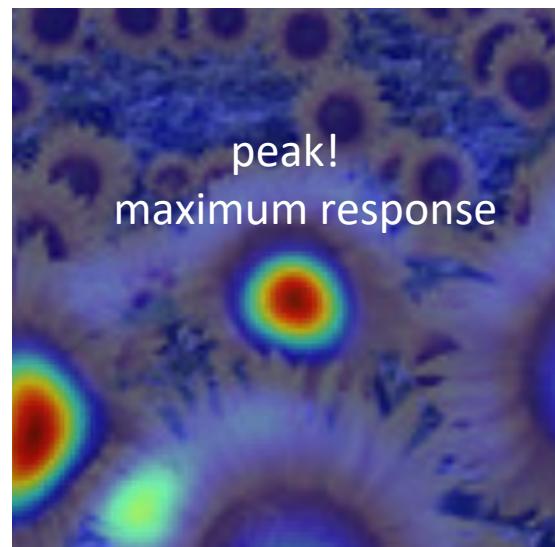
4.2



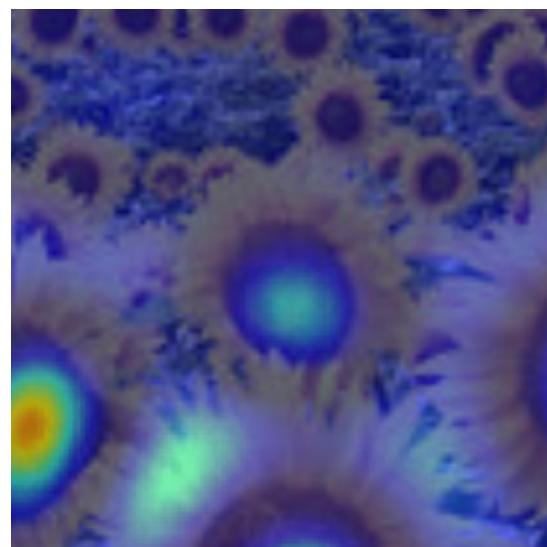
6.0



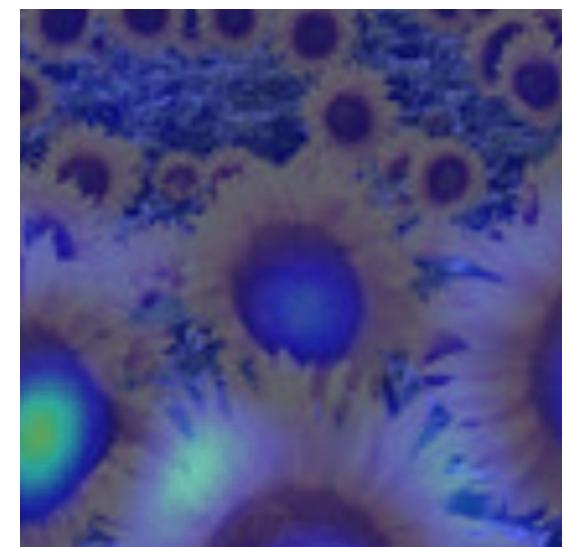
9.8



15.5

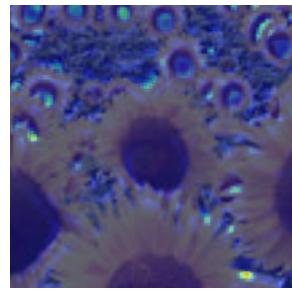


17.0

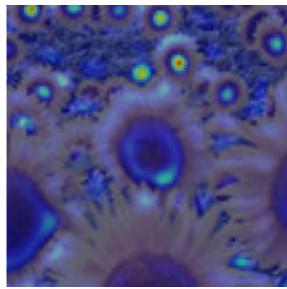


# Example: optimal scale

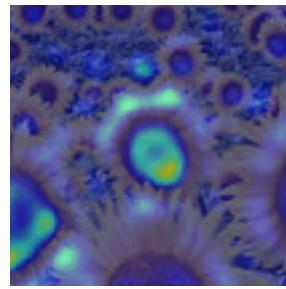
2.1



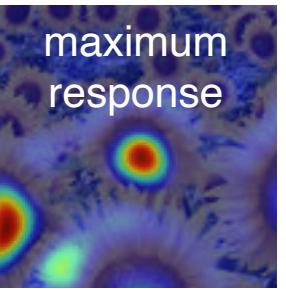
4.2



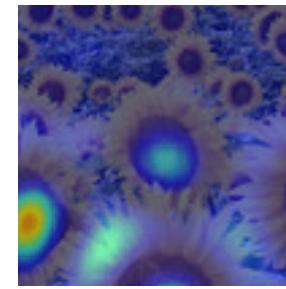
6.0



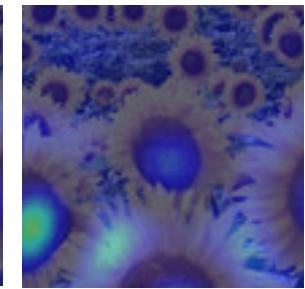
9.8



15.5

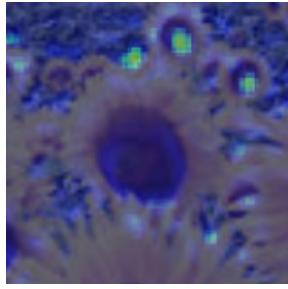


17.0

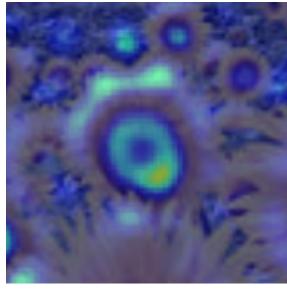


Full size image

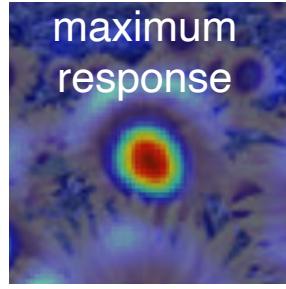
2.1



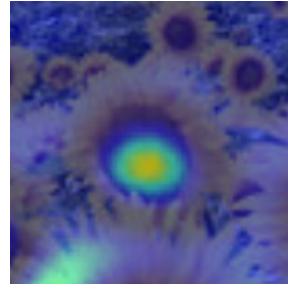
4.2



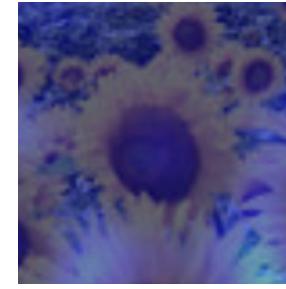
6.0



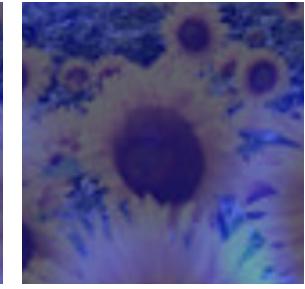
9.8



15.5



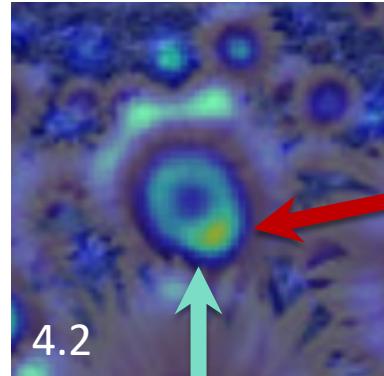
17.0



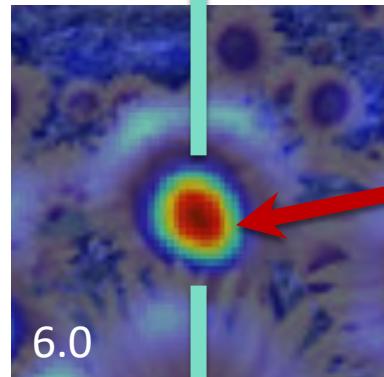
3/4 size image

# Example:

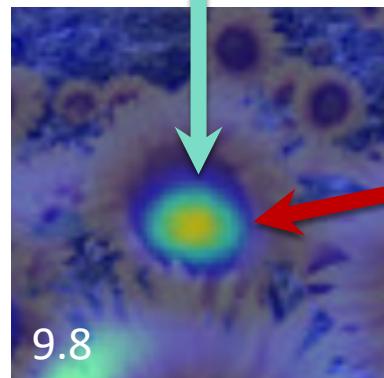
Maximum across scales  
3/4 size image



local maximum



local maximum



local maximum

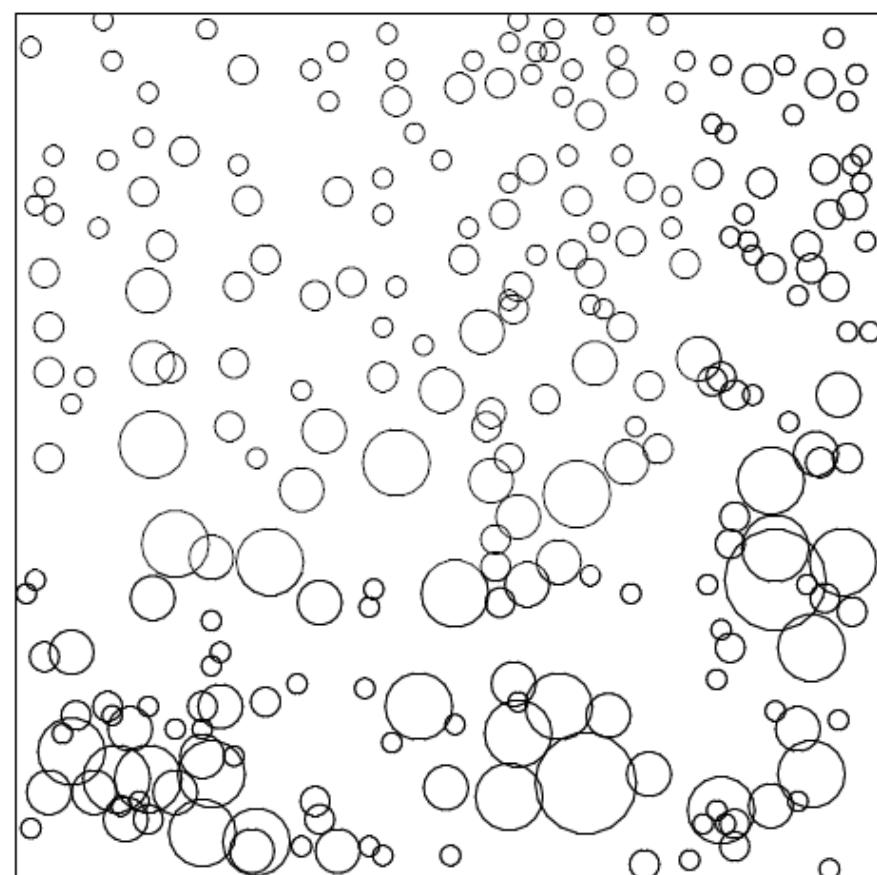


## Example: result

*original image*



*scale-space maxima of  $(\nabla_{norm}^2 L)^2$*

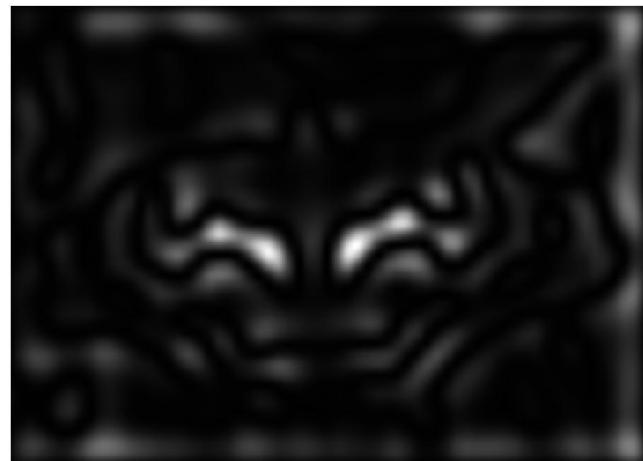
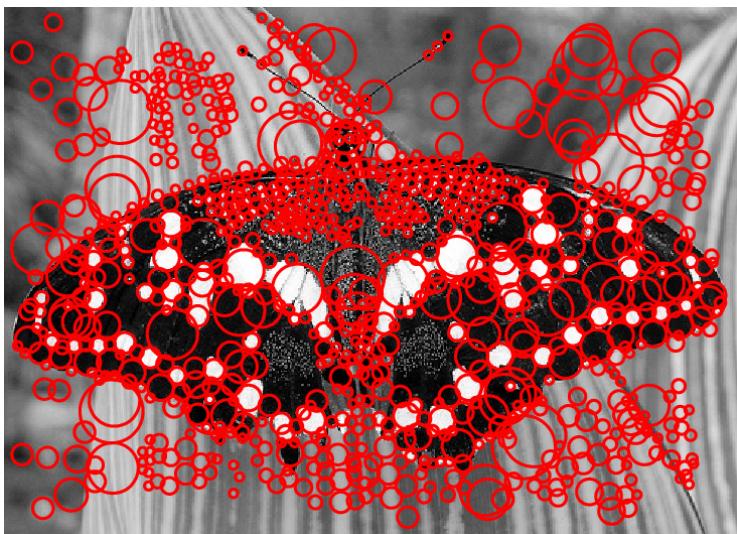


Source: S. Lazebnik

# Scale-space blob detector: example



sigma = 2



sigma = 15

# Efficient implementation

- Function for determining scale  $f = \text{kernel} * \text{image}$
- Kernels:

Laplacian of Gaussian

$$L = \sigma^2 \left( G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

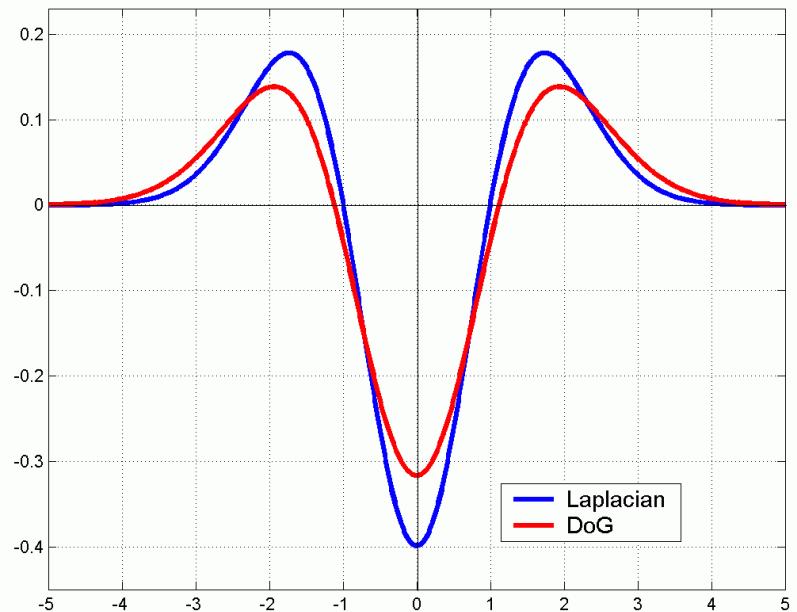
- We can efficiently approximate the Laplacian with a difference of Gaussians

Difference of Gaussians

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

$$G(x, y, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$k = \sqrt{2}$$

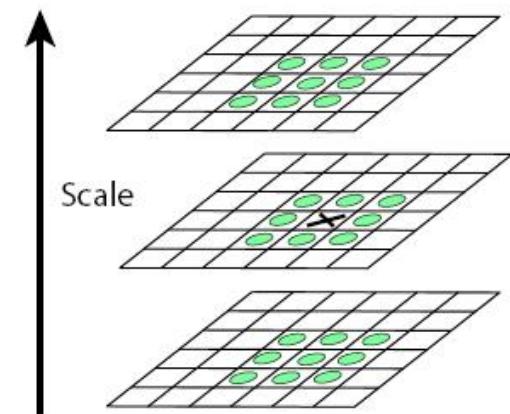
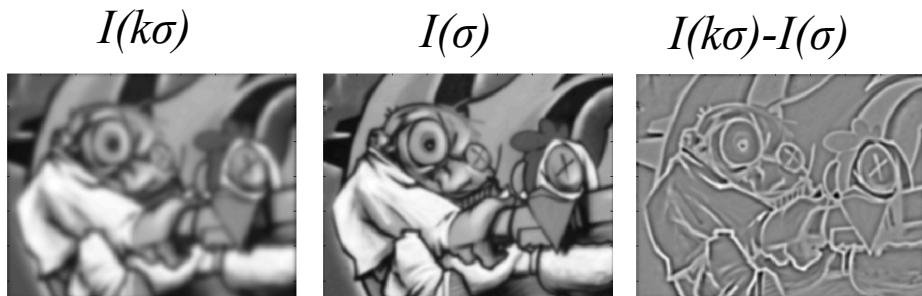


Note: both kernels are invariant to *scale* and *rotation*

Source: B. Leibe

# Difference of Gaussians (DoG) [Lowe 04]

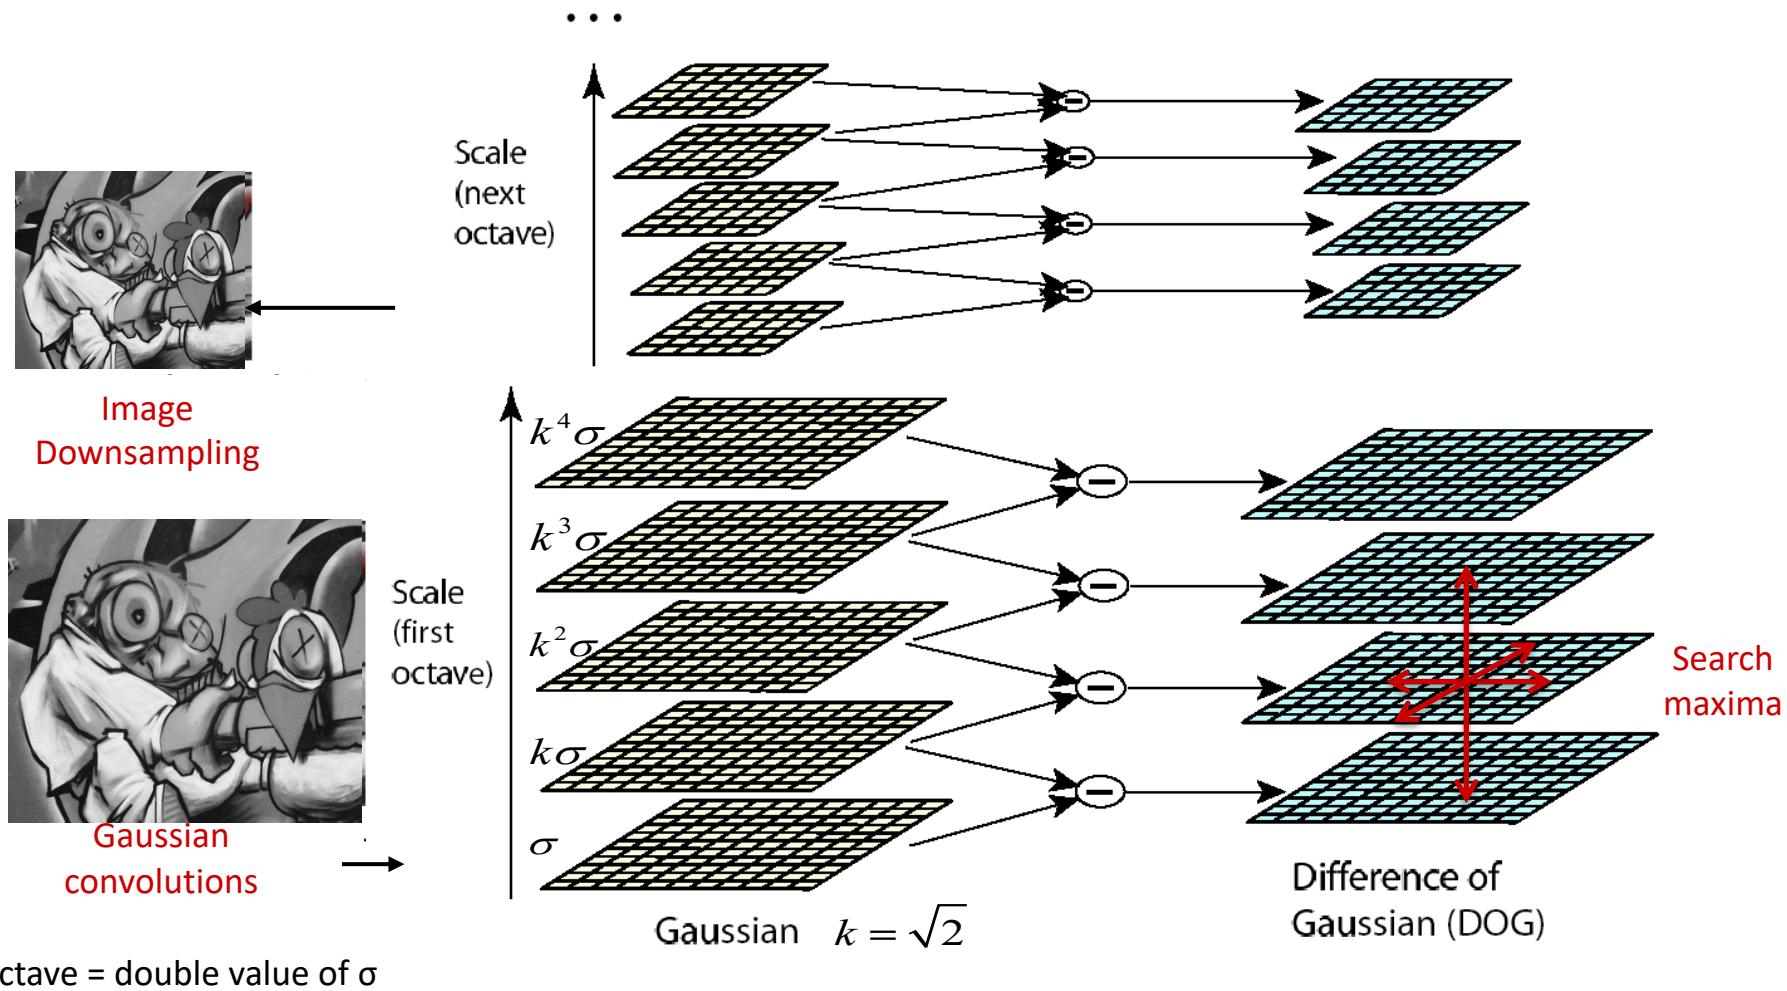
- This is used in Lowe's SIFT pipeline for keypoint detection
- Advantage
  - No need to compute second derivatives
- Key point localization with DoG
  - Detect maxima of DoG in scale space
  - Reject points with low contrast (threshold)
  - Eliminate responses along edges



Candidate keypoints:  
list of  $(x, y, \sigma)$

# Difference of Gaussians (DoG) [Lowe 04]

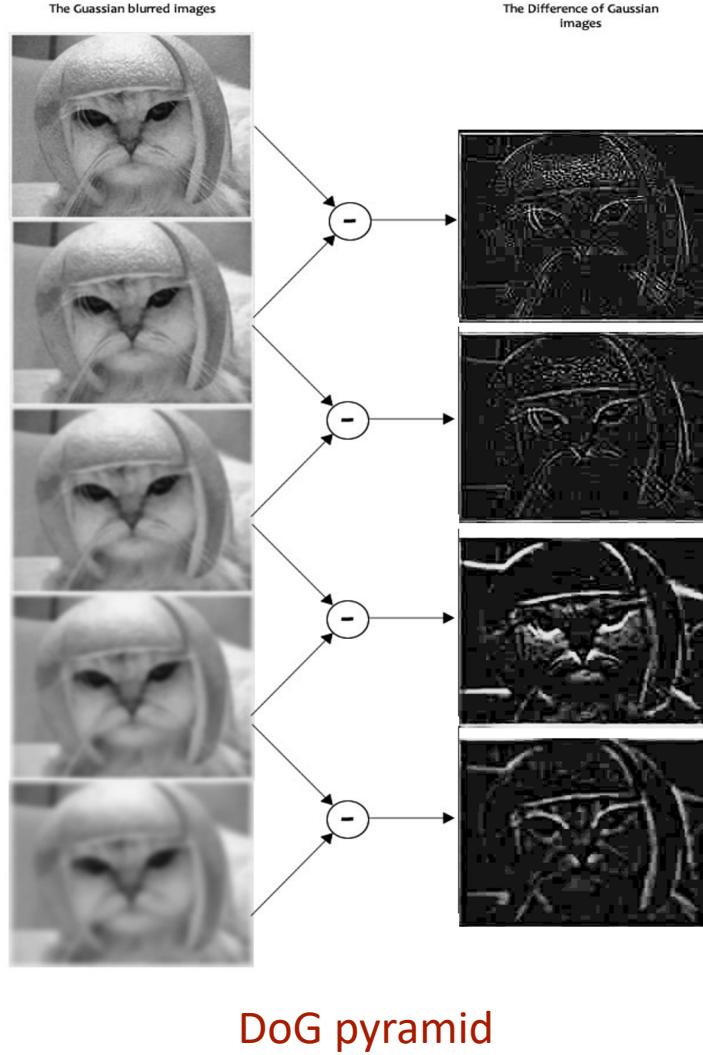
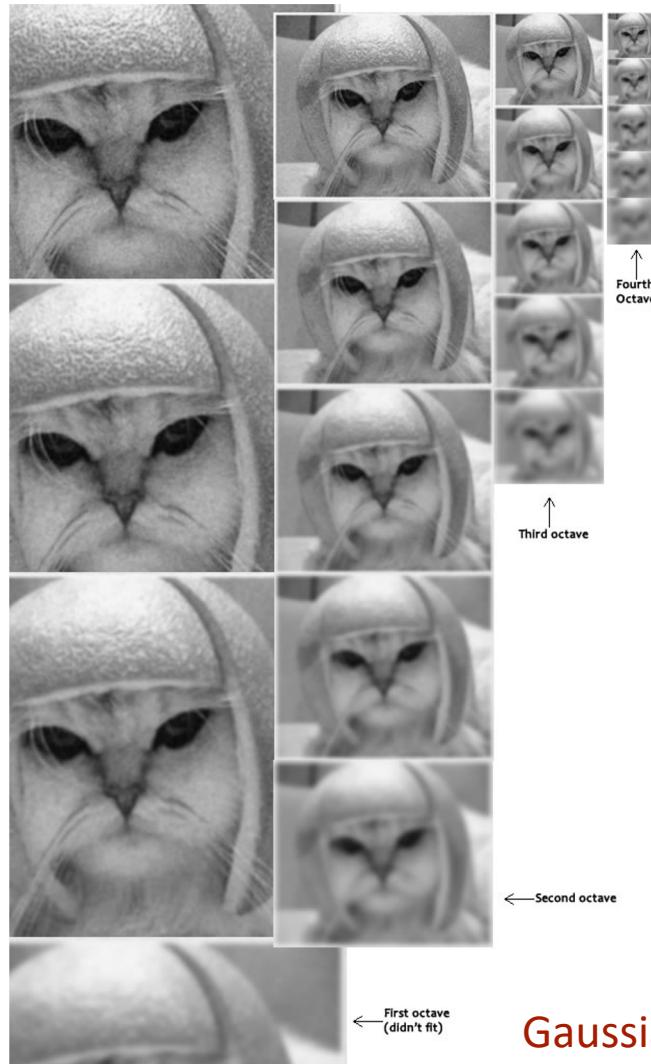
- Efficient Computation in Gaussian scale pyramid



David G. Lowe. "Distinctive image features from scale-invariant keypoints." *IJCV*, 2004

Source: Mikolajczyk

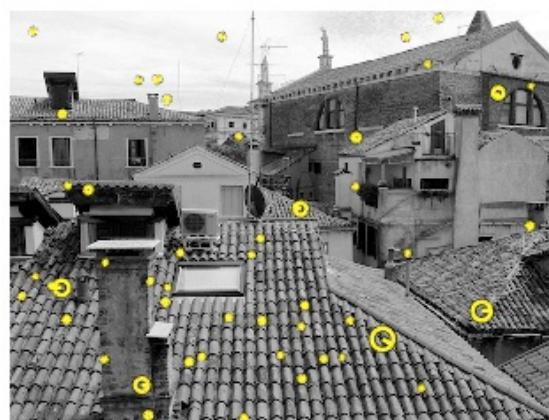
# Difference of Gaussians (DoG) [Lowe 04]



# From feature detection to feature description

- Feature description: extract vector feature descriptor surrounding each interest point
- To recognize the same pattern in multiple images, we need to match appearance “signatures” in the neighborhoods of extracted keypoints
  - But corresponding neighborhoods can be related by a scale change or rotation
  - We want to *normalize* neighborhoods to make signatures invariant to these transformations

Scale invariant feature transform (SIFT) descriptor [Lowe 2004]



Interest points and their scales  
and orientations  
(random subset of 50)



SIFT descriptors

Image: <http://www.vlfeat.org/overview/sift.html>

# SIFT detector / descriptor [Lowe 04]

## Scale Invariant Feature Transform

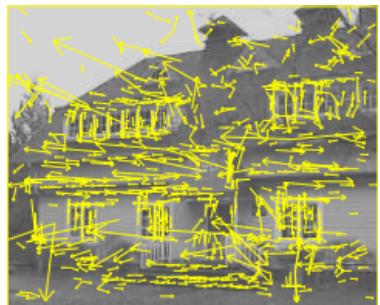
- **Detector**

1. Find maxima of DoG in scale-space
2. Keypoint localization and filtering
  - Improve localization (pixel accuracy at best, corresponds to several pixels in base image at higher scales)
  - Remove points with low contrast (threshold)
3. Remove points with strong responses along edges

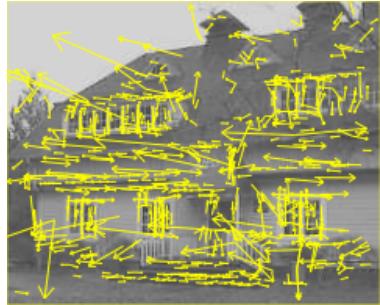
233x189  
image



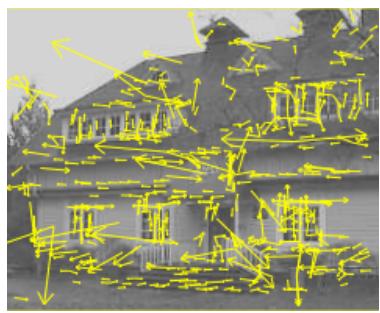
832 DoG  
extrema



729 after  
peak value  
threshold



536 after  
removing  
edge  
responses

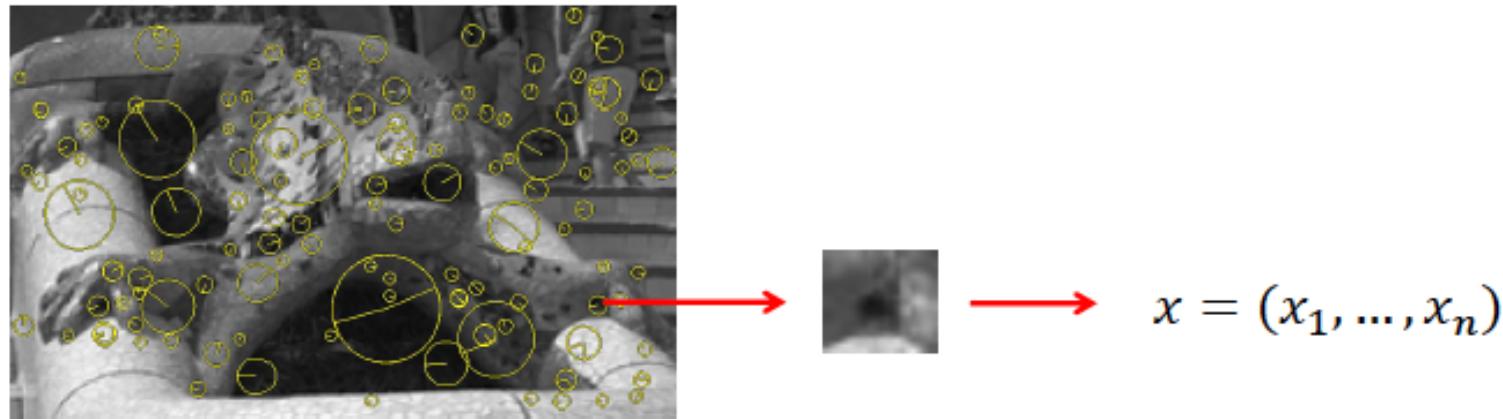


- **Descriptor**

- Orientation assignment
  - Remove effects of orientation and scale
- Create descriptor
  - Using histograms of orientations

# Sift descriptor

- Based on gradient information of a patch defined around every keypoint at a given scale

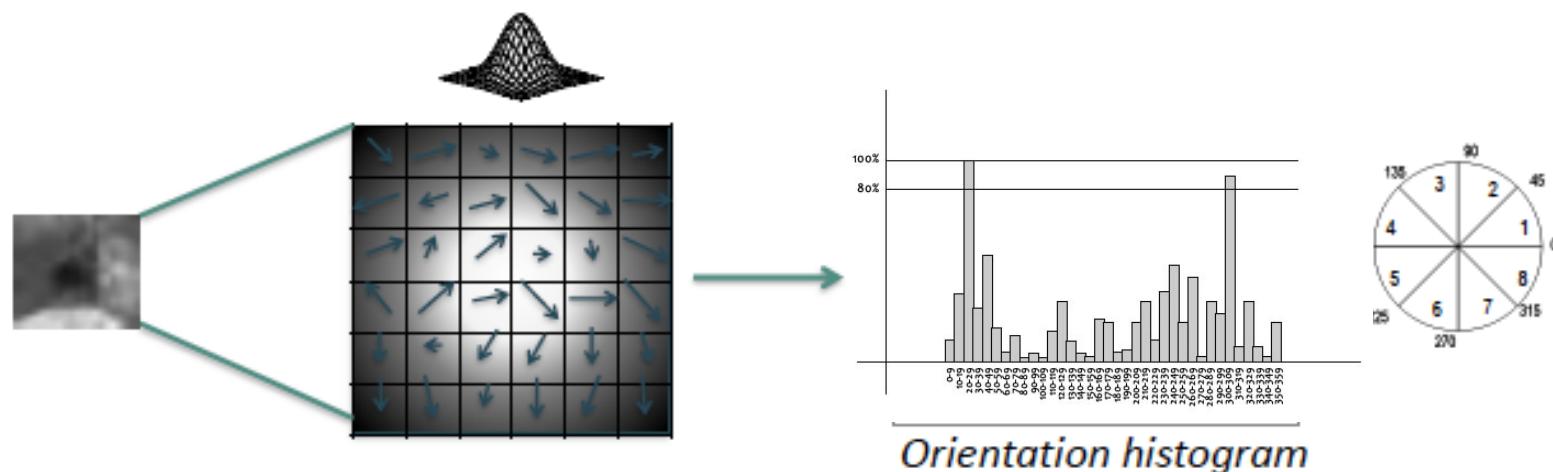


- Orientation histogram
- Rotation invariance: estimating keypoint dominant orientation
- Feature vector
  - Basic computation
  - Trilinear interpolation
  - Normalization

Slide credit: E. Valveny

# Sift descriptor

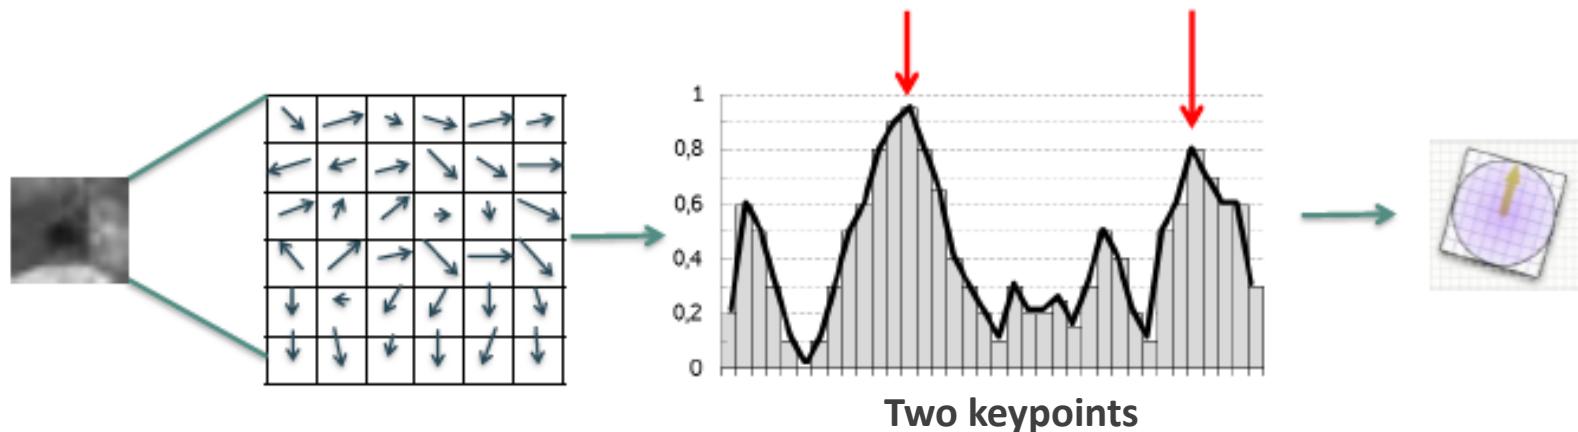
- Orientation histogram:
  - Gradient (orientation and magnitude) for each pixel
  - Every bin of the orientation histogram accumulates gradient magnitudes weighted by a gaussian centered on the patch (36 bins)



Slide credit: E. Valveny

# Sift descriptor

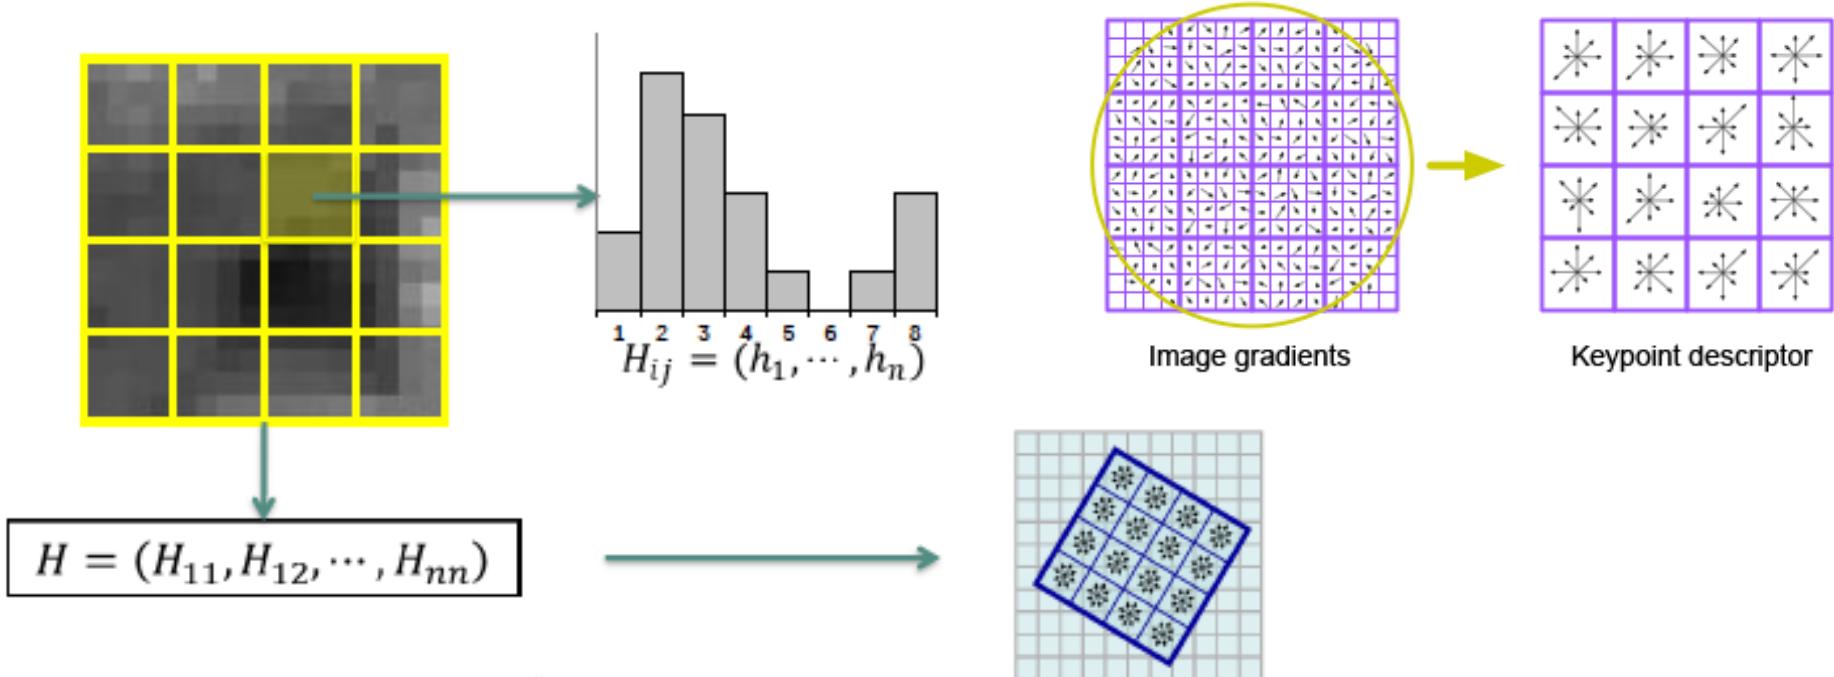
- Rotation invariance: estimating keypoint dominant orientation
  - Orientation histogram with 36 bins
  - Dominant orientation: histogram maximum
  - Additional orientations: local máxima over 80% of global máximo
  - The patch is rotated according to the dominant orientation before obtaining the feature vector



Slide credit: E. Valveny

# Sift descriptor

- Feature vector: basic computation
  - The patch is divided into  $n \times n$  cells (usually  $n=4$ )
  - One orientation histogram per cell (usually 8 bins in each histogram)
  - Feature vector: concatenation of all cell histograms ( $4 \times 4 \times 8 = 128$  dimensions)

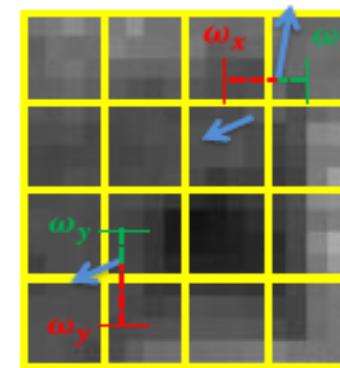
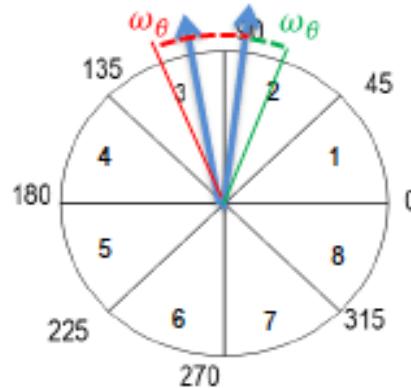


Slide credit: E. Valveny

# Sift descriptor

- Feature vector: trilinear interpolation
  - Sensitivity to noise: gradients with very similar orientation / location could be assigned to different bins /cells
  - Interpolation: each gradient magnitude is assigned to the two closest bins /cells with a weight proportional to the distance to the center of the bin /cell
  - Final contribution of a gradient to a bin of a cell

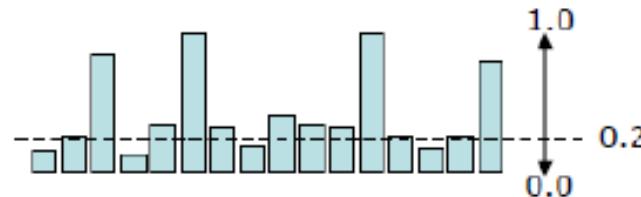
$$w = g(x, y) \cdot \omega_\theta \cdot \omega_x \cdot \omega_y$$



Slide credit: E. Valveny

# Sift descriptor

- Feature vector normalization
  - L2 normalization to cancel out contrast variations
  - Clip values greater than 0.2 to minimize the effect of very large gradients
  - Final L2-renormalization



- SIFT descriptor is composed of: Location, Scale, Orientation, Descriptor

Slide credit: E. Valveny

# What have we learned today?

- **Edges**
  - Image gradients and the effect of noise
  - Derivative of Gaussian and Laplacian of Gaussian
  - Canny detector
- **Corners and blobs**
  - to match the same scene or object under different viewpoint, it's useful to first detect interest points (keypoints)
  - we discussed these detectors:
    - Harris corner detector: translation and rotation but not scale invariant
    - Scale invariant interest points: Laplacian of Gaussian and Lowe's DoG
- **Harris' approach** computes  $I_x^2$ ,  $I_y^2$  and  $I_x I_y$ , and blurs each one with a Gaussian. The matrix  $M$  characterizes the shape of E for a window around  $(x,y)$ . Compute 'cornerness' scores for each  $(x,y)$  as  $R = \det(M) - \alpha \text{trace}(M)^2$ . Find  $R >$  threshold and do non-maxima suppression to find corners
- **Lowe's approach** creates a Gaussian pyramid with s blurring levels per octave, computes difference between consecutive levels and finds local extrema in space and scale