

Developing a CNN from scratch

Group members: Jose Manuel López, Alex Martín, Marcos V. Conde





INDEX

- Statement of the problem
- Base model
- Fixed parameters
- Variable parameters
- Regularization
- Literature model
- Summary of results
- Conclusions



Statement of the problem

We want to use deep learning techniques for image classification.

Given a dataset, how can we a CNN model that optimizes the tradeoff *model-complexity vs vs model performance (accuracy)* ? Are always literature/SOTA architectures the best option?

To find this answer we run manual **Architecture Search** (AS) experiments, and try +15 different CNN configurations using well-known techniques like: Convolutional Blocks, Batch Normalization, Dropout, etc.

We use a common experimentation setup in order to compare fairly all the models: same validation split and validation metrics, same data-loader and augmentations, same number of epochs and optimizer etc.

At the end of this experiment, we analyze the **ablation study** results and take conclusions.

Base model

All the tested models have the same configuration in common:

- N Convolutional Layers/Blocks → Feature extractors
- Global Average Pooling (GAP) layer → produces the feature vector (aka embedding).
- M dense layers → conforms the classification head.



Fig: Dissecting AS into three steps: 1) exploring the search space, where operations (e.g. pooling, convolution etc) live; 2) have a good detective strategy to pick the right bits and pieces for a specific tasks; 3) have an evaluator system for any created architecture.

[Image from <https://towardsdatascience.com/neural-architecture-search-a-model-creation-company-602c6ba4f576>]



Fixed parameters

After starting to test different parameters for the CNN we fix some of them in order to focus more on the other ones:

- The **activation functions** used were always ReLu and Softmax for the last dense layer.
- The **optimizer** used was Adam with the default parameters.
- The models were trained with learning rate decay depending on the validation accuracy during 100 Epochs.

To set up our initial model we used 64 filters per convolutional layer, a kernel size of 5 and 1 dense layer with 1024 neurons before the last layer.

| | Train Accuracy | Test Accuracy | Num. Params | Ratio |
|---------------|----------------|---------------|-------------|-----------------------|
| Initial model | 0,614 | 0,656 | 79624 | $7,71 \cdot 10^{-07}$ |



Variable parameters: Number of layers

The first variable tested was the **optimal number of convolutional layer** and dense layer to be used:

| Nº of conv layers | Train Accuracy | Test Accuracy | Num. Params | Ratio |
|-------------------|----------------|---------------|-------------|-----------------------|
| 2 | 0,742 | 0,774 | 182088 | $4,08 \cdot 10^{-07}$ |
| 4 | 0,860 | 0,850 | 387016 | $2,22 \cdot 10^{-07}$ |
| 8 | 0,867 | 0,843 | 796872 | $1,07 \cdot 10^{-07}$ |

| Nº of dense layers | Train Accuracy | Test Accuracy | Num. Params | Ratio |
|--------------------|----------------|---------------|-------------|-----------------------|
| 1 | 0,897 | 0,850 | 387016 | $2,22 \cdot 10^{-07}$ |
| 2 | 0,897 | 0,854 | 1436616 | $6,17 \cdot 10^{-08}$ |
| 4 | 0,898 | 0,869 | 3535816 | $2,54 \cdot 10^{-08}$ |



Variable parameters: Number of filters

After setting the number of convolutional layers and dense layer to 4 and 3 respectively, we started to seek the optimal parameters for the number of filters in each convolutional layer:

| Nº of filters | Train Accuracy | Test Accuracy | Num. Params | Ratio |
|---------------|----------------|---------------|-------------|-----------------------|
| 32 | 0,862 | 0,844 | 1170920 | $7,36 \cdot 10^{-08}$ |
| 64 | 0,897 | 0,854 | 1436616 | $6,25 \cdot 10^{-08}$ |
| 128 | 0,883 | 0,846 | 2428808 | $3,64 \cdot 10^{-08}$ |
| 256 | 0,870 | 0,862 | 6256392 | $1,39 \cdot 10^{-08}$ |



Regularization

After the previous testing, the best model was like this:

- 4 convolutional layers with 256 filters
- 2 dense layers with 1024 neurons
- Kernel size 5

The next step was to **introduce regularization techniques** into the model. We tested the following methods one by one and combining them:

- **MaxPooling:** Reduce spatial dimension → compact features maps
- **Batch Normalization:** train models faster and more stable through normalization of the layers' inputs by re-centering and re-scaling. Work by Ioffe and Szegedy et al.
- **Dropout:** randomly sets input units to 0 with a frequency of rate at each step during training time, which helps prevent overfitting.



Regularization

| Regularization method | Train Accuracy | Test Accuracy | Num. Params | Ratio |
|--|----------------|---------------|-------------|-----------------------|
| MaxPooling | 0,920 | 0,864 | 6256392 | $1,47 \cdot 10^{-08}$ |
| Dropout (0.2) | 0,886 | 0,861 | 6256392 | $1,42 \cdot 10^{-08}$ |
| BatchNorm | 0,987 | 0,879 | 6268680 | $1,57 \cdot 10^{-08}$ |
| MaxPooling + Dropout (0.2) | 0,927 | 0,867 | 6256392 | $1,48 \cdot 10^{-08}$ |
| MaxPooling + BatchNorm | 0,970 | 0,887 | 6268680 | $1,55 \cdot 10^{-08}$ |
| MaxPooling + BatchNorm + Dropout (0.2) | 0,960 | 0,890 | 6268680 | $1,53 \cdot 10^{-08}$ |



Variable parameters: Kernel size

Kernel size ablation studies:

| Kernel size | Train Accuracy | Test Accuracy | Num. Params | Ratio |
|-------------|----------------|---------------|-------------|-----------------------|
| 3 | 0,900 | 0,844 | 3098376 | $2,90 \cdot 10^{-08}$ |
| 5 | 0,886 | 0,861 | 6256392 | $1,42 \cdot 10^{-08}$ |
| 7 | 0,830 | 0,829 | 10993416 | $7,55 \cdot 10^{-09}$ |
| 9 | 0,758 | 0,763 | 17309448 | $4,38 \cdot 10^{-09}$ |

Examples of prediction


- Our model has a fantastic **performance**. However, there are some scenarios that can lead to a misclassification.
- **Street**, **tallbuilding** and **inside_city** are often confused because they all relate to urban landscapes.
- **Street** and **highway** because they both rely on roads.
- **Forest** and the rest of classes if there are high presence of vegetation





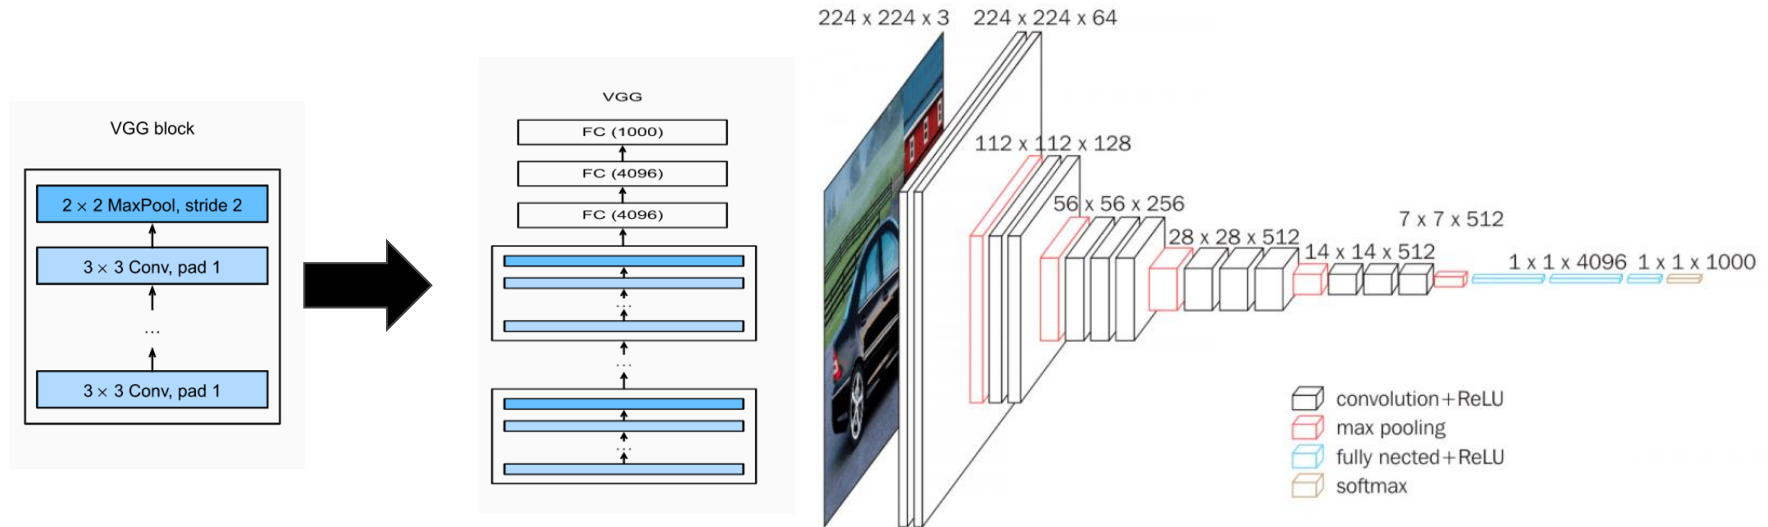
Examples of prediction

We can check more deeply those kind of errors by looking at the **confusion matrix**.



Literature model

We compare the obtained model with well-known architectures in the literature like VGG by Karen Simonyan and Andrew Zisserman, 2014.





Our model vs Literature model

We compare the obtained model with well-known architectures in the literature, like VGG by Karen Simonyan and Andrew Zisserman, 2014.



Our custom model achieves competitive results in comparison with the well-known VGG19.

We think the gap between training-validation is due to the difference in the model complexity, and the Dropout that we used. Our model does not overfit training and therefore can generalize decently.

| Model | Train Accuracy | Test Accuracy | Train Loss | Test Loss | Num, Params | Ratio |
|-------|----------------|---------------|------------|-----------|-------------|-----------------------|
| VGG | 0,952 | 0,879 | 0,141 | 0,460 | 24 M | $3,92 \cdot 10^{-09}$ |
| Ours | 0,886 | 0,861 | 0,303 | 0,437 | 6 M | $1,42 \cdot 10^{-08}$ |

Summary


| | | | | | | | |
|--|----------|----------|----------|----------|----------|---------|--------------|
| 4 Conv 256 filters 5 kernel size 2 dense 1024 neurons with BatchNorm | Model 11 | 0.987241 | 0.878563 | 0.061328 | 0.440767 | 6268680 | 1.574878e-08 |
| Conv 4 filters 256 kernel size 5 dense 2 neurons with MaxPooling and Dropout | Model 12 | 0.926635 | 0.867410 | 0.196407 | 0.408125 | 6256392 | 1.481101e-08 |
| 4 Conv 256 filters 5 kernel size 2 dense 1024 neurons with MaxPooling and BatchNorm | Model 13 | 0.970229 | 0.887237 | 0.102440 | 0.398179 | 6268680 | 1.547740e-08 |
| 4 Conv 256 filters 5 kernel size 2 dense 1024 neurons with MaxPooling and Dropout | Model 14 | 0.922382 | 0.864932 | 0.226596 | 0.437865 | 6256392 | 1.474303e-08 |
| 4 Conv 256 filters 5 kernel size 2 dense 1024 neurons with BatchNorm, MaxPooling and Dropout | Model 15 | 0.960128 | 0.889715 | 0.127559 | 0.357219 | 6268680 | 1.531626e-08 |






Summary


Train Loss




Train Accuracy



Test Loss



Test Accuracy





Summary

| | Model | Train Accuracy | Test Accuracy | Train Loss | Test Loss | Num. Params | Ratio |
|--|----------|----------------|---------------|------------|-----------|-------------|--------------|
| 1 Conv 64 filters 5 kernel size 1 dense 1024 neurons | Model 0 | 0.613503 | 0.655514 | 1.075764 | 1.001122 | 79624 | 7.705007e-07 |
| 2 Conv 64 filters 5 kernel size 1 dense 1024 neurons | Model 1 | 0.736310 | 0.774473 | 0.741201 | 0.683470 | 182088 | 4.043707e-07 |
| 4 Conv 64 filters 5 kernel size 1 dense 1024 neurons | Model 2 | 0.856991 | 0.850062 | 0.419730 | 0.494306 | 387016 | 2.214355e-07 |
| 8 Conv 64 filters 5 kernel size 1 dense 1024 neurons | Model 3 | 0.867092 | 0.842627 | 0.386379 | 0.490809 | 796872 | 1.088119e-07 |
| 4 Conv 64 filters 5 kernel size 2 dense 1024 neurons | Model 4 | 0.897395 | 0.853779 | 0.303307 | 0.460579 | 1436616 | 6.246589e-08 |
| 4 Conv 64 filters 5 kernel size 4 dense 1024 neurons | Model 5 | 0.897927 | 0.868649 | 0.305947 | 0.442636 | 3535816 | 2.539517e-08 |
| 4 Conv 32 filters 5 kernel size 2 dense 1024 neurons | Model 6 | 0.862307 | 0.843866 | 0.408355 | 0.513012 | 1170920 | 7.364357e-08 |
| 4 Conv 128 filters 5 kernel size 2 dense 1024 neurons | Model 7 | 0.883041 | 0.862454 | 0.315263 | 0.407398 | 2428808 | 3.635697e-08 |
| 4 Conv 256 filters 5 kernel size 2 dense 1024 neurons | Model 8 | 0.870282 | 0.846344 | 0.382219 | 0.477707 | 6256392 | 1.391028e-08 |
| 4 Conv 256 filters 5 kernel size 2 dense 1024 neurons with MaxPooling | Model 9 | 0.920255 | 0.863693 | 0.223434 | 0.507280 | 6256392 | 1.470904e-08 |
| 4 Conv 256 filters 5 kernel size 2 dense 1024 neurons with 0.2 Dropout on FC | Model 10 | 0.886231 | 0.861214 | 0.303123 | 0.437434 | 6256392 | 1.416520e-08 |

Fig. Our log for the ablation studies.



Conclusions

- We have explored the fundamental blocks of CNNs and found an optimal setup empirically using manual Neural Architecture Search.
- We used different techniques from the most influential computer vision literature “Deep Residual Learning for Image Recognition” by *He et al.*
- A proper **experimentation setup is key** for tracking results, save/log meaningful information and take decisions about the modelling.
- Our ablation study shows that the models with: **deep feature representations** (256-filters) and **strong classification heads** (1024-space MLPs), **plus**, Batch Normalization and Dropout for training stability, achieved the most competitive performance in the **tradeoff model complexity vs accuracy**.
- Depending on the dataset and task (in this case simple classification) shallow and “compact” models, like the ones we have developed, can achieve the same results or even outperform SOTA architectures.