



Master in Computer Vision *Barcelona*

Module 3: Machine learning for computer vision


Project: Bag of Visual Words Image Classification

Lecturer: Ramon Baldrich, ramon.baldrich@uab.cat

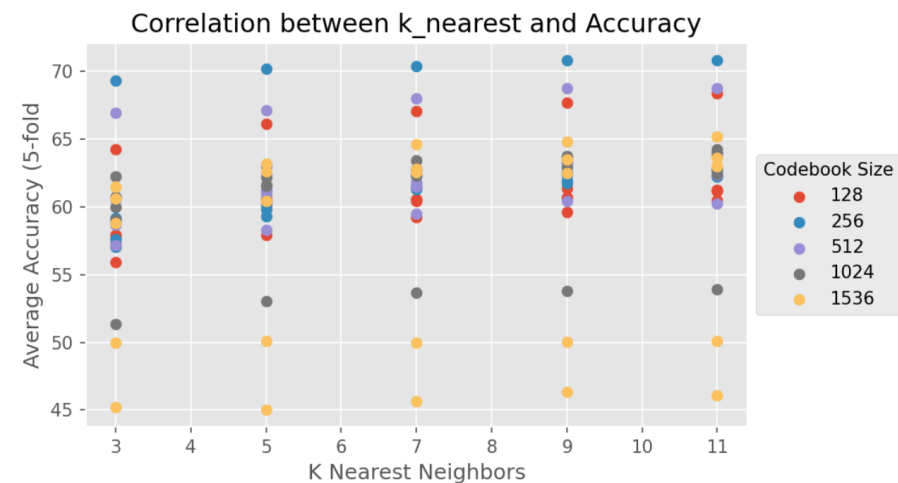
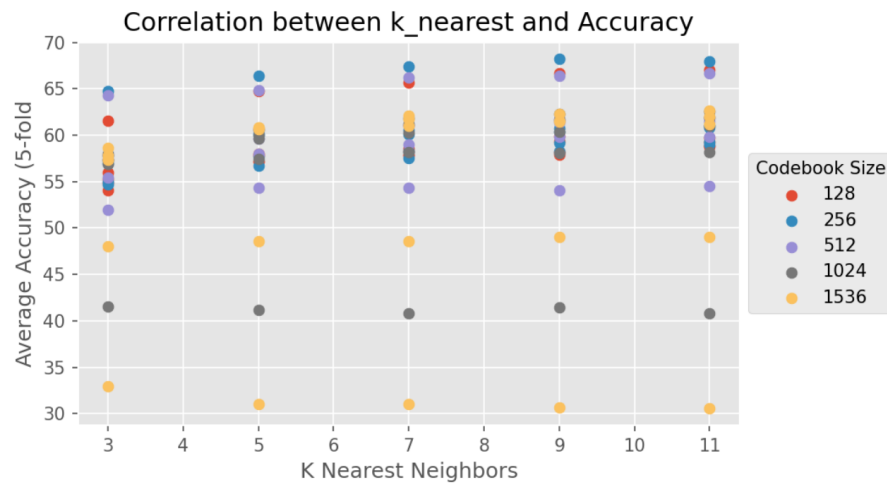
Credit to Marçal Rossinyol

S01 discussion


- Number of keypoints
 - The more the better
- Dense SIFT
 - Nearly nobody tried the role of scales!
- Codebook sizes / k-nn value
- k-nn and distances
 - Just slight differences found between point-wise distances
 - Which distance would work better for HISTOGRAMS?
- Dimensionality reduction
- Precompute stuff, store to disk!




- Amount of points (kaze)






Dense vs Non Dense

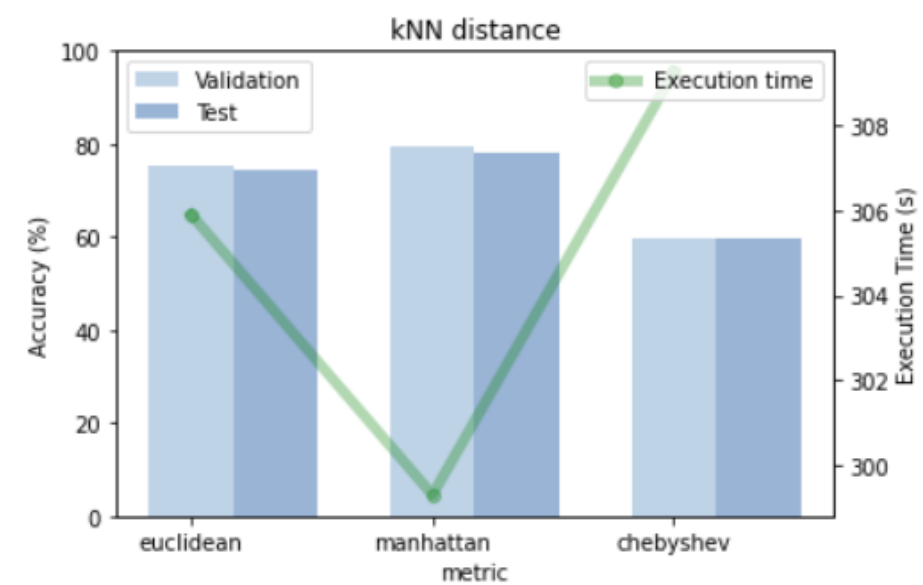
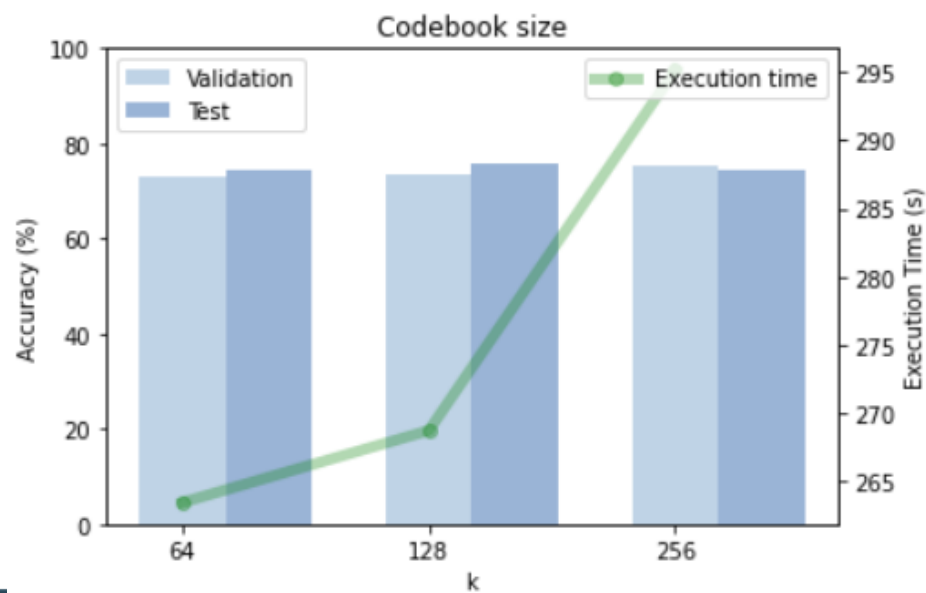
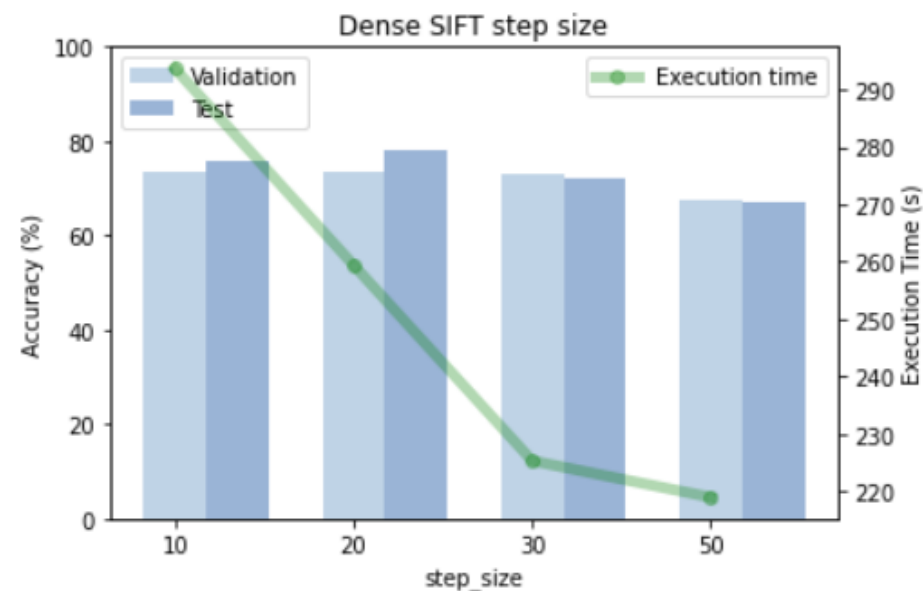
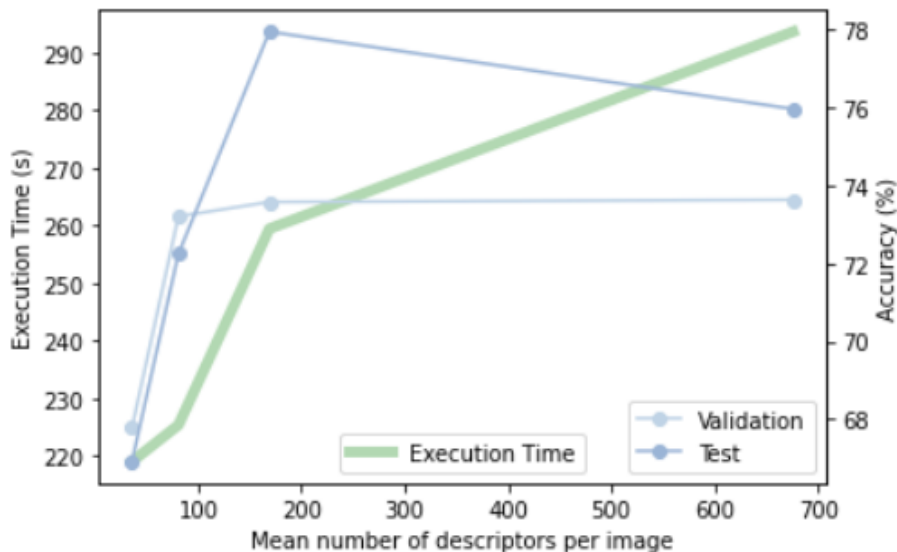


5000 kp

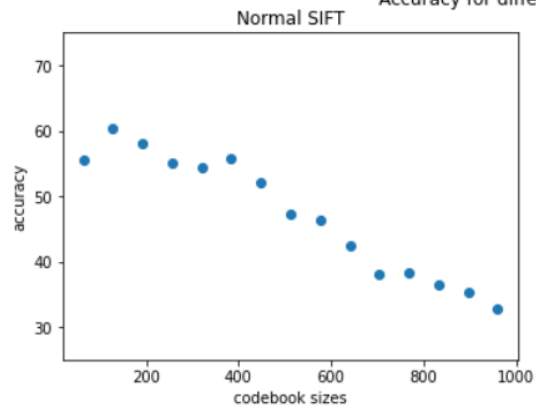


KNN
















Accuracy for different codebook sizes




Dense SIFT







Some images from coast class



Train accuracy for LogReg: 0.9946836788942052
 Test accuracy for LogReg: 0.7757125154894672

Train Accuracy: KNN = % 84, Logistic Regression = %92
Test Accuracy: KNN = % 79, Logistic Regression = %84


Image

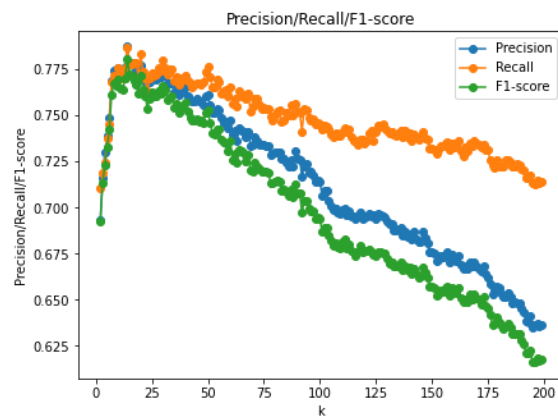


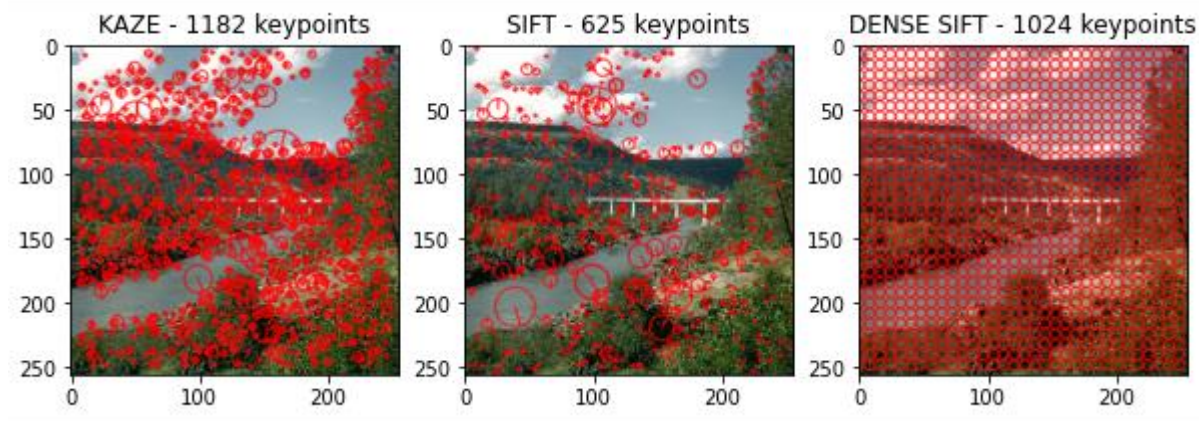


Keypoints:450




Keypoints:710





What I missed




Group	grade
1	8
2	9
3	6
4	10
5	9
6	9
7	8
8	9

S02

- We'll start with BoVW computed with Dense SIFT with a large enough codebook size
- We'll normalize descriptors
 - L2-norm, Power-norm, etc..
- Cross-validation
 - Sklearn functions: StratifiedkFold, GridsearchCV
- Spatial Pyramids
- SVM and kernels
 - Use sklearn standardScaler to project every dimension to [0, 1]!
 - linear kernel
 - RBF kernel
 - our own histogram intersection kernel
- OPTIONAL: Fisher Vectors (http://yael.gforge.inria.fr/tutorial/tuto_imgindexing.html)

Cross Validation




We use K-fold cross validation to pick suitable parameters for the classifier while **avoiding overfitting**.



For instance, in the KNN classifier we want to find the most suitable K (number of neighbours used to classify each descriptor)

We plotted the average of the accuracies (Y axis) obtained with 5 folds with respect to varying K from 1 to 100 (X axis).



Best average accuracy with $K = 30$!

So we obtain that individual SIFT descriptors are best classified using $K = 30$ (for another kind of descriptor the 'best' K will vary)

Hyperparameter optimization

Journal of Machine Learning Research 13 (2012) 281-305

Submitted 3/11; Revised 9/11; Published 2/12

Random Search for Hyper-Parameter Optimization

James Bergstra
Yoshua Bengio

Département d'Informatique et de recherche opérationnelle
Université de Montréal
Montréal, QC, H3C 3J7, Canada

JAMES.BERGSTRA@UMONTREAL.CA
YOSHUA.BENGIO@UMONTREAL.CA

Editor: Leon Bottou


Abstract

Grid search and manual search are the most widely used strategies for hyper-parameter optimization. This paper shows empirically and theoretically that randomly chosen trials are more efficient for hyper-parameter optimization than trials on a grid. Empirical evidence comes from a comparison with a large previous study that used grid search and manual search to configure neural networks and deep belief networks. Compared with neural networks configured by a pure grid search, we find that random search over the same domain is able to find models that are as good or better within a small fraction of the computation time. Granting random search the same computational budget, random search finds better models by effectively searching a larger, less promising configuration space. Compared with deep belief networks configured by a thoughtful combination of manual search and grid search, purely random search over the same 32-dimensional configuration space found statistically equal performance on four of seven data sets, and superior performance on one of seven. A Gaussian process analysis of the function from hyper-parameters to validation set performance reveals that for most data sets only a few of the hyper-parameters really matter, but that different hyper-parameters are important on different data sets. This phenomenon makes grid search a poor choice for configuring algorithms for new data sets. Our analysis casts some light on why recent "High Throughput" methods achieve surprising success—they appear to search through a large number of hyper-parameters because most hyper-parameters do not matter much. We anticipate that growing interest in large hierarchical models will place an increasing burden on techniques for hyper-parameter optimization; this work shows that random search is a natural baseline against which to judge progress in the development of adaptive (sequential) hyper-parameter optimization algorithms.

Keywords: global optimization, model selection, neural networks, deep learning, response surface modeling

1. Introduction

The ultimate objective of a typical learning algorithm \mathcal{A} is to find a function f that minimizes some expected loss $\mathcal{L}(x; f)$ over i.i.d. samples x from a natural (grand truth) distribution \mathcal{G}_x . A learning algorithm \mathcal{A} is a functional that maps a data set $\mathcal{X}^{(\text{train})}$ (a finite set of samples from \mathcal{G}_x) to a function



Continuous hyperparameter: distribution over possible values

generate random variable

Discrete hyperparameter: list of discrete choices
random selection
(without replacement if all discrete)

Set the number of trials



Bergstra, James, and Yoshua Bengio. "Random search for hyper-parameter optimization." *Journal of Machine Learning Research* 13.Feb (2012): 281-305.

```
10 params = {
11     "C": np.arange(0.001, 1, 0.01),
12     "max_iter": np.arange(50, 550, 50)
13 }
14
15 lg = LogisticRegression(solver='liblinear')
16
17 lg_grid = GridSearchCV(lg, params, cv=8,
18     scoring=["accuracy", "f1_macro"], refit="accuracy", return_train_score=True)
19 lg_grid.fit(visual_words, train_labels)
```


Spatial Pyramids



Spatial Pyramids



Histogram Intersection kernel

```
def histogramIntersection(M, N):
```

$$K_{int}(A, B) = \sum_{i=1}^m \min\{a_i, b_i\}.$$

```
    return K_int
```

Tasks to do

Improve the BoVW code with:

- Dense SIFT (with tiny steps and different scales!)
- L2-norm - power norm
- SVM classifier
- StandardScaler
- Cross-validation
- Linear, RBF and histogram intersection kernels
- Spatial Pyramids
- Fisher Vectors (OPTIONAL)

Deliverable

- A **single Python notebook file per group** reporting all the work done,
 - with the different experiments,
 - code,
 - plots,
 - explanations, etc.
 - **EVERYTHING EXECUTED!**
- To deliver by Monday 17th @ 10 A.M.
 - Please, state clearly your group.

Warning: provided code might not work out of the box depending on the used versions (OpenCV, numpy, sklearn...) do not panic, and read the documentation