







# Neural networks for graph-structured data

Master in Computer Vision Barcelona

Adriana Romero

adrianars@fb.com

# Motivation



Graphs are everywhere!


# Outline

- Notation & problem formulation
- Node Embeddings:
  - Per-node classifier
  - Indirectly injecting structural information
- Graph Neural Networks
  - Neural Message Passing
  - Basic GNN
  - GCN
  - GAT
  - Updates and over-smoothing
- Computer vision applications
- Wrap Up

# Notation & problem formulation

# Notation (1)

A graph  $\mathcal{G}$  is defined as a set of nodes (vertices)  $\mathcal{V}$  and a set of edges  $\mathcal{E}$  between nodes:  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$




Graphs are often represented by means of an **adjacency matrix  $A$** :

- $A$  can represent directed and undirected graphs (symmetric).
- $A \in \{0, 1\}^{|\mathcal{V}| \times |\mathcal{V}|}$  to represent presence/absence of edges.
- $A \in \mathbb{R}^{|\mathcal{V}| \times |\mathcal{V}|}$  to represent graphs with weighted edges.
- $A \in \mathbb{R}^{|\mathcal{V}| \times |\mathcal{V}| \times |\mathcal{R}|}$  to represent multi-relation graphs (with edges representing types of interactions)

# Notation (2)

A graph  $\mathcal{G}$  is defined as a set of nodes (vertices)  $\mathcal{V}$  and a set of edges  $\mathcal{E}$  between nodes:  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$




Graphs are often represented by means of an **adjacency matrix**  $A$ .

Nodes may be characterized by a matrix of node features  $F \in \mathbb{R}^{|\mathcal{V}| \times m}$ , representing e.g. attribute information.

# Problem formulation

Machine learning with graphs can also follow supervised, unsupervised or reinforcement learning paradigms.




Among the most common tasks, we can find:

- Node classification
- Relation prediction
- Community detection
- Graph classification

# Problem formulation

Machine learning with graphs can also follow supervised, unsupervised or reinforcement learning paradigms.



Among the most common tasks, we can find:


- **Node classification (transductive vs inductive)**
- Relation prediction
- Community detection
- Graph classification



Node embeddings

# Per-node classifier


Classify single node features, ignoring graph structure.



**How can we leverage structural information?**

# Node embeddings (1)


Represent each node with a feature vector that preserves the structure of the graph.



# Node embeddings (2)

Node embedding architectures:


- **encode** nodes into an **embedding vector** (feature vector).
- **enforce** some **structural constraint** in the **embedding space** or **decode** node embeddings into information about node and/or its neighborhood



Embeddings are used for downstream tasks.

# Node embeddings (3)


Constraining the learnt node embeddings to be close/distant depending on the presence of edges  
(Weston et al., 2008, extending Belkin and Niyogi, 2006 to neural networks)



Training: node classification  
loss + similarity constraint


# Node embeddings (4)

Decoding structural information (LINE – Tang et al. 2015)




# Node embeddings (5)

Alternatively to augmenting the loss with structural constraints, learn node structural features from random walks (DeepWalk – Perozzi et al. 2014, Node2Vec – Grover et al. 2016)



# Node embeddings (5)


Alternatively to augmenting the loss with structural constraints, learn node structural features from random walks (DeepWalk – Perozzi et al. 2014, Node2Vec – Grover et al. 2016)






# Node embeddings (5)

Alternatively to augmenting the loss with structural constraints, learn node structural features from random walks (DeepWalk – Perozzi et al. 2014, Node2Vec – Grover et al. 2016)



# Node embeddings (5)

Alternatively to augmenting the loss with structural constraints, learn node structural features from random walks (DeepWalk – Perozzi et al. 2014, Node2Vec – Grover et al. 2016)




1. Start with random node features  $\phi_u$
2. Sample a random walk  $\mathcal{W}_u$  from  $u$
3. Update  $\phi_u$  to maximize  $P(v|\phi_u)$ , where  $v \in \mathcal{W}_u$ .

A good representation of a node should allow us to easily predict the nodes that *surround* it.

# Node embeddings (6)

What if we had access to node features and labels? (Planetoid – Yang et al. 2016)



$\hat{\gamma}_{uv}$  predicts whether node  $u$  and  $v$  have the same label or different label.

# Node embeddings (?)

All methods covered so far used a **classifier** that classifies each node **independently**, with **graph structure** injected only **indirectly**, through the learnt embeddings.



**How can we directly leverage the graph structure when computing node features?**

# Deep Learning toolbox (1)

**CNNs ?** They work well on data defined in n-D grids



2D grid



1D grid



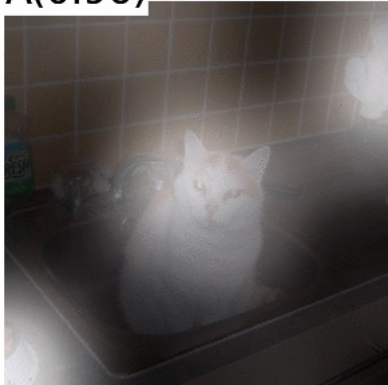
*"I like cats and dogs."*

# Deep Learning toolbox (2)

**RNNs?** They are well-defined over sequences.

image captioning

A(0.96)




from <http://kelvinxu.github.io/projects/capgen.html>

video frame classification



from <http://cs.stanford.edu/people/karpathy/deepvideo/>


image generation  
(sequential processing of output)



from <https://github.com/jbornschein/draw>

# Deep Learning toolbox (3)

MLPs ?



Depends on ordering of nodes in the adjacency matrix, the model would not be permutation equivariant.

# Graph Neural Networks (GNN)





# GNN (1)

**Key idea:** compute node representations that depend on the structure of the graph as well as node features. (Gori et al. 2005, Scarselli et al. 2009)

*Note: The same model has been derived as a generalization of convolutions to non-Euclidean data (Bruna et al., 2014).*

## Neural Message Passing



# GNN (2)

Formally, message passing updates are expressed as

Note that we can reformulate with self loops by considering  $\mathbf{H}_u^{(k-1)}$  within  $AGGREGATE^{(k)}$ .

$$\mathbf{H}_u^{(k)} = UPDATE^{(k)}(\mathbf{H}_u^{(k-1)}, \underbrace{AGGREGATE^{(k)}(\{\mathbf{H}_v^{(k-1)}, \forall v \in \mathcal{N}(u)\})}_{\text{message}})$$

where  $UPDATE$  and  $AGGREGATE$  are arbitrary differentiable functions (e.g. neural nets).

As iterations progress, each node contain more information from further reaches of the graph.

# GNN (3)

Formally, message passing updates are expressed as

$$\mathbf{H}_u^{(k)} = \text{UPDATE}^{(k)}(\mathbf{H}_u^{(k-1)}, \underbrace{\text{AGGREGATE}^{(k)}(\{\mathbf{H}_v^{(k-1)}, \forall v \in \mathcal{N}(u)\})}_{\text{message}})$$

where *UPDATE* and *AGGREGATE* are arbitrary differentiable functions (e.g. neural nets).

For basic GNNs:

$$\mathbf{H}_u^{(k)} = \sigma(\mathbf{W}_{self}^{(k)} \mathbf{H}_u^{(k-1)} + \underbrace{\mathbf{W}_{neigh}^{(k)} \sum_{v \in \mathcal{N}(u)} \mathbf{H}_v^{(k-1)}}_{\text{Sum of neighbor features}} + \mathbf{b}^{(k)})$$

# GNN (3)

Formally, message passing updates are expressed as

$$\mathbf{H}_u^{(k)} = \text{UPDATE}^{(k)}(\mathbf{H}_u^{(k-1)}, \underbrace{\text{AGGREGATE}^{(k)}(\{\mathbf{H}_v^{(k-1)}, \forall v \in \mathcal{N}(u)\})}_{\text{message}})$$

where *UPDATE* and *AGGREGATE* are arbitrary differentiable functions (e.g. neural nets).

For basic GNNs:  $\mathbf{H}^{(k)} = \sigma((\mathbf{A} + \mathbf{I})\mathbf{H}^{(k-1)}\mathbf{W}^{(k)})$

# GNN (3)

Formally, message passing updates are expressed as

$$\mathbf{H}_u^{(k)} = \text{UPDATE}^{(k)}(\mathbf{H}_u^{(k-1)}, \underbrace{\text{AGGREGATE}^{(k)}(\{\mathbf{H}_v^{(k-1)}, \forall v \in \mathcal{N}(u)\})}_{\text{message}})$$

where *UPDATE* and *AGGREGATE* are arbitrary differentiable functions (e.g. neural nets).

For basic GNNs:

$$\mathbf{H}_u^{(k)} = \sigma(\mathbf{W}_{self}^{(k)} \mathbf{H}_u^{(k-1)} + \underbrace{\mathbf{W}_{neigh}^{(k)} \sum_{v \in \mathcal{N}(u)} \mathbf{H}_v^{(k-1)}}_{\text{Sum of neighbor features}} + \mathbf{b}^{(k)})$$

highly sensitive to node degree

# Graph Convolutional Networks (GCN)

**Key idea:** neighborhood normalization motivated by spectral graph theory, resulting in first order approximation of a spectral graph convolution. (Kipf and Welling, 2017)

Message defined as:

$$\sum_{v \in \mathcal{N}(u)} \frac{\mathbf{H}_v^{(k-1)}}{\sqrt{|\mathcal{N}(u)| |\mathcal{N}(v)|}}$$

and so,

$$\mathbf{H}_u^{(k)} = \sigma \left( \mathbf{W}^{(k)} \left( \sum_{v \in \mathcal{N}(u) \cup \{u\}} \frac{\mathbf{H}_v^{(k-1)}}{\sqrt{|\mathcal{N}(u)| |\mathcal{N}(v)|}} \right) \right)$$

This formulation gives **less weight** to messages from neighbors with **higher degrees**.

# Other aggregate functions

**Key idea:** use any permutation invariant function as aggregator (e.g. set aggregators).

In particular, the message could be defined by means of a **set pooling** operation (Zaheer et al., 2017)

$$MLP_{\theta} \left( \sum_{v \in \mathcal{N}(u)} MLP_{\phi} (\mathbf{H}_v^{(k-1)}) \right)$$

by combining a set of node features into a single embedding.

Note that the **sum** could be replaced by **max** or **min** operation. (Qi et al., 2017)

# Graph Attention Networks - GAT (1)

**Key idea:** weigh the neighbors influence at aggregation time, by learning attention weights. (Velickovic et al., 2018)

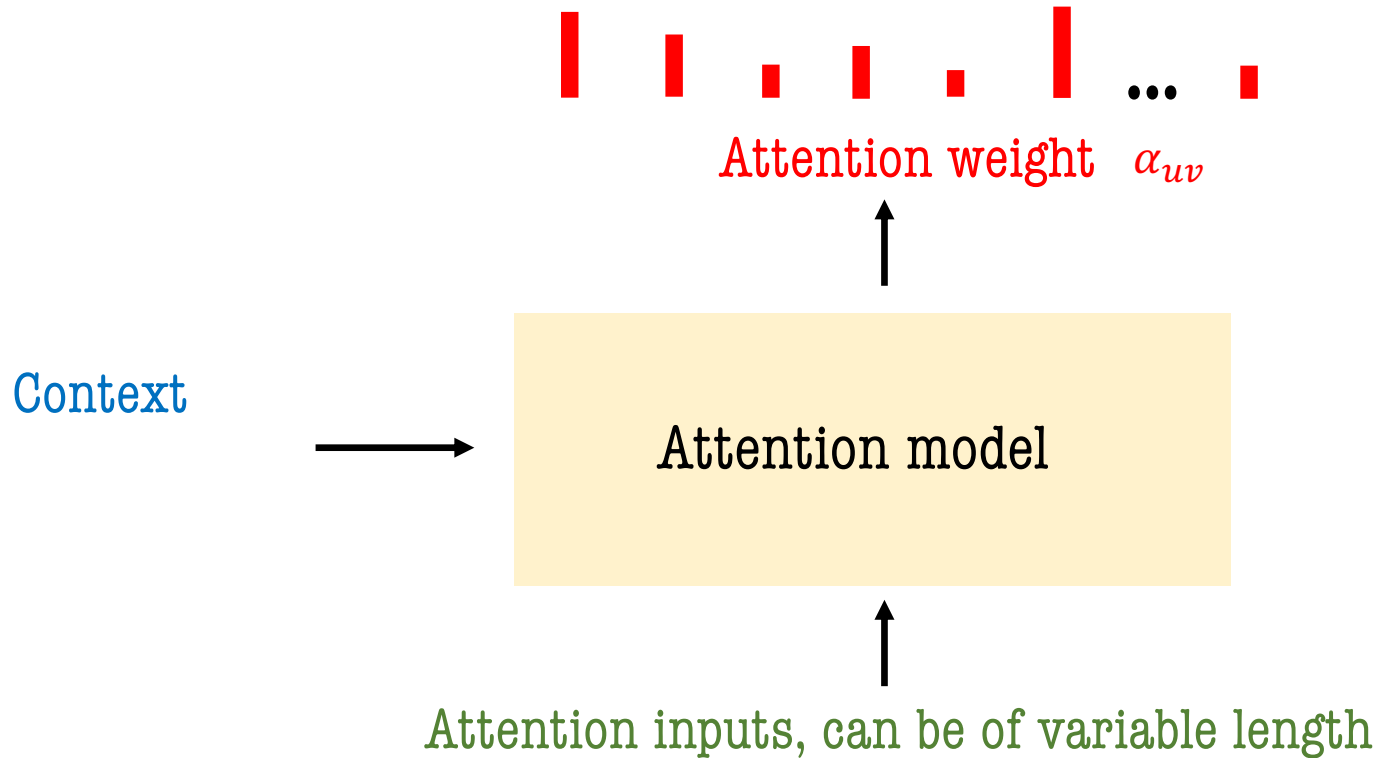
Message defined as:  $\sum_{v \in \mathcal{N}(u) \cup \{u\}} \alpha_{uv} \mathbf{H}_v^{(k-1)}$

**How do we compute  $\alpha_{uv}$  ?**



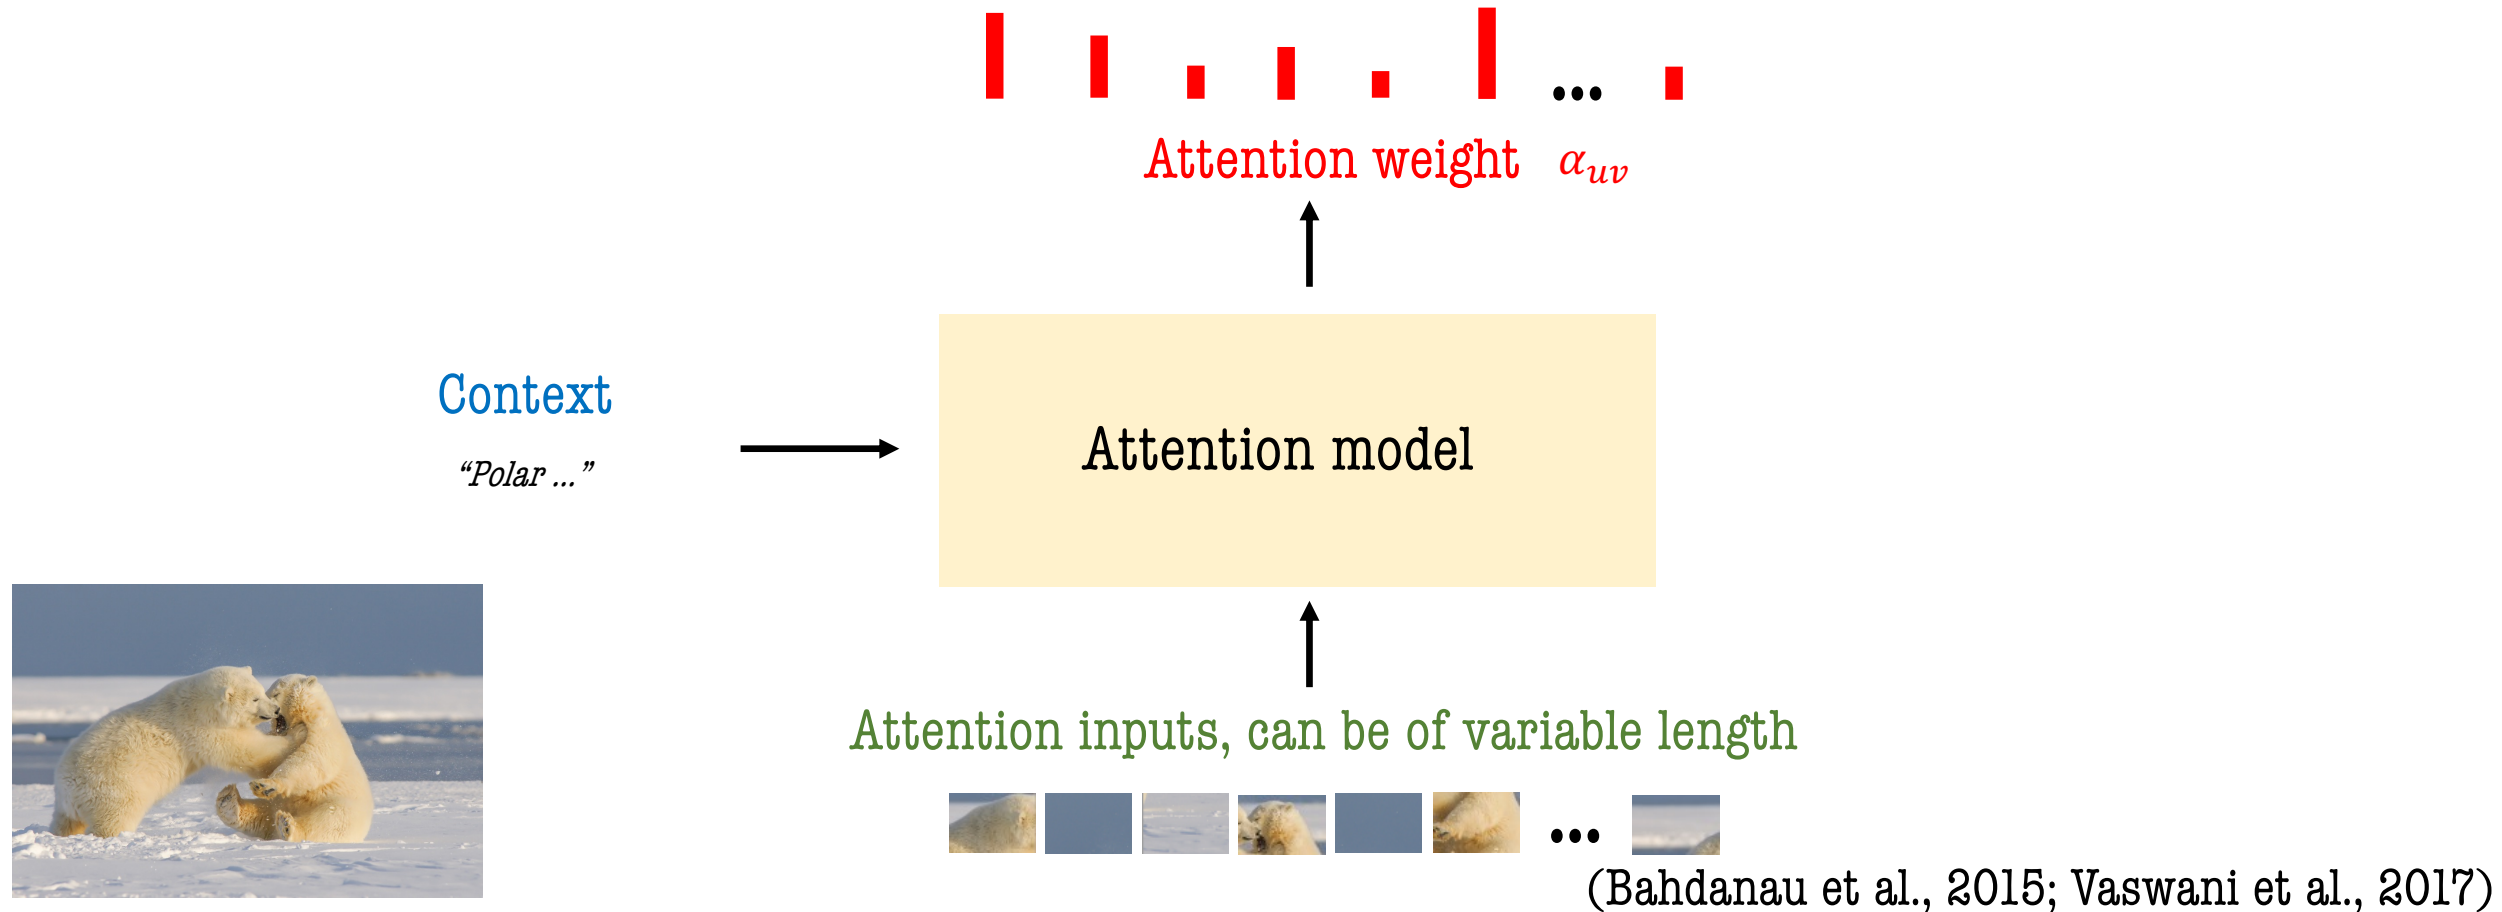
# GAT (2)

Attention allows us to focus and select the most pertinent pieces of information.



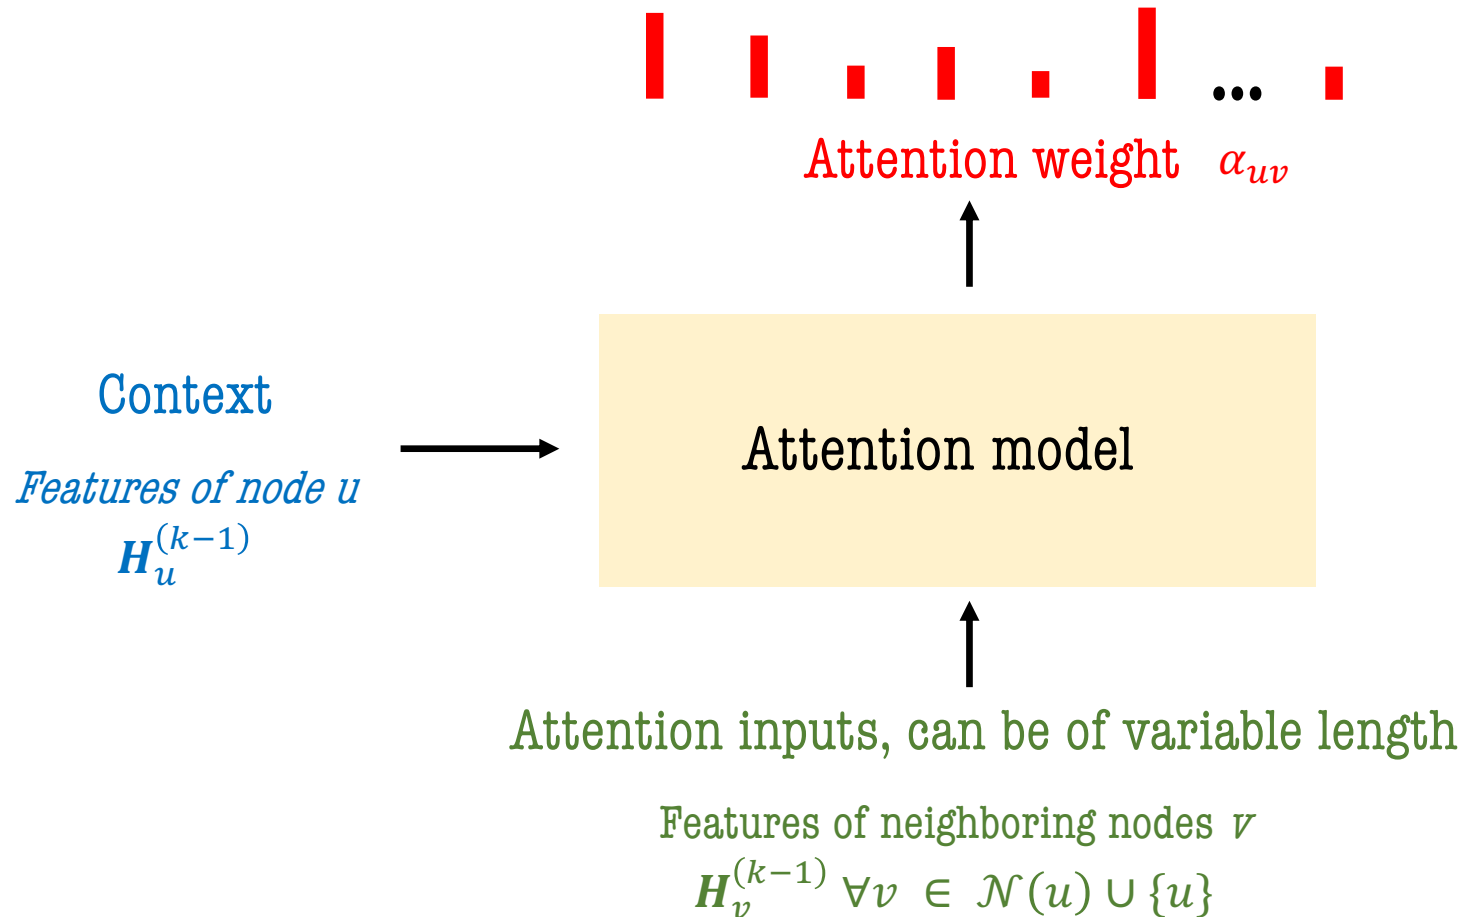
# GAT (2)

Attention allows us to focus and select the most pertinent pieces of information.



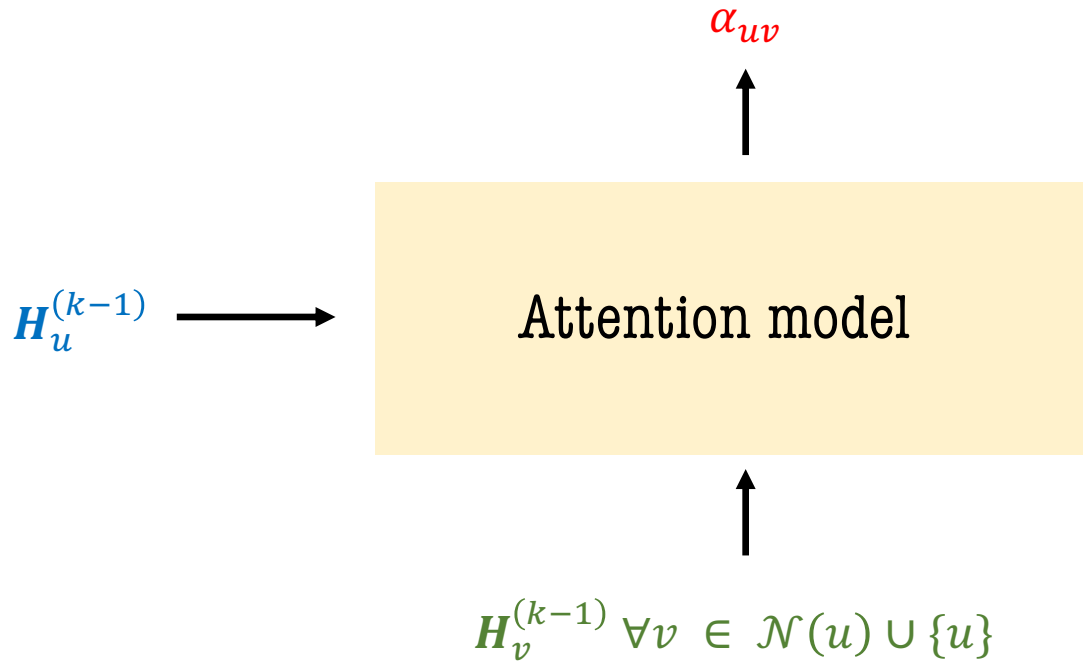
# GAT (2)

Attention allows us to focus and select the most pertinent pieces of information.



(Velickovic et al., 2018)

# GAT (3)



1. Combine attention input and context information

$$s_{uv} = f_{\theta} \left( H_u^{(k-1)}, H_v^{(k-1)} \right)$$

2. Apply softmax to obtain attention weights  $\alpha_{uv}$

$$\alpha_{uv} = \frac{\exp(s_{uv})}{\sum_{v' \in \mathcal{N}(u) \cup \{u\}} \exp(s_{uv'})}$$


3. Apply attention coefficients to input and aggregate

$$\sum_{v \in \mathcal{N}(u) \cup \{u\}} \alpha_{uv} H_v^{(k-1)}$$

(Velickovic et al., 2018)

# GAT (4)

Following the Transformer (Vaswani et al., 2017), GAT employs multiple attention heads:




$$\mathbf{H}_u^{(k)} = \sigma \left( \sum_{v \in \mathcal{N}(u) \cup \{u\}} \alpha_{uv} \mathbf{H}_v^{(k-1)} \right)$$

For one attention head.

The output of multiple heads can then be combined via e.g. concatenation.

# GAT-POS

Note that GAT formulation treats nodes in a neighborhood as a set, and so structural information of the nodes is lost. GAT-POS augments GAT with positional information.



# UPDATE & over-smoothing

Recall the basic GNN update:

$$\mathbf{H}_u^{(k)} = \sigma(\underbrace{W_{self}^{(k)} \mathbf{H}_u^{(k-1)} + W_{neigh}^{(k)} \sum_{v \in \mathcal{N}(u)} \mathbf{H}_v^{(k-1)}}_{\text{Linear combination of the aggregated neighboring information and the previous representation of the node}}) + \mathbf{b}^{(k)}$$

Linear combination of the aggregated  
neighboring information and the previous  
representation of the node

May result in over-smoothing.

**How can we address over-smoothing?**

# **Skip connections (1)**

To preserve information from previous message passing iterations in the update function.



# Skip connections (2)


Review of image classification architectures with skip connections

## CNNs




$$F(x) = \left( \begin{array}{c} \text{Conv.} \\ + \\ \text{ReLU} \end{array} \right) \times n$$

## ResNets




$$G(x) = \left( \begin{array}{c} \text{BN} \\ + \\ \text{ReLU} \\ + \\ \text{Conv.} \end{array} \right) \times n$$

## Highway Nets



$$H(x) = \left( \begin{array}{c} \text{Conv.} \\ + \\ \text{ReLU} \\ \text{or sigmoid} \end{array} \right)$$

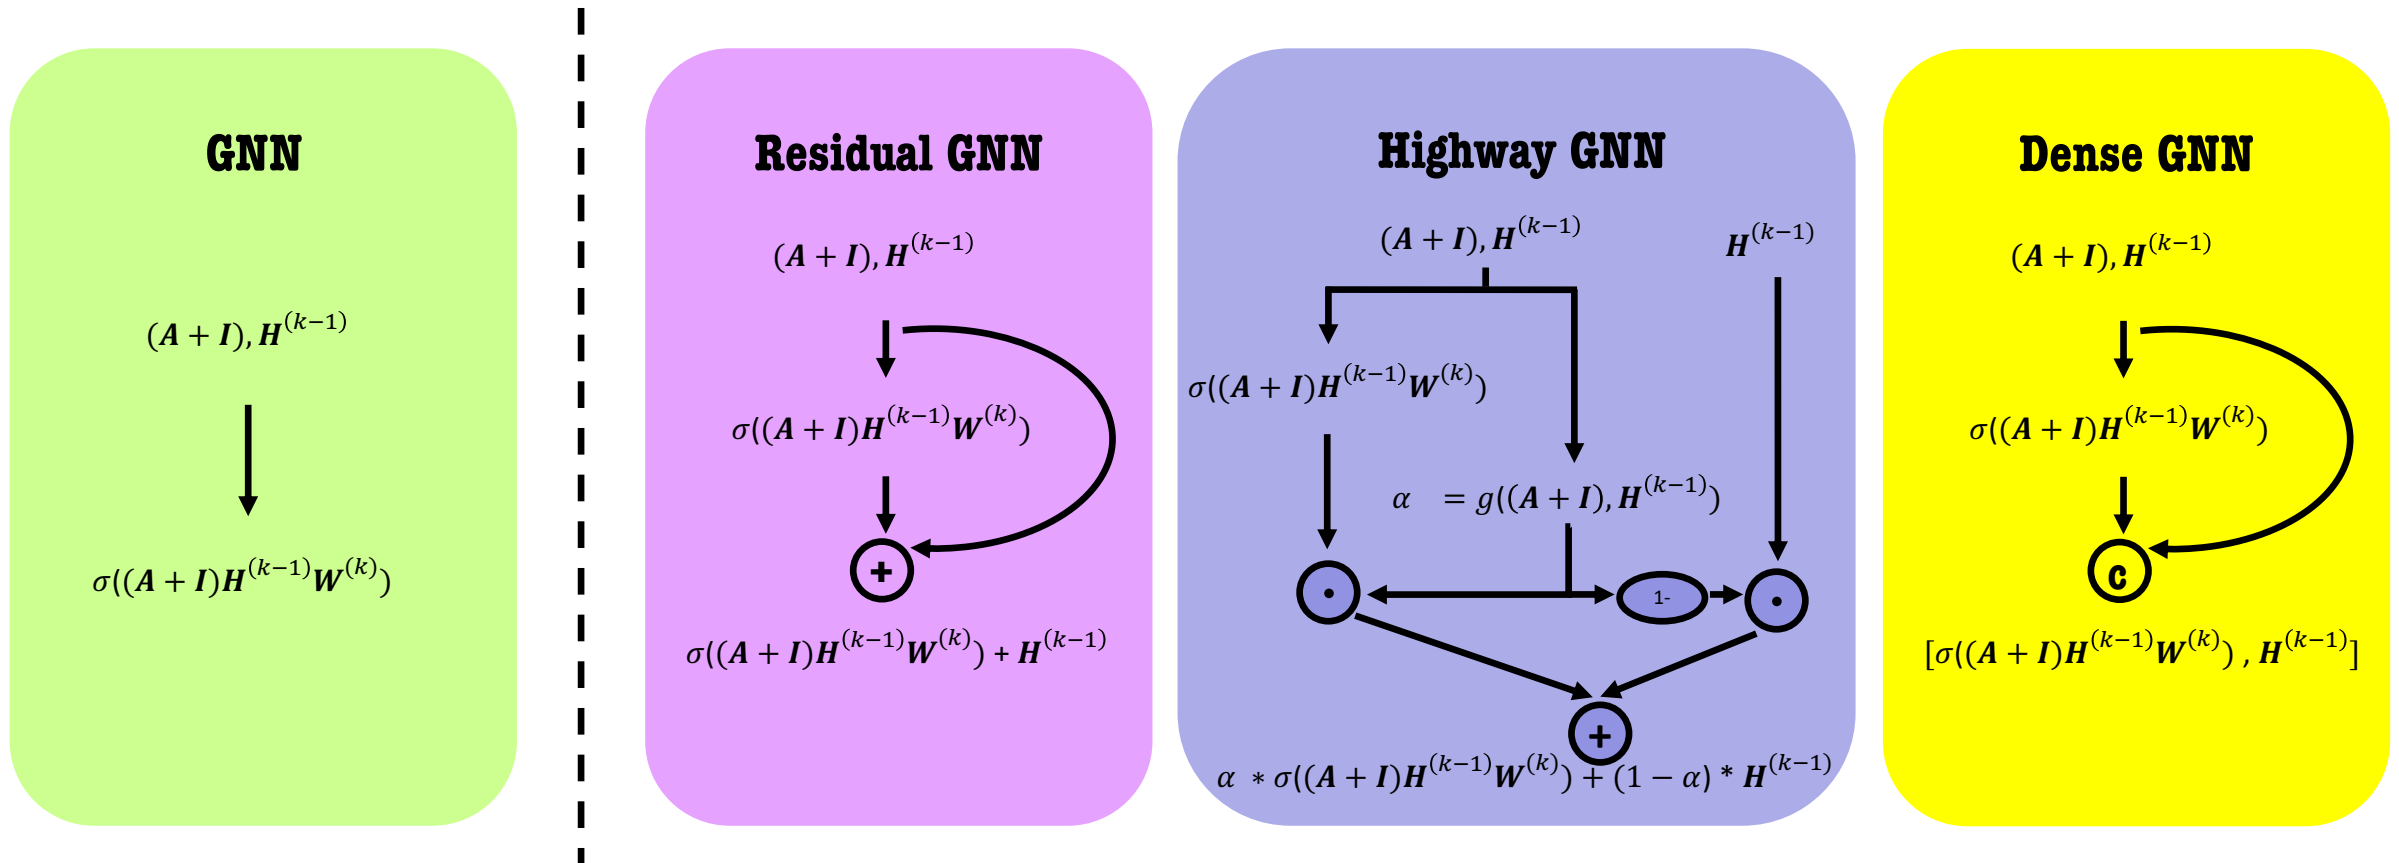
## DenseNets



$$H(x) = \left( \begin{array}{c} \text{BN} \\ + \\ \text{ReLU} \\ + \\ \text{Conv.} \\ + \\ \text{Dropout} \end{array} \right)$$

# Skip connections (3)

GNN also have their skip connection-based counterparts.




(Pham et al., 2017) (Hamilton et al., 2017 a& Xu et al., 2018)

# Zero-neighbor update

**Key idea:** Update only a part of the node's features with the neighbors' information.

Update defined as:  $\mathbf{H}^k = \sigma([\mathbf{A}\mathbf{H}'_{0:i}^k \parallel \mathbf{I}\mathbf{H}'_{i:}^k] + \mathbf{b}^k)$

where  $\mathbf{H}'^k = \mathbf{H}^{(k-1)}\mathbf{W}^{(k)}$



Providing a soft middle ground between **full** exchange of information and **no communication**

# Graph Isomorphism Networks (GIN)

**Key idea:** maps different node neighborhoods to different feature vectors to obtain high discriminative/representational power.

Update defined as:  $\mathbf{H}_u^{(k)} = \text{MLP}^{(k)}((1 + \epsilon^{(k)})\mathbf{H}_u^{(k-1)} + \sum_{v \in \mathcal{N}(u)} \mathbf{H}_v^{(k-1)})$

 Learnable parameter.

# RNN-based updates

**Key idea:** GNN updates formulated in terms of RNN-based updates, where

- the RNN hidden state  $h_t$  becomes the node's hidden state
- the RNN observation (input)  $x_t$  becomes the aggregated message

(Li et al., 2016 and Selsam et al., 2018)

$$\mathbf{H}_u^{(k)} = \text{RNN}^{(k)}(\mathbf{H}_u^{(k-1)}, \text{AGGREGATE}^{(k)}(\{\mathbf{H}_v^{(k-1)}, \forall v \in \mathcal{N}(u)\}))$$

- RNN update function – eg LSTM/GRU update
- Parameters are shared across nodes (same RNN cell used to update each node)
- Update function shared across (message passing) layers

# Notes


- All GNNs presented thus far **iteratively update node features** that can be used for any **downstream task**.
- In the case of **node classification**:
  - the non-linearity  $\sigma$  used in the last update (layer) is the **softmax**.
  - the **loss function** to train the model is analogous to the semantic segmentation one: **per node cross-entropy loss**.
- In the case of **graph classification**,
  - a **pooling** operation is often used to **aggregate** the information across **all nodes**.
  - the **pooled features** are fed to a **classifier**.
  - the **loss function** to train the model is analogous to the one of image classification: **per graph cross-entropy loss**.
- Most of the presented models can be augmented to handle **multi-relational graphs**, by separately **aggregating information across different edge types**.

# Computer Vision Applications

# 3D vision (1)

**Key idea:** Exploit mesh representations in the context of 3D visual understanding.


Why meshes?




(a) Voxels  
(262, 144 units)



(b) Point cloud  
(30, 000 points)



(c) Uniform mesh  
(2416 vertices)




(d) Adaptive mesh  
(120 vertices)

GNNs for mesh reconstructions: see e.g. Wang et al. 2018, Smith et al. 2019, Gkioxari et al. 2019, Smith et al. 2021.



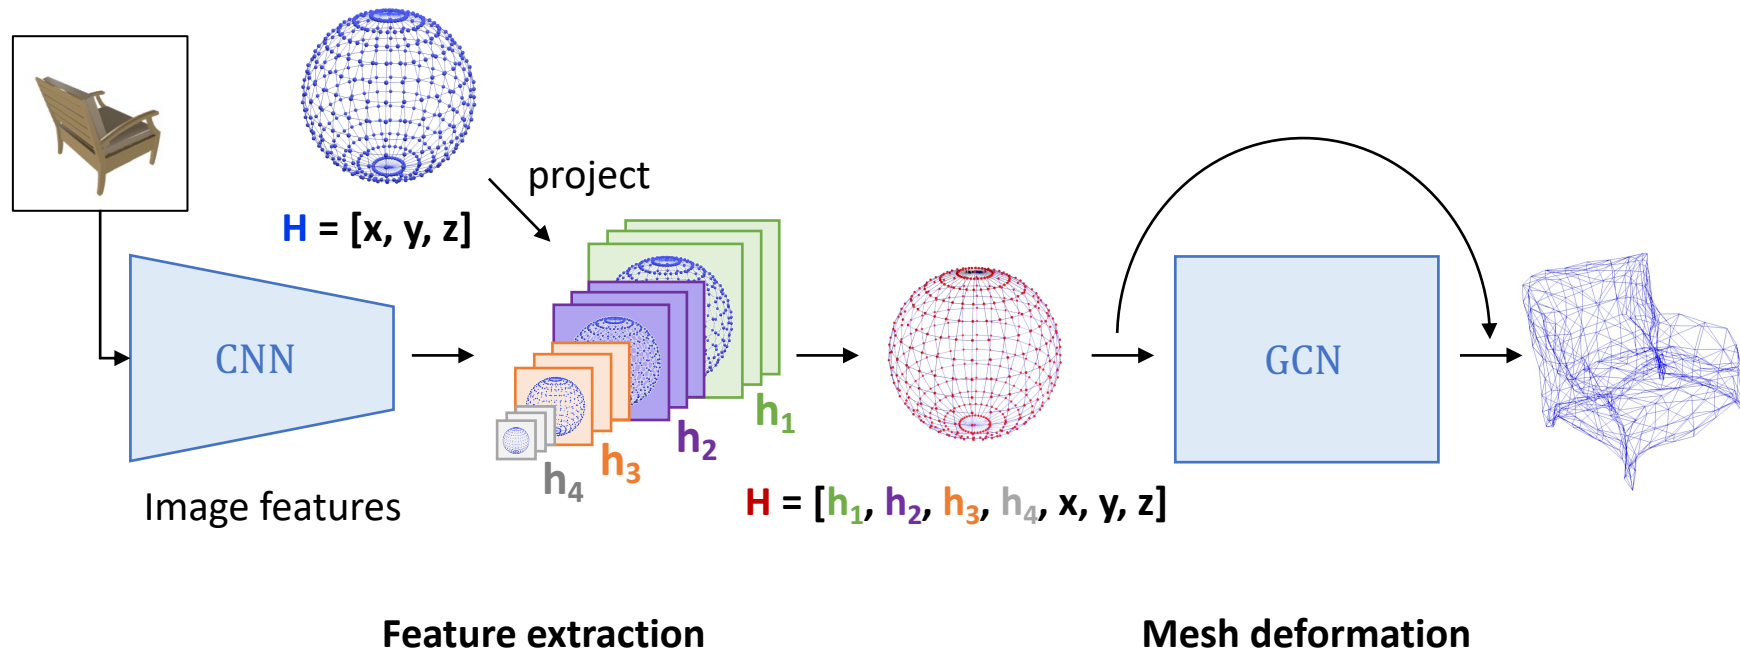
# 3D vision (2)

Often posed as shape reconstruction from a **single view** image:



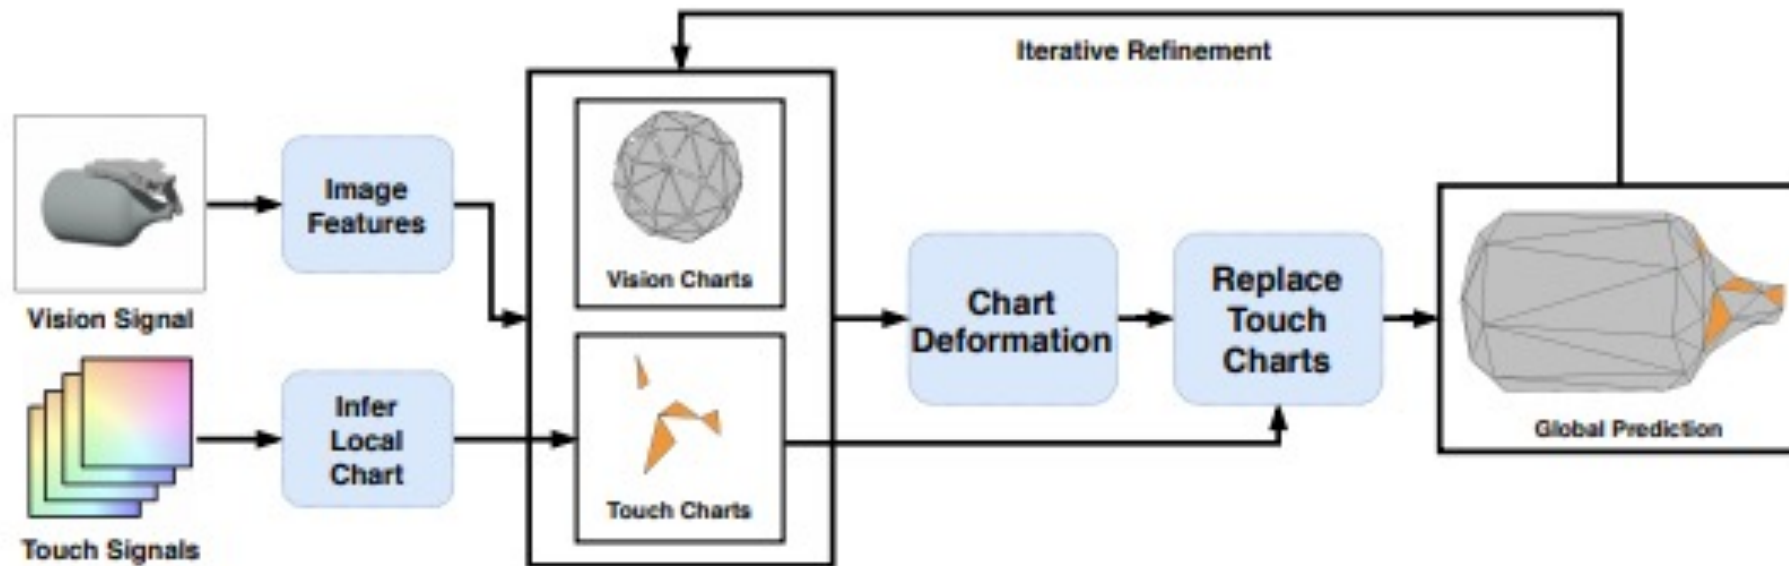
Goal: Inferring 3D shape from a single view image of an object.

# 3D vision (3)



Trained to minimize the Chamfer loss: 
$$\sum_{p \in S} \min_{q \in \hat{S}} \|p - q\|_2^2 + \sum_{q \in \hat{S}} \min_{p \in S} \|p - q\|_2^2$$



# 3D vision (4)



Goal: Inferring 3D shape from a single view image of an object and tactile information.

# Relational representations

**Key idea:** Represent the visual world relations (e.g. positional, conceptual or from attributes) by means of graph.



**Tasks:** Infer relations (graph) from images or videos, exploit relations to improve model generalization, and generate images from graphs, among others.

Wrap Up

# Wrap Up

## Problem formulation

- What are graphs?
- How do we represent them?
- What kind of problems can we tackle?

## Node embeddings

- Per-node classifier (ignoring graph structure)
- Indirect injection of structural information

## Graph Neural Networks

- Neural Message Passing
  - Basic GNN
  - GCN
  - GAT

*Different aggregate functions.*

- Skip connections
- 0-neighbor updates
- GIN
- RNN updates

*Different update functions.*

- How to train those models?
- Multi-relational graphs

## Computer Vision Applications

- 3D Vision
- Relational representations
  - Inferring visual knowledge graph from data
  - Generating visual content from scene graphs

# References

(Weston et al., 2008) J. Weston, F. Ratle & R. Collobert; **Deep Learning via Semi-Supervised Embedding**, ICML, 2008.

(Perozzi et al., 2014) B. Perozzi, R. Al-Rfou & S. Skiena; **DeepWalk: Online Learning of Social Representations**, KDD, 2014.

(Tang et al., 2015) J. Tang, M. Qu, M. Wang, J. Yan & Q. Mei; **LINE: Large-scale Information Network Embedding**, WWW, 2015.

(Grover et al., 2016) A. Grover & J. Leskovek; **node2vec: Scalable Feature Learning for Networks**, KDD, 2016.

(Yang et al., 2016) Z. Yang, W. W. Cohen & R. Salakhutdinov; **Revisiting Semi-Supervised Learning with Graph Embeddings**, ICML, 2016.

(Gori et al., 2005) M. Gori, G. Monfardini & F. Scarselli; **A new model for learning in graph domains**, IJCNN, 2005.

(Scarselli et al., 2009) F. Scarselli, M. Gori, A. C. Tsoi, M. Hagenbuchner & G. Monfardini; **The Graph Neural Network Model**, TNNLS, 2009.

(Li et al., 2016) Y. Li, D. Tarlow, M. Brockschmidt & R. Zemel; **Gated Graph Neural Networks**, ICLR, 2016.

(Bruna et al., 2014) J. Bruna, W. Zaremba, A. Szlam & Y. LeCun; **Spectral Networks and Locally Connected Networks on Graphs**, ICLR, 2014.

(Kipf & Welling, 2017) T. Kipf and M. Welling; **Semi-Supervised Classification with Graph Convolutional Networks**, ICLR, 2017.

(Belkin and Niyogi, 2002) M. Belkin and P. Niyogi. **Laplacian eigenmaps and spectral techniques for embedding and clustering**, NIPS, 2002



(Hamilton et al., 2017) W. L. Hamilton, R. Ying & J. Leskovec; **Inductive Representation Learning on Large Graphs**, NIPS, 2017.

(Bahdanau et al., 2015) D. Bahdanau, K. Cho & Y. Bengio; **Neural Machine Translation by Jointly Learning to Align and Translate**, ICLR, 2015.

(Vaswani et al., 2017) A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser & I. Polosukhin; **Attention is all you need**, NIPS, 2017.

(Velickovic et al., 2018) P. Velickovic, G. Cucurull, A. Casanova, A. Romero. P. Lio & Y. Bengio; **Graph Attention Networks**, ICLR, 2018.

(Qi et al., 2017) C.R. Qi, H. Su, K. Mo, and L .J. Guibas; **Pointnet: Deep learning on point sets for 3d classification and segmentation**, In CVPR, 2017.

(Zaheer et al., 2017) M. Zaheer, S. Kottur, S.. Ravanbakhsh, B. Póczos, R. Salakhutdinov, and A. J. Smola; **Deep sets**, NIPS, 2017.

(Pham et al., 2017) T. Pham, T. Tran, D. Phung, and S. Venkatesh; **Column networks for collective classification**, AAAI, 2017.

(Xu et al, 2018) K. Xu, C. Li, Y. Tian, T. Sonobe, K. Kawarabayashi, and S. Jegelka; **Representation learning on graphs with jumping knowledge networks**, arXiv preprint arXiv:1806.03536, 2018.

(Defferrard et al., 2016) M. Defferrard, X. Bresson & P. Vandergheynst; **Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering**, NIPS, 2016.

(Duvenaud et al., 2015) D. K. Duvenaud, D. Maclaurin, J. Iparraguirre, R. Bombarell, T. Hirzel, A. Aspuru-Guzik & R. P. Adams; **Convolutional Networks on Graphs for Learning Molecular Fingerprints**, NIPS, 2016.

- (Smith et al., 2019) E. J. Smith, S. Fujimoto, A. Romero, and D. Meger; **GEOMetrics: Exploiting Geometric Structure for Graph-Encoded Objects**, ICML, 2019.
- (Selsam et al., 2018) D. Selsam, M. Lamm, B. Bunz, P. Liang, L. de Moura, and D. L. Dill; **Learning a sat solver from single-bit supervision**, arXiv preprint arXiv:1802.03685, 2018.
- (Wang et al., 2018) N. Wang, Y. Zhang, Z. Li, Y. Fu, W. Liu, and Y.-G. Jiang; **Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images**, ECCV, 2018.
- (Gkioxari et al., 2019) G. Gkioxari, J. Malik, and J. Johnson; **Mesh R-CNN**, ICCV, 2019.
- (Ashual and Wolf, 2019) O. Ashual and Lior Wolf; **Specifying Object Attributes and Relations in Interactive Scene Generation**, ICCV, 2019.
- (Johnson et al., 2018) J. Johnson, A. Gupta, and L. Fei-Fei; **Image Generation from Scene Graphs**, CVPR, 2018.
- (Xu et al., 2017) D. Xu, Y. Zhu, C. B. Choy, and L. Fei-Fei; **Scene Graph Generation by Iterative Message Passing**, CVPR, 2017.
- (Smith et al., 2020) E. Smith, R. Calandra, A. Romero, G. Gkioxari, D. Meger, J. Malik, M. Drozdal; **3D Shape Reconstruction from Vision and Touch**, NeurIPS, 2020.
- (Casanova et al., 2020) A. Casanova, M. Drozdal, A. Romero-Soriano; **Generating unseen complex scenes: are we there yet?**, arXiv preprint arXiv:2012.04027, 2020.
- (Ma et al., 2021) L. Ma, R. Rabbany, and A. Romero-Soriano; **Graph Attention Networks with Positional Embeddings**, PAKDD, 2021.
- (Xu et al., 2019) K. Xu, W. Hu, J. Leskovec, and S. Jegelka; **How powerful are graph neural networks?**, ICLR, 2019.

# **Neural networks for graph-structured data**

Master in Computer Vision Barcelona

Adriana Romero

[adrianars@fb.com](mailto:adrianars@fb.com)